Independent reconstruction of the coastal adaptation model of Wong et al. (2025): continuing research report

Research-agent record prepared for Sophia Ludtke

13 October 2026

Abstract

Coastal cost distributions depend on the coupled climate response and the uncertainty assigned to exposure and adaptation. We independently reconstruct the written model lineage of Wong et al. (2025), without consulting its implementation or computational outputs. Scholarly source recovery supplies climate and carbon equations, Antarctic geometry and paleo constraints, coastal cost functions, and independent observational and geographic inputs. The resulting Julia calculations verify heat and carbon conservation, ice-response convergence, flood integration and economic accounting. A temperature-only calibration passes sampling diagnostics but fails held-out observations. Joint temperature–heat calibration improves prediction, and substituting the documented HadCRUT5 version raises median SSP2-4.5 warming from 1.97 to 2.22 K; the median remains below the AR6 best estimate of 2.7 K. Carbon fits are audited for optimization and predictive error separately. Three 250-member emissions-driven alternatives propagate through sea-level components; the HadCRUT4 and HadCRUT5 heat-constrained alternatives reach sampled coastal costs. The base warming-conditioned median incremental costs span 0.208–1.217 trillion 2019 USD PPP under annual 2010–2100 accounting. The regional ranking differs from the target's written findings, and the highest warming category contains only five members. HadCRUT5 increases this count to 22, with conditional effective size 13.6, without recovering the target distribution. Numerical coast-sampling errors and alternative closure sensitivities are reported separately from climate variation. Independent economic marginal draws change paired costs by -35% to $+85\%$ across six purposive paths. The historical heat-constrained sea-level sum rises at about 2.03–2.05 mm/year, compared with 2.60 mm/year in the observational record; that dependent consistency check is not independent validation. None of the thirteen numerical target figures is established as independently reproduced: the target joint probability model, geographic inputs and accounting remain unmatched. The original workflow schematic, verified mechanisms, explicitly labeled numerical alternatives and complete figure inventory document substantial continued progress without converting missing specifications into invented published results.

1 Objective and evidentiary boundary

Wong et al. [3] estimate global coastal adaptation costs and losses conditional on warming in 2100. Their model connects emissions, carbon cycling, temperature and ocean heat, global and local sea level, and segment-level coastal decisions. The scientific objective here is to reproduce the resulting conditional distributions from the documented mechanisms. It is not sufficient to draw curves resembling the published figures or to reproduce marginal parameter ranges while discarding their dependence.

The independence restriction permits scientific papers, their supplementary methods, and exogenous inputs from independent providers. Published figures and printed summary statistics are

comparison evidence only. No source code, processed input bundle, posterior sample, simulation file, or plotting script from the target authors was opened or used. An initial broad literature search displayed incidental repository-index snippets, including a data-loading excerpt. Those results were discarded, and subsequent searches focused on scientific documents and provider records. This exposure is retained in the research log rather than represented as complete isolation from search results.

The recovered sources include the target article and its ten-figure supplement, the SNEASY climate description [5], its carbon-cycle reference [6], BRICK and MimiBRICK descriptions [7, 8], the DAIS paper [9], and the CIAM mathematical supplement [10]. The December 2024 revision of the cited emissions preprint [4] was recovered through public scholarly file-version metadata. Its current DOI landing page serves a later 2025 version. Since the target citation does not pin a revision, this historical manuscript remains a version-qualified specification. Stable acquisition URLs and hashes are recorded in `docs/SOURCES.csv`.

2 Original experiment and recovered interfaces

The target source ensemble contains 100,000 states of the world, with 5,000 used for the coastal analysis. Global warming in 2100, relative to 1850–1900, defines four bins: at most 1.5 °C, above 1.5 and at most 2.5 °C, above 2.5 and at most 3.5 °C, and above 3.5 °C. Global sea level in 2100 uses a 1995–2014 reference. The baseline socioeconomic scenario is SSP4; SSP2 provides a sensitivity case. Monetary outcomes are in 2010 US dollars with a 4% annual discount rate [3].

The coastal model uses 12,148 DIVA segments and decadal calculations from 2010 through 2150. Planning intervals are 2010–2050, 2060–2100, and 2110–2150. Protection and retreat standards are associated with return periods of 1, 10,100, 1,000, and 10,000 years, alongside no adaptation. The paper specifies nine adaptation combinations. The complete 2022 scholarly supplement resolves the menu: protection at 10–10,000 years and retreat at 1–10,000 years. The nominal one-year level is annual maximum high tide, obtained for the GTSR experiment by scaling the DINAS one-year level by the GTSR/DINAS 100-year ratio. It is not the GEV quantile at probability zero.

The 2022 main methods specify first-period cost minimization, persistence of the selected strategy and return level, and later height augmentation. Actual adaptation already in place is not globally observed; the original calculation initializes an optimal adaptation state. Reconstructing this spin-up requires the same exposure and extreme-water inputs as the forward model. It cannot be inferred from a global cost total.

Figure1 identifies the implemented kernels and the missing links. These components can be exercised independently. The continuing experiments also connect independently calibrated climate and ice responses under explicitly alternative probability and accounting assumptions.

3 Independent mathematical reconstruction

3.1 Emissions and climate kernels

For annual emissions q_t in $\mathrm{Gt\,CO_2\,yr^{-1}}$, the target's printed equation is

$$q_t = \begin{cases} q_{t-1} + \gamma_g(\tau - t), & t \le \tau, \\ q_{t-1} - 2q_\tau\gamma_d e^{\gamma_d(t-\tau)}/[1 + e^{\gamma_d(t-\tau)}]^2, & t > \tau. \end{cases} \tag{1}$$

The implementation evaluates the second branch with negative exponential arguments to avoid overflow while preserving the algebra. The growth coefficient is interpreted in the annual recurrence's

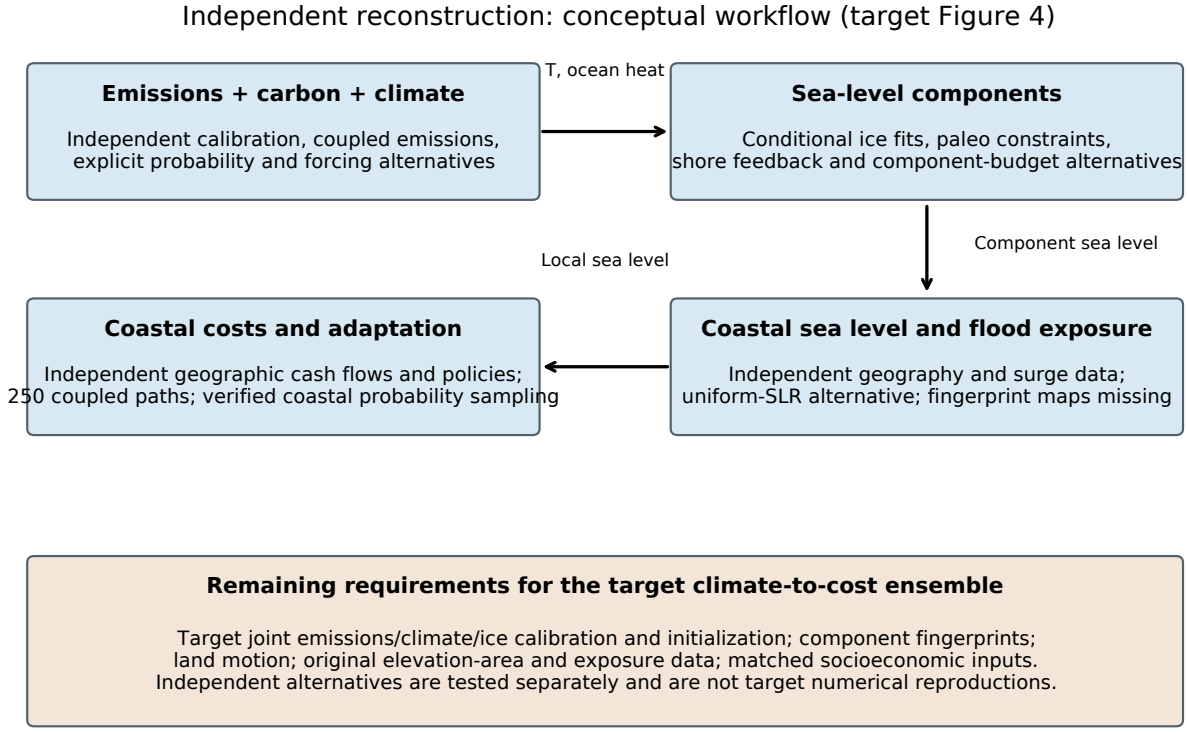


Figure 1: Original conceptual schematic corresponding to target Figure 4. Blue boxes identify implemented component calculations and their physical interfaces. The lower box lists unresolved inputs and model closure. The arrows describe physical exchanges implemented in explicitly alternative coupled experiments; unresolved target inputs are listed separately. This diagram is not counted as a reproduced numerical result.

time convention; a dimensional continuous-time growth coefficient would have different units. An integer peak year is required explicitly. The emissions prior allows continuous peak years, but the treatment of q_τ for a noninteger peak was not recovered.

The 2024 manuscript gives independent truncated-normal priors for growth, peak year, and decline: means $(0.004, 2070, 0.07)$, standard deviations $(0.0075, 25, 0.05)$, and bounds $[0.001, \infty)$, $[2030, 2200]$, and $[0.001, 0.2]$, respectively [4]. These are not the posterior distributions used in the target ensemble. The subsequent inversion uses expert cumulative-emissions information whose complete probability-density prescription was not recovered. Sampling these priors would therefore not reproduce the target emissions distribution.

Recoverable carbon kernels are net primary production $N = N_0[1 + \beta \ln(C/C_0)]$, with $N_0 = 60\,\mathrm{Gt\,C\,yr^{-1}}$ and $C_0 = 280\,\mathrm{ppm}$, and heterotrophic respiration $(a_1 D + a_2 S) Q_{10}^{\Delta T/10}$ [6]. Here D and S are detrital and soil carbon, and a_1, a_2 are turnover rates. The continuation recovers terrestrial allocation and carbonate equilibrium, allowing the closed, explicitly alternative calculation in Section7. The implemented CO_2 forcing is $3.7 \log_2(C/C_0)\,\mathrm{W\,m^{-2}}$ [5].

The ocean verification operator solves $\partial_t T = \kappa \partial_{zz} T$ on a one-dimensional column with prescribed surface temperature and zero bottom heat flux. Cell-centred finite volumes use half-cell distance at the upper boundary. Backward Euler time integration ensures that the column heat change equals the top flux evaluated at the new time, to roundoff. Multiplication by volumetric heat capacity converts the tested temperature integral to heat per unit area. This independently chosen discretization is not asserted to match the original SNEASY solver.

3.2 Sea-level components and spatial transformation

The glacier contribution S satisfies the documented reservoir relation

$$\frac{dS}{dt} = \beta (T - T_{\mathrm{eq}})(1 - S/V_0)^n, \qquad T_{\mathrm{eq}} = -0.15\,^\circ\mathrm{C}. \tag{2}$$

Temperature is referenced to 1850–1870, V_0 is initial glacier volume expressed as metres of sea-level equivalent, and β has units $\mathrm{m\,K^{-1}\,yr^{-1}}$ [7]. The numerical domain requires $S \le V_0$ for fractional n. An integration step that leaves this domain raises an error; no undocumented volume clipping is introduced.

Greenland volume V, also in metres of sea-level equivalent, follows

$$V_{\mathrm{eq}} = cT + b, \qquad \frac{dV}{dt} = (V_{\mathrm{eq}} - V)(\alpha T + \beta). \tag{3}$$

The temperature reference is 1961–1990, and the volume loss relative to its initial state yields sea-level rise. The rate $\alpha T + \beta$ must be nonnegative. Constant temperature gives an exponential relaxation and therefore an independent verification solution.

The target lineage uses a heat-driven thermal-expansion response, $\Delta S_{\mathrm{TE}} = k_H \Delta H$, not the earlier temperature-relaxation emulator [11, 4]. The coefficient k_H in $\mathrm{m\,J^{-1}}$ remains an explicit input. If the density expansion parameter α_ρ has units $\mathrm{kg\,m^{-3}\,K^{-1}}$, dimensional mass and heat accounting gives $k_H = \alpha_\rho / (A C_p \rho^2)$. This is a dimensional derivation, not recovery of a calibrated numerical coefficient. Supplying a conventional volumetric coefficient in $\mathrm{K^{-1}}$ in that expression would be a unit error.

The cited physical reference supplies spatial ocean averages, rather than the missing calibrated prior. Roquet et al. Table 1 [2] gives density expansion coefficients of 0.16, 0.18, 0.24 and 0.05 $\mathrm{kg\,m^{-3}\,K^{-1}}$ for the global ocean, horizontal-density-gradient-weighted ocean, surface and polar

surface, respectively. With explicitly assumed $A = 3.62 \times 10^{14}\ \mathrm{m}^2$, $\rho = 1030\ \mathrm{kg\,m^{-3}}$ and conservative-temperature heat capacity $3991.87\ \mathrm{J\,kg^{-1}\,K^{-1}}$, our empirical 2000 m heat-to-sea-level fit corresponds to about $0.170\ \mathrm{kg\,m^{-3}\,K^{-1}}$. That dimensional comparison does not establish a full-depth global response: the effective coefficient depends on where heat enters the ocean. The table's spatial variability is not treated as a Gaussian uncertainty on a global mean, and no new prior or sea-level correction is inferred from it.

The documented Antarctic ocean-temperature transformation is

$$T_o = T_f + \frac{x}{1 + e^{-x/a}}, \qquad x = aT + b - T_f, \qquad T_f = -1.8\,^{\circ}\mathrm{C}. \tag{4}$$

An additional threshold-triggered loss rate can be evaluated when its threshold and rate are supplied. The continuation implements DAIS accumulation, marine discharge and differentiated geometry from Shaffer[9], with sign and unit corrections derived explicitly. Present-day equilibrium and warming tests verify these mechanisms. The target paleo calibration, initial state and full coupling remain distinct recovery tasks. Land-water increments remain externally specified.

Local relative sea level is represented by $L_i(t) = \sum_c F_{ic} S_c(t) + v_i(t - t_0)$, where F contains component fingerprints and positive v_i increases relative sea level. This sign convention includes subsidence as positive. The matrix operation is verified on known area-weighted factors; actual fingerprint grids and coast matching are absent. Reference-period routines require complete annual coverage and refuse silent rebasing of an incomplete series.

3.3 Coastal damage and costs

The independently implemented flooding calculation uses a generalized extreme value (GEV) distribution for annual extreme water level and piecewise uniform area density within supplied elevation bands [10]. For water height w and exposed elevation e, property damage is proportional to $\phi(w - e) = (w - e)/(h_0 + w - e)$ with $h_0 = 1\,\mathrm{m}$. The integral over an exposed band $[a, b]$, where $b \le w$, is

$$\int_a^b \phi(w - e)\, de = (b - a) - h_0 \ln\!\left(\frac{h_0 + w - a}{h_0 + w - b}\right). \tag{5}$$

Multiplication by area density and capital density yields event loss. Mortality uses flooded population, a supplied mortality fraction and value of statistical life. Both contributions are multiplied by $1 - r$, where r denotes economic resilience. Integration of this loss against the GEV density produces expected annual loss. The surge integration bound is explicit, and the truncated density is not renormalized. Retreat changes the lowest retained elevation; protection changes the overtopping threshold.

Relocation combines income-related moving cost, mobile-capital transfer cost, and demolition of immobile capital. The documented benchmark mobile share is one quarter; immobile capital is three quarters. Reactive relocation costs five times the proactive counterpart in the target model. Separate kernels calculate inundation with explicit depreciation, wetland loss with a critical rise rate, and quadratic-height seawall construction. Annual maintenance is returned as a separate cash flow. The numerical coefficient in the construction formula must have units consistent with its quadratic height dependence; the quoted historical reference cost is not silently substituted where its unit convention is ambiguous.

For annual payments C_y, the explicit discount operator is

$$\mathrm{PV}_{y_0} = \sum_y \frac{C_y}{(1 + r)^{y - y_0}}. \tag{6}$$

A 2010 base year and 4% rate are used in verification. This base-year convention does not resolve the target caption's wording about discounting through/to 2150. The optimizer chooses the least-cost member of a supplied finite candidate menu. It does not construct the original adaptation trajectories, because their physical and socioeconomic inputs are missing.

4 Verification results

The expanded verification suite passes 486 Julia assertions across fourteen test groups and twelve Python checks covering diagnostic statistics, coastal probability sampling and expert-density inversion. These include the new carbonate-equilibrium comparison and numerical-library interface. Additional trajectory, calibration, conservation and artifact audits are separate from this unit-test count. Passing these checks does not resolve the predictive failures reported below.

All 81 original-kernel Julia assertions pass across seven groups: emissions/carbon/forcing; ocean conservation and convergence; sea-level limits; spatial accounting and references; extreme levels and damage integration; costs and optimization; and ensemble/regional arithmetic (the sea-level group includes ODE accuracy). These original tests are numerical verification tests. Additional mechanism tests and observational checks are described in Section7. The tests require no third-party Julia packages. Configuration values are explicitly marked as manufactured in `config/validation.toml`. Table1 is generated directly from the diagnostic files and the independent SciPy checks.

Verification quantity	Calculated value
Greenland finest absolute error (m SLE)	5.398e-10
Ocean finest maximum temperature error (K)	6.759e-06
Ocean maximum heat-balance residual (K m)	1.021e-16
Expected flood loss, SciPy reference (USD/year)	442422.479342
Expected flood loss, Julia relative difference	1.271e-11
Depth integral maximum difference (USD)	9.313e-10
Annual emissions residual at gamma=0.07 (% of peak)	1.750000
Minimum sampled Antarctic ocean transform (degree C)	-2.078464

Table 1: Computed verification metrics. These quantities describe manufactured test cases, not target global projections. The heat-balance residual is reported before multiplication by volumetric heat capacity.

4.1 Printed emissions recurrence and its limit

Figure2 compares the literal recurrence with the continuous logistic decline $q(t)/q_\tau = 2/[1 + e^{\gamma_d t}]$. The decrement in Equation1 is the right-endpoint sample of the continuous decline rate. Summing endpoint samples does not exactly integrate that rate. For small step Δt, the residual fraction is $\gamma_d \Delta t/4$ to leading order. At an annual step and decline rate $0.07\,\mathrm{yr}^{-1}$, the calculated residual is approximately 1.75% of peak emissions. The tested rates 0.01 and $0.2\,\mathrm{yr}^{-1}$ give approximately 0.25% and 5%, respectively. Refining the quadrature step reduces the residual. These are normalized mathematical cases extending 2000 years after an arbitrary peak, not calendar-year emissions projections.

The discrepancy separates two interpretations of the published specification: literal annual accumulation and a continuous logistic trajectory that vanishes asymptotically. The subsequent emissions-driven alternatives retain the literal annual recurrence. This manufactured test alone does

not quantify its effect on calibrated temperature or coastal costs; the continuous interpretation has not been propagated through a matched coastal experiment.

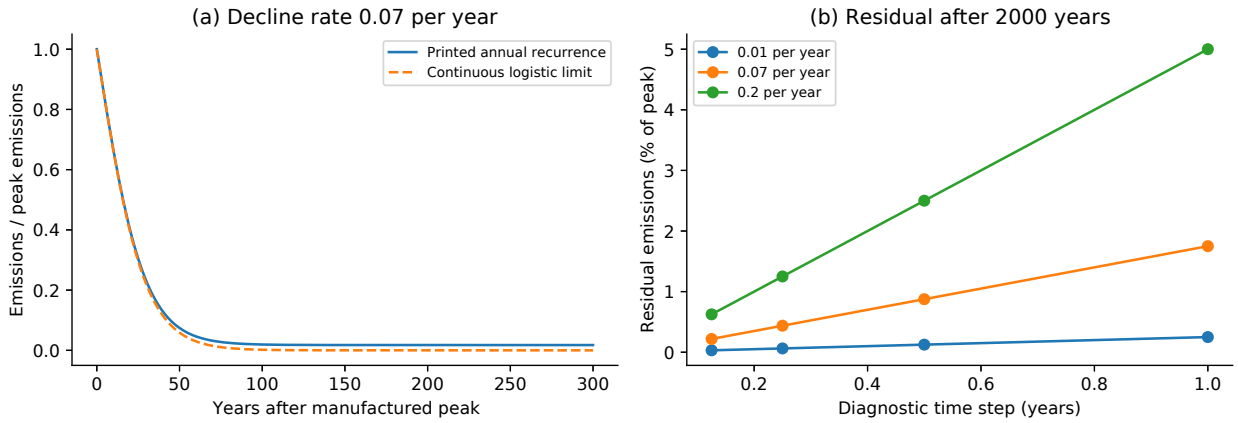


Figure 2: Independent diagnostic of the printed emissions equation. (a) Annual endpoint accumulation for normalized peak emissions and decline $0.07\,\mathrm{yr}^{-1}$, compared with the continuous logistic expression. (b) Remaining emissions after 2000 years for three decline rates and four diagnostic step sizes. The residual decreases with step size; this does not establish which interpretation the original software used.

4.2 ODE and ocean-diffusion convergence

Figure3 tests the Greenland relaxation against its exact constant-temperature solution. The manufactured case has $T = 2\,\mathrm{K}$, $c = -0.2\,\mathrm{m\,K}^{-1}$, $b = 1\,\mathrm{m}$, $\alpha = 0.005\,\mathrm{K}^{-1}\,\mathrm{yr}^{-1}$, $\beta = 0.01\,\mathrm{yr}^{-1}$ and initial volume 1.2 m. The equilibrium is 0.6 m and the integration lasts 100 years. Fourth-order convergence is obtained when the RK4 step is halved.

The ocean case uses a 1 m column, diffusivity $0.1\,\mathrm{m}^2\,\mathrm{s}^{-1}$ and initial temperature $\sin[\pi z/(2L)]\,\mathrm{K}$. The upper temperature is zero and the bottom flux vanishes. The exact amplitude decays as $\exp[-\kappa(\pi/2L)^2 t]$. Time steps decrease with squared grid spacing, so the first-order temporal and second-order spatial errors converge together at second order in spacing. These deliberately artificial dimensions isolate the numerical operator from Earth-system calibration.

4.3 Flood-loss integration and decision sensitivity

Figure4 uses a manufactured $5\,\mathrm{km}^2$ segment distributed uniformly over 0–5 m elevation. Capital density is one million USD km^{-2}, population density $100\,\mathrm{km}^{-2}$, mortality fraction 0.01, value of life 100,000 USD and resilience 0.2. These values are verification fixtures, not recommended valuations. The GEV parameters are location 1 m, scale 0.3 m and shape 0.1. The outer integral extends from -2 to 20 m. SciPy uses its opposite-sign shape convention and independent adaptive integration for both elevation and surge. This verifies the Julia closed-form depth integral and composite-Simpson expectation.

Table2 quantifies two accounting ambiguities with manufactured costs. A two-period menu has candidate costs $(1, 9)$ and $(6, 2)$ at years 2010 and 2060. Keeping one candidate differs from allowing a switch. This menu has no inherited infrastructure state; it is a test of the decision rule only. Likewise, ten annual unit payments during 2010–2019 differ from paying ten units at either

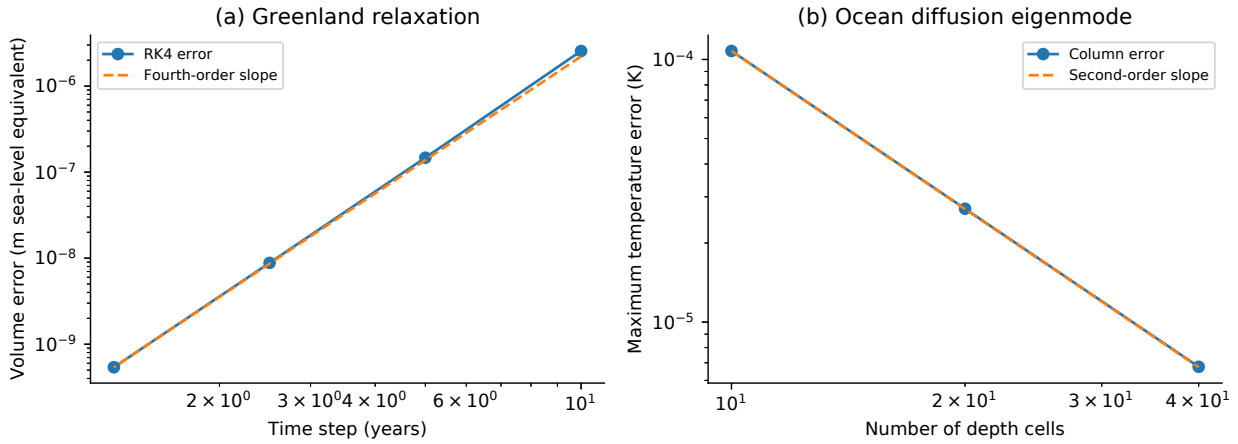


Figure 3: Independent convergence tests. (a) Greenland volume error at 100 years against the exact exponential relaxation; the dashed line has fourth-order slope. (b) Maximum ocean temperature error at 0.2 s for 10, 20 and 40 cells; the dashed line has second-order slope. Time step is 0.2 divided by the square of the cell count, in seconds.

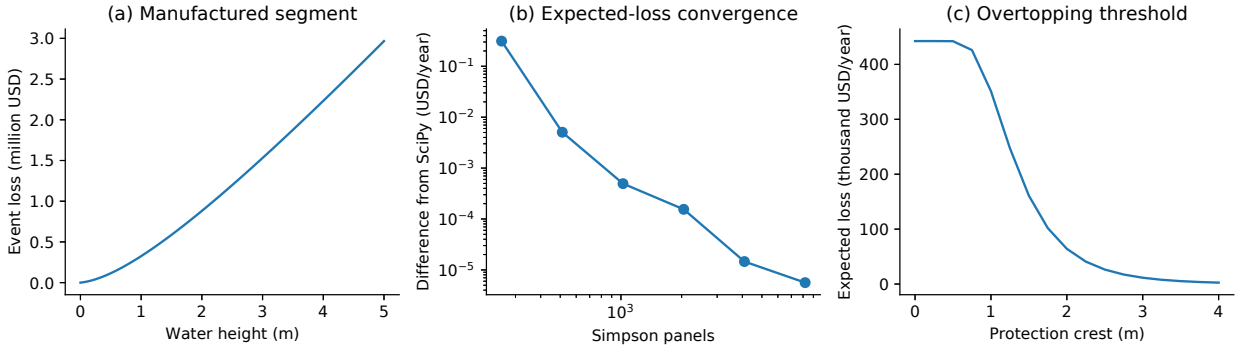


Figure 4: Manufactured coastal damage verification. (a) Event damage increases with water height. (b) Expected annual loss approaches an independent nested SciPy integral as the number of Simpson panels increases. (c) Expected overtopping loss decreases with protection crest. Panel(c) omits construction cost and therefore is not an optimal seawall-height calculation.

end. A coarse timestep must retain the annual-flow interpretation if it is to reproduce the same present value.

Manufactured accounting choice	Present value
fixed strategy pv	2.266414
revisable strategy pv	1.281425
ten annual unit payments pv	8.435332
ten units paid at start pv	10.000000
ten units paid at end pv	7.025867

Table 2: Sensitivity to strategy persistence and payment timing at 4% annual discounting. Units are arbitrary currency; these are not coastal estimates.

4.4 Antarctic transform and its asserted floor

For $a > 0$ and $x < 0$, the added term in Equation4 is negative. The literal transformation therefore falls below T_f, despite its description as bounded below by freezing [7]. For the test values $a = 1, b = 0$, the minimum is approximately $-2.07846\,^{\circ}\mathrm{C}$. In dimensionless form, $z = x/a$, the extremum satisfies $1 + e^{-z} + ze^{-z} = 0$, giving $z \simeq -1.27846$ and a minimum increment $-0.27846a$. Figure5 shows this algebraic behavior. No freezing clamp was added. Whether the calibrated historical trajectory visits this regime remains unknown.

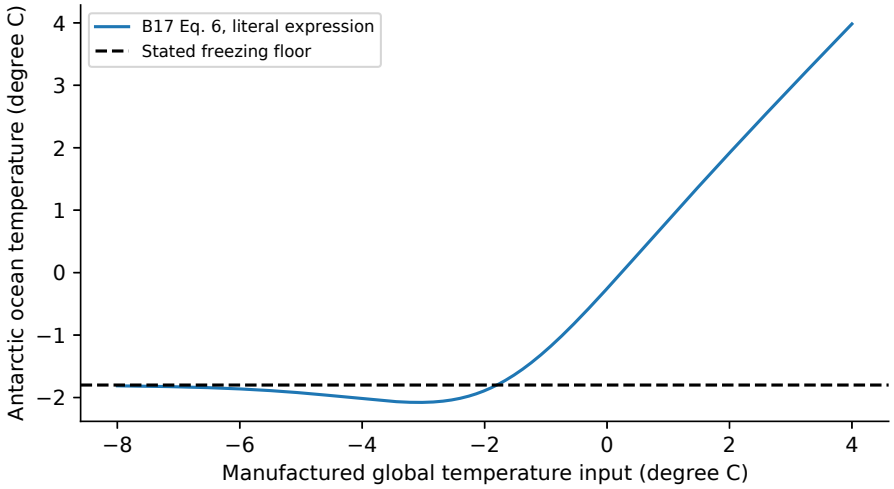


Figure 5: Literal BRICK Antarctic ocean-temperature transformation with manufactured coefficients $a = 1, b = 0$ and stated freezing temperature $-1.8\,^{\circ}\mathrm{C}$. The curve falls below the dashed floor for negative x. This is a specification diagnostic, not an Antarctic projection.

5 Quantitative comparison boundary

The target reports median costs of 1.2, 1.5, 2.1 and 2.6 trillion 2010 USD across the four warming bins, with fifth percentiles 1.0, 1.2, 1.6 and 2.2 trillion USD [3]. The prose supplies 95th percentiles 1.3 and 2.3 trillion USD for the first two bins; unreported numerical endpoints were not digitized or

invented. No independent global costs exist here, so neither relative errors nor agreement intervals can be calculated for those targets.

An arithmetic audit of the printed Table 1 is possible. Its four counts sum to 5,000. Among 250 simulations above 2.7 trillion USD, the text gives 125 in the highest bin and 109 in the third, with none in the lowest. The remaining 16 belong to the second bin. Dividing these counts by bin size yields the rounded probabilities in the paper (Table3). This verifies internal consistency of printed summaries; it does not reproduce the modeled probability distribution or establish an independently calculated tail threshold.

Bin	Printed count	Tail count	Calculated%	Printed%
1	202	0	0.0000	0
2	2264	16	0.7067	1
3	2145	109	5.0816	5
4	389	125	32.1337	32

Table 3: Arithmetic audit using published counts only. The second-bin tail count 16 is derived by subtraction. These data never enter the model kernels.

6 Complete figure-by-figure record

Table4 covers all main and supplementary figures and the main table. Unlettered joint and marginal views are listed separately. Every entry's dependencies, source provenance, assumptions, commands, comparison status and failure explanation are also recorded in `docs/FIGURE_STATUS.csv`. The inventory gives separate commands and saved outputs for the implemented alternatives. Their execution does not resolve the target experiment's remaining input and calibration differences. Supplementary Figure 10 is one histogram containing all-member and Antarctic-trigger-conditioned distributions of proactive migrants.

Table 4: Target reproduction status. "Not reproduced" means no independently generated target numerical result.

Figure/view	Quantity	Status
1: global interval plot	Conditional median and 5th-95th cost quantiles	Not reproduced
2: joint scatter	Sea level 2100 versus discounted costs	Not reproduced
2: top marginal	Sea level distribution by warming bin	Not reproduced
2: right marginal	Cost distribution by warming bin	Not reproduced
3: regional bars	Regional cost/GDP minus global cost/GDP	Not reproduced
4: workflow	Model components and exchanges	Conceptual schematic
S1: interval plot	5000 versus 4000 subsample comparison	Not reproduced
S2: interval plot	202 per bin versus full sample	Not reproduced
S3: trajectory plot	GMST median and 90% intervals plus AR6 boxes	Not reproduced
S4: trajectory plot	GMSL median and 66% intervals plus AR6 boxes	Not reproduced
S5: map	Median local sea level 2100	Not reproduced
S6: interval plot	GTSR versus DINAS cost intervals	Not reproduced
S7: joint scatter	Sea level versus costs above 2.7 trillion	Not reproduced
S7: top marginal	Upper-tail sea level distribution	Not reproduced
S7: right marginal	Upper-tail cost distribution	Not reproduced
S8: regional bars	SSP2 regional cost/GDP deviations	Not reproduced
S9: regression plot	Regional SSP income contrast versus cost sensitivity	Not reproduced

Figure/view	Quantity	Status
S10: histograms	Migration conditional on AIS trigger	Not reproduced
Table1: table	Warming-bin counts and tail probabilities	Not reproduced

Main Figure 1 requires full conditional cost distributions. Figure 2 requires joint sea-level/cost realizations, and Figure 3 additionally requires regional economic aggregation. Supplementary Figures 1 and 2 require controlled subsampling of a legitimate ensemble, not merely a sample-size calculation. Figures 3 and 4 require calibrated time series plus AR6 comparisons. Figure 5 requires spatial fingerprints and site matching. Figure 6 requires matched GTSR and DINAS extreme levels. Figure 7 requires the original type of upper-tail joint distribution. Figures 8 and 9 require paired SSP2/SSP4 economics. Figure 10 requires retained segment migration histories and Antarctic-trigger states. Later sections generate these quantities for explicitly alternative experiments while retaining the distinction between model execution and independent target reproduction.

7 Continued recovery, implementation and observational experiments

The earlier assessment identified missing mechanisms too broadly. Reading the complete Tanaka thesis and the 2022 scholarly preprint recovered the DOECLIM constants, the coastal decision menu and the socioeconomic priors. Independent providers supplied forcing, temperature, carbon emissions, ocean heat and coastal-hazard inputs [20]. The active audit in `docs/BLOCKER_AUDIT.md` distinguishes recovered specifications, derived corrections, explicit alternatives and remaining data gaps.

7.1 Coupled energy balance and numerical heat transport

The land–troposphere and ocean–troposphere boxes exchange heat and lose energy through separate radiative feedbacks [16, 5]. For land fraction $f = 0.29$, their temperatures T_L, T_S satisfy

$$C_L \dot{T}_L = Q - \lambda_L T_L - \frac{k}{f}(T_L - bT_S), \tag{7}$$

$$C_S \dot{T}_S = Q - \lambda_S T_S + \frac{k}{1-f}(T_L - bT_S) - F_O. \tag{8}$$

Here $b = 1.3$, $C_L = 0.52$ and $C_S = 7.8\,\mathrm{W\,yr\,m^{-2}\,K^{-1}}$. The global surface-air observable is $fT_L + (1-f)bT_S$. The equilibrium land/SST ratio is 1.43 and $k = 1.59 - 0.31\lambda_L$. Together with an equilibrium warming of ECS under $3.7\,\mathrm{W\,m^{-2}}$, these relations determine both feedbacks and the exchange coefficient. The explicit formulas are recorded in `docs/RECOVERED_SPECIFICATIONS.md`.

An 80-cell, 4000 m interior ocean exchanges heat continuously with the mixed layer and has an insulated bottom. Finite-volume fluxes cancel under a global heat-capacity weighted sum; this is tested independently of equilibrium fitting. Crank–Nicolson integration uses four steps per year. Refining the vertical grid from 80 to 160 cells changes the year-70 warming under 1% annual CO_2 growth by 0.000787 K; the next doubling changes it by 0.000197 K. This approximately fourfold error reduction agrees with second-order spatial convergence.

The 60 m effective mixed-layer capacity contains about 2.5 m of atmospheric-equivalent capacity. Comparisons with NOAA integrate only 57.5 m of water and the appropriate fractional interior cells down to the stated observational depth. Whole-column heat is not compared to a 700 m record.

HadCRUT4 blends land air temperature with SST, whereas the legacy DOECLIM observable includes enhanced marine-air warming. Treating that blended record as the global-air observable follows the simplified legacy calibration and is an additional representation approximation, not an exact observation operator.

7.2 A closed, explicitly alternative carbon cycle

The recovered terrestrial stocks are 100, 500, 120 and 1500 Gt C in leaves, wood, detritus and soil [17]. Baseline NPP of 60 Gt C yr^{-1} allocates 35 to leaves and 25 to wood. Leaf turnover sends 35 to detritus; wood sends 20 to detritus and 5 to soil. Detritus respires 45 and transfers 10 to soil; soil respires 15. Linear turnover rates are these fluxes divided by their source stocks. Fertilization modifies NPP and the documented Q_{10} factor modifies the two respiration fluxes [6]. All internal transfers cancel in the global carbon balance.

Ocean carbon uses an instantaneously equilibrated atmosphere–mixed-layer reservoir and three deeper reservoirs. Constant alkalinity, total boron and carbonate equilibrium determine atmospheric CO_2 and mixed-layer carbon from their combined inventory [18]. The nonlinear scalar root is solved by bisection in log hydrogen concentration. At 280 ppm the reconstructed DIC is 0.002080909 mol L^{-1}, consistent with the printed H99 value to its reported precision. The mixed-layer depth and deep exchange parameters are from the later Tanaka description. Mixing these documented versions, and holding carbonate chemistry temperature fixed, defines an alternative closure; it does not establish exact recovery of the target SNEASY version.

Second-order splitting couples carbon RK4 steps to climate CN steps. Twenty-six checks establish baseline equilibrium, chemical partition, stock positivity, second-order coupled convergence and total-carbon conservation. Across 18 historical sensitivities, the largest carbon-budget residual is below 10^{-11} Gt C. Conservation alone does not validate the response: the uncalibrated default predicts 461.38 ppm in 2022, compared with NOAA's 417.08 ppm. The GCB historical table has no pre-1850 land-use emissions; the input retains blanks, while separate calculations prescribe zero or the 1850 rate before 1850. These are alternatives, not recovered observations.

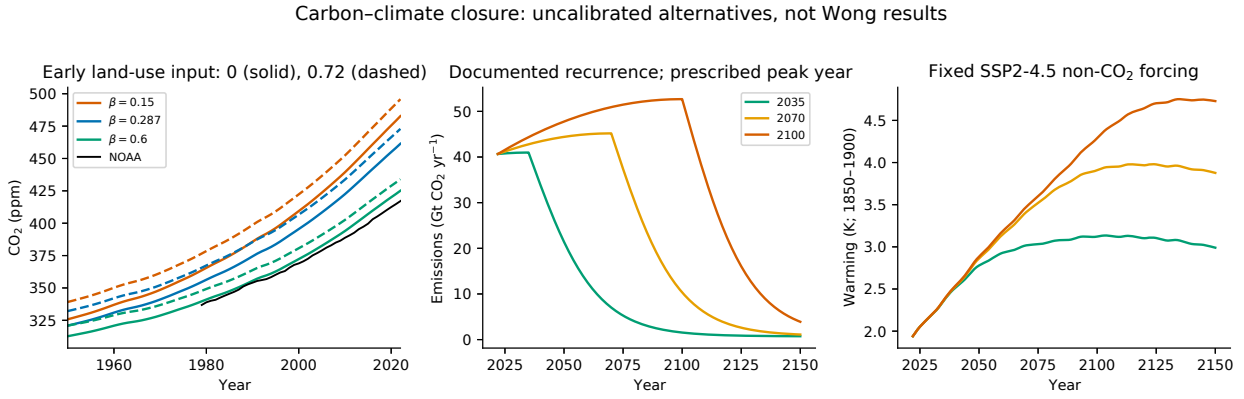


Figure 6: Uncalibrated carbon–climate experiments, not target reproductions. Left: historical CO_2 sensitivities to fertilization and the explicitly assumed early land-use input, with independent NOAA observations. Centre: documented emissions recurrence initialized from GCB2023, with three prescribed peak years. Right: coupled warming under fixed SSP2-4.5 non-CO_2 forcing. The historical mismatch limits interpretation of these future curves.

7.3 Temperature-only calibration and failed holdout checks

The first observational experiment estimates ECS, heat diffusivity, aerosol multiplier, temperature offset, AR1 innovation scale and AR1 correlation. HadCRUT4.6.0.0 years 1850–2005 form the training data; years 2006–2021 are held out. Known observational standard deviations are approximated by the total 95% interval widths divided by 2(1.96). This approximation omits the record's full temporally correlated measurement covariance. A scalar Kalman filter integrates the latent AR1 residual exactly under these assumptions; tests match the dense Gaussian covariance likelihood to 10^{-12}.

Uniform bounds follow the B17 climate table, except that ECS is restricted to 0.7–10 K: the recovered exchange relation becomes negative below about 0.68 K. CEDA effective forcing supplies historical and future inputs, with both aerosol components multiplied by the same parameter. This differs from the older indirect-aerosol convention and from Darnell's uncertain emissions and nearest-SSP forcing construction. Neither the posterior nor the six scenario experiments is claimed as the Wong ensemble.

An initial short run failed convergence checks. The retained calculation uses four dispersed chains, each with 6000 warmup and 30,000 production steps; every fifth state is retained. Proposal adaptation stops before production. Log ECS and log diffusivity proposals include their Jacobian, preserving the physical-coordinate prior. Independently implemented rank/folded split $\widehat{R}$ and bulk/tail ESS follow Vehtari et al. [19]. The largest $\widehat{R}$ is 1.00427 and the smallest ESS is 1894, satisfying the prespecified limits of 1.01 and 400.

Conditional forecasts propagate the final training-period residual state and include both AR1 innovation and observational uncertainty. Despite sampler convergence, holdout temperature RMSE is 0.209 K, with mean bias +0.189 K. Only 10 of 16 held-out annual temperatures lie inside the nominal 90% predictive intervals. After fitting only an additive heat datum offset over 1955–2005, the 700 m heat parameter-median RMSE in 2006–2021 is 16.46×10^{22} J. These failures motivate a separate temperature–heat calibration; they are not removed from the record.

One thousand complete joint parameter rows generate six SSP trajectories; 250- and 500-row subsets test sampling variation. The physical temperature intervals exclude future residual noise. SSP2-4.5 mean warming over 2081–2100 is 2.924 K, with a parameter 90% interval of 1.976–4.188 K. The corresponding printed AR6 assessment is 2.7 [2.1, 3.5] K. This comparison uses the same averaging/reference periods but compares different inference systems; it is not a numerical reproduction of Supplementary Figure S3.

7.4 Joint temperature–heat calibration

The next experiment adds independent NOAA 0–700 m heat observations over 1955–2005, with a datum offset and separate AR1 residual parameters. The first joint run failed convergence ($\widehat{R} = 1.0218$ and a heat-correlation ESS of about 285). Its diagnostics remain in the repository. A revised sampler uses logarithmic innovation scales, logit correlations and a recent warmup covariance window; Jacobians preserve the same physical priors. Four chains of 8000 warmup and 50,000 production iterations retain 20,000 states. All acceptance checks pass: the maximum $\widehat{R}$ is 1.00582 and the minimum bulk/tail ESS is 589.

An exact modal solver reduces computation while retaining all finite-volume degrees of freedom. A diagonal similarity transforms the climate operator into a symmetric tridiagonal matrix. Each mode is integrated analytically under linearly interpolated annual forcing. Eighteen checks agree with an independent augmented matrix exponential and establish second-order Crank–Nicolson convergence. Four substeps per year differ from the exact solution by at most 0.00133 K in three

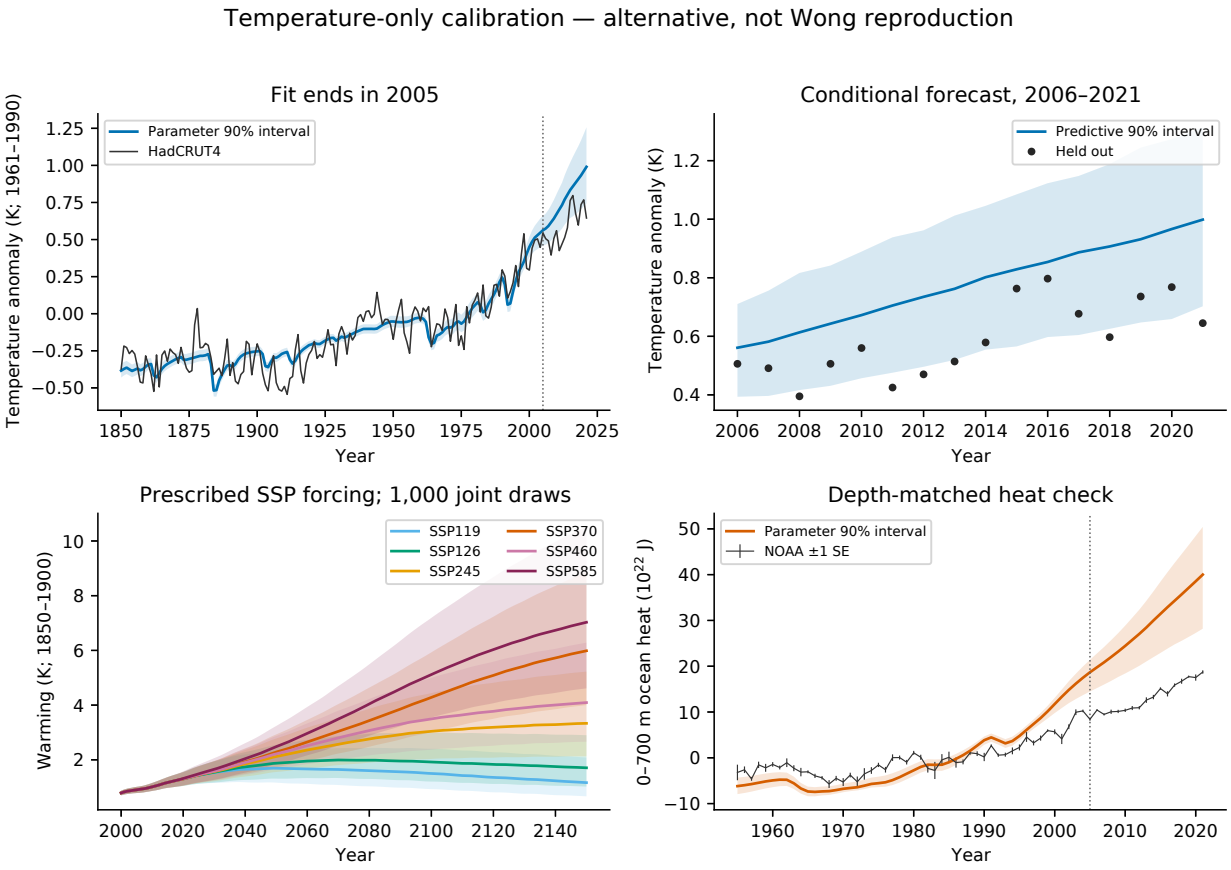


Figure 7: Temperature-only calibration and independent observational checks. The upper-right forecast includes residual and measurement uncertainty. The lower-left scenario bands contain parameter uncertainty only. Ocean heat is compared at the actual observational depth, with a training-period datum offset. The temperature and heat holdout discrepancies prevent treating this converged alternative posterior as validation of the target model.

historical comparisons; the exact solver is used consistently for the new joint fit and projections.

The joint fit reduces held-out temperature RMSE to 0.0870 K and bias to +0.0402 K. Heat predictive RMSE is 0.879×10^{22} J. All 16 held-out years lie in both nominal 90% predictive intervals. These short, correlated holdouts provide limited validation; coverage is not a proof of correct uncertainty. ECS has a median of 1.485 K and a parameter 90% interval of 1.093–1.958 K. Future warming is substantially below the AR6 assessment: SSP2-4.5 warming over 2081–2100 is 1.973 [1.518,2.483] K, compared with 2.7 [2.1,3.5] K; SSP5-8.5 is 3.319 [2.569,4.160] K, compared with 4.4 [3.3,5.7] K. This discrepancy remains unresolved. Improved historical prediction does not justify replacing the target's joint climate–carbon–ice calibration with this alternative.

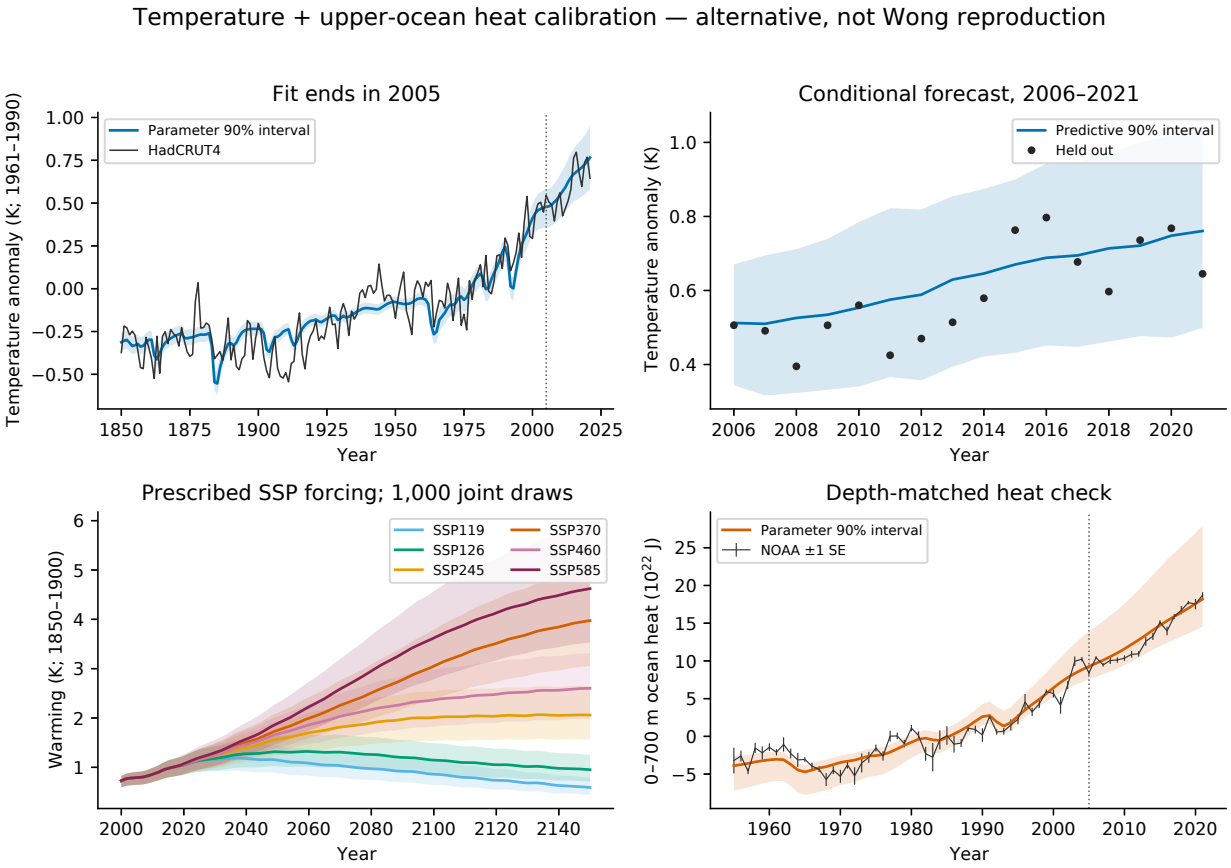


Figure 8: Alternative joint temperature and upper-ocean-heat calibration. Observations after 2005 are withheld. Future bands propagate parameter uncertainty under prescribed SSP forcing. The improved holdout performance and the low late-century warming relative to AR6 are both retained.

7.5 Antarctic events and observational thermal expansion

DAIS geometry, accumulation and marine discharge are now implemented. The discharge water depth follows sea level minus bed elevation, correcting the printed sign from geometry; coefficient units follow dimensional consistency. The equilibrium volume of $2.478479 \times 10^{16}\,\mathrm{m}^3$ differs from Shaffer's rounded $2.478 \times 10^{16}\,\mathrm{m}^3$ by 0.0193%. This is a reproduced methodological benchmark, not a reproduced Wong figure.

The printed runoff gate activates at $h_0 + cT_A = 0$. Over marine bed this can produce a

discontinuous accumulation rate. Splitting time steps at that event and evaluating endpoint stages on the appropriate side removes the earlier irregular refinement. In the 6 K, century-long warming ramp, the event occurs at year 41.92982456. Two-year and one-year steps differ in year-200 radius by 1.57×10^{-7} m; finer steps reach floating-point limits. The dashed curves in Figure9 instead activate runoff when the runoff line exceeds local edge height, and remain explicit alternative physics. These tests express sea level by mass conservation using ice/water densities and ocean area, whereas B17 scales its initial Antarctic volume to 57 m. The two observable conventions are not silently interchanged.

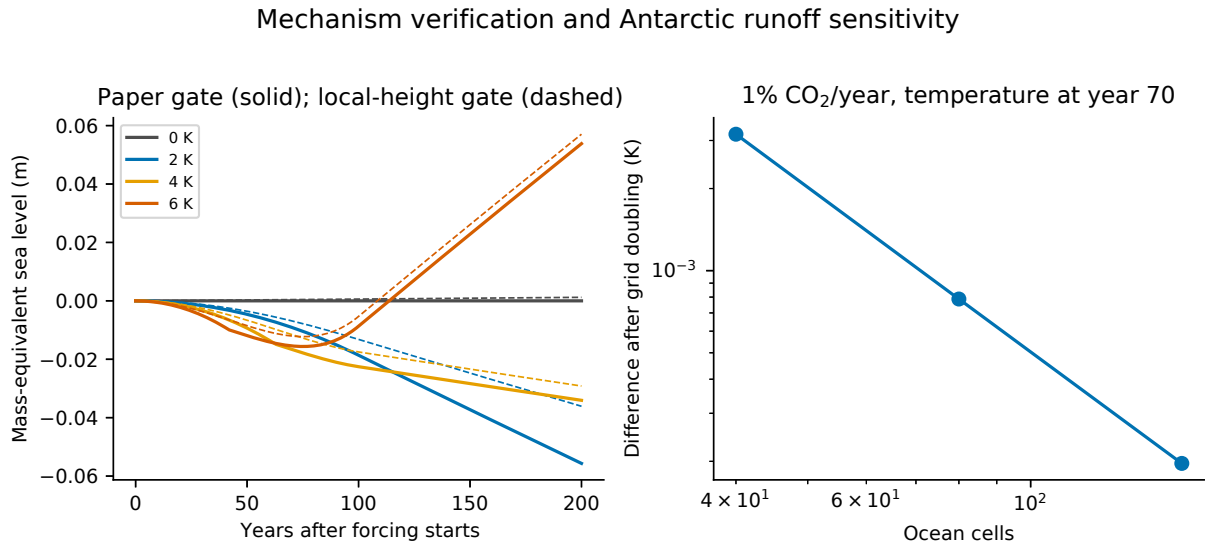


Figure 9: Independent mechanism tests. Left: prescribed Antarctic warming, with fixed ocean temperature and sea level, isolates the runoff convention. Right: climate grid refinement under 1% annual CO_2 growth. Neither panel is a global calibrated projection.

An independent coefficient check uses NOAA heat and thermosteric sea level from the same depth range. OLS over 1955–2005 at 0–700 m gives 1.392×10^{-25} m J^{-1}; fitting 2005–2015 at 0–2000 m gives 1.107×10^{-25} m J^{-1}. Held-out thermosteric RMSEs through 2021 are 0.414 and 0.757 mm, respectively. First-difference slopes provide a check against merely fitting a common trend. Because heat and sea level share hydrographic observations, these are descriptive constraints rather than independent-error coefficient posteriors. They are not substituted silently for the unavailable V19 prior or a full-depth expansion coefficient.

7.6 Coastal cash flows and input quality

The annual coastal experiment retains the documented ten-choice menu and first-period decision rule. Sixty checks establish nonnegative cash-flow components, first-period policy selection, no repeated relocation at fixed sea level and telescoping incremental dike construction costs. Upgrades at 2010, 2051 and 2101, annual payment timing and a paired no-SLR counterfactual are explicit alternatives to the target's decadal accounting and initial-step subtraction. Forty-eight synthetic scenarios compare income/exposure, sea-level rise, retreat land accounting and linear versus quadratic maintenance. All exposure values are labeled fixtures; no global cost estimate is inferred. Doubling flood quadrature from 512 to 1024 panels changes the verification-case NPV by about 0.00024 USD and leaves its policy unchanged.

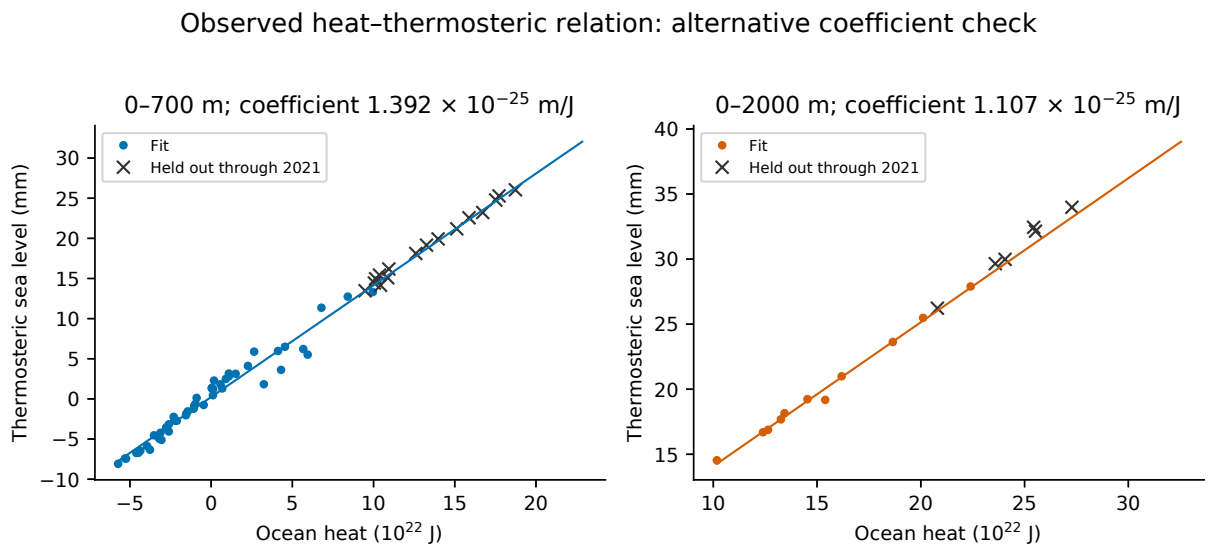


Figure 10: Independent NOAA heat and thermosteric observations with fitted depth-specific linear relations. Later years through 2021 were withheld. This coefficient check is an explicit alternative to the target calibration.

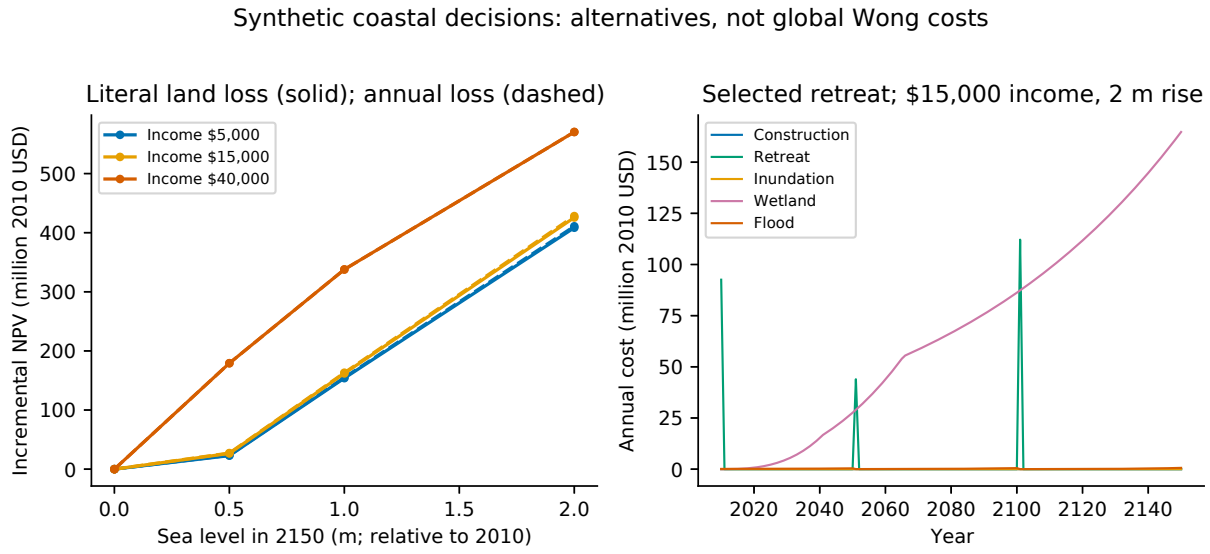


Figure 11: Synthetic coastal alternatives. Left: paired-counterfactual incremental costs under two explicit land-accounting rules. Income and exposure both vary across the three fixtures, so this is not an isolated income effect. Right: annual components for one selected retreat policy, including discrete upgrades. These outputs do not estimate a named coast or a global cost.

The independent GTSR shapefile has 12,148 coastal feature rows but only 7868 distinct station labels. The initial interpretation of these labels as unique DIVA identifiers was corrected by this audit. A second public release provides 16,611 station-wise Gumbel fits. Geographic mapping leaves 133 nearest matches with invalid parameters; these remain missing, not filled from distant stations. The two input products also disagree in their 100-year levels. These findings prevent a silent merge or assumption of exact input equivalence. The full target elevation-area, exposure and DINAS tide records still require recovery. The next section uses a separately sourced geographic input product for an explicit alternative, with different segmentation and monetary conventions.

Every alternative here has its own input, script, output and figure record. No alternative curve is entered into the target inventory as a recovered published numerical result. Ongoing calibration and data-recovery tasks remain active rather than being reclassified as irreducible merely because the first implementation lacked them.

8 Conditional ice and geographic coastal alternatives

8.1 Carbon constraints beyond atmospheric concentration

Atmospheric concentration alone does not identify terrestrial fertilization and ocean transport separately. We prescribe the observed temperature history and fit fertilization against NOAA CO_2 in 1979–2005, retaining 2006–2021 as a holdout. The nine conditional fits vary respiration $Q_{10} = 1, 2, 3$ and mixed-layer exchange $\eta = 10, 16.87, 30\,\mathrm{m\,yr^{-1}}$. A bounded scalar search over $\beta \in [0, 2]$ is a new numerical experiment, not the wide R08 Bayesian prior. Every fit begins in 1850 at the alternative closure's 280 ppm equilibrium. This avoids the missing pre-1850 land-use series but omits earlier perturbations.

Figure 12 shows that fits with similar atmospheric residuals predict different ocean inventories. At $Q_{10} = 2$ and $\eta = 16.87\,\mathrm{m\,yr^{-1}}$, the fitted $\beta = 0.6680$ yields a 0.535 ppm holdout RMSE and 118.06 Gt C ocean uptake by 1994. R08 cites the independent estimate 118 ± 19 Gt C; its nominal 1800–1994 interval differs from our 1850 initialization, as in R08's equilibrium approximation. Predicted 1980s and 1990s ocean fluxes are 1.98 and $2.32\,\mathrm{Gt\,C\,yr^{-1}}$, compared with the written constraints 1.6 ± 0.4 and 2.0 ± 0.4. These checks do not condition the fit. Across all nine cases, time-step doubling changes CO_2 by at most 0.000333 ppm and carbon accounting closes within 7×10^{-12} Gt C. The remaining chemistry-version and prior uncertainties prevent interpreting these conditional fits as the target carbon posterior.

8.2 Independent ice observations and identifiable parameters

Independent IMBIE and GlaMBIE observations allow a conditional test of the documented Greenland and glacier equations[21, 22]. IMBIE monthly mass-change rates are averaged annually. Their errors are temporally dependent; we use the arithmetic mean of monthly standard deviations as an upper bound on annual-mean uncertainty, rather than dividing by $\sqrt{12}$. GlaMBIE supplies annual rates and standard errors directly. All rates remain in $\mathrm{Gt\,yr^{-1}}$ until the explicit conversion 362500 Gt per metre of sea-level equivalent. Observed HadCRUT4 temperature is prescribed, so held-out ice years are a conditional hindcast, not a forecast made without knowledge of temperature.

The glacier equation is separable. For $n \neq 1$, define $u = (1 - S/V)^{1-n}$. Then

$$u(t + \Delta t) = u(t) - \frac{(1 - n)\beta}{V} \int_t^{t+\Delta t} [T(s) + 0.15]\,\mathrm{d}s. \tag{9}$$

The integral is exact for linearly interpolated annual temperature. For $n = 1$ the remaining volume decays exponentially with the same temperature integral. Zero remaining volume is absorbing,

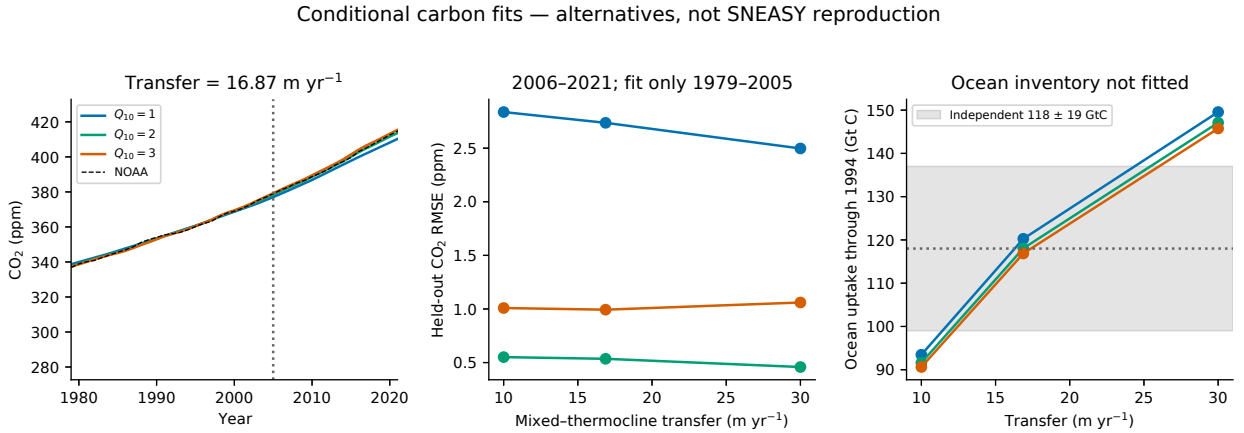


Figure 12: Conditional carbon alternatives. Left: fitted concentration histories for three respiration choices at a fixed transfer rate; the vertical line ends the fit. Centre: held-out concentration errors. Right: independently predicted ocean inventory, with the observational interval cited by R08. The latter was not used in optimization. These fits are not target-model reproductions.

including extinction before the end of an interval whose temperature later falls. This explicit convention resolves the nonunique mathematical continuation for $n < 1$.

For Greenland, write sea-level contribution $S = V_0 - V$ and use the documented linear equilibrium volume. The equation becomes

$$\dot{S} = [d + kT - S]10^{-3}(aT + b), \qquad d = V_0 - b_{\text{eq}}, \quad k = -c. \tag{10}$$

Annual rates cannot distinguish V_0 and b_{eq} separately. Their B17 independent uniform priors induce a trapezoid density for d, which we derive from the overlap of the two intervals. We retain that density instead of inventing a uniform prior on the difference. Negative inverse response times are excluded on the modeled temperature domain. Four RK4 steps per year agree with refined and analytical cases.

We fit rates through 2014 and reserve 2015–2020. Glacier initialization in 1990 and Greenland initialization in 1961 are explicit alternatives. Physical bounds follow B17; residual innovation standard deviation is newly assigned $[0.1, 250]\,\text{Gt yr}^{-1}$ and AR1 correlation $[0, 0.999]$. These rate-error priors differ from the original cumulative-state calibration. The first chains failed the tail effective-sample criterion and were retained as failed attempts. Only chains passing the recorded diagnostics enter subsequent projections.

The longer runs pass: the maximum $\widehat{R}$ is 1.0012 and the minimum tail effective sample size exceeds 1100. Conditional holdout RMSE is $84.8\,\text{Gt yr}^{-1}$ for glaciers and $128.9\,\text{Gt yr}^{-1}$ for Greenland. Both 90% predictive intervals cover all six held-out years, but that short record does not establish predictive accuracy. A straight-line glacier trend fitted only to training data has lower holdout RMSE, $73.6\,\text{Gt yr}^{-1}$. The glacier model therefore does not improve on this simple baseline despite its physical structure and converged sampler. Greenland improves on the tested training-mean, last-year and linear-trend baselines (144–$161\,\text{Gt yr}^{-1}$).

Future component calculations pair complete ice-parameter rows with complete joint climate rows independently. They do not recover the target's cross-model posterior correlations. Integration restarts from the observed-temperature hindcast state in 2021; its boundary temperature is observed, followed by modeled temperatures on the same reference datum. Future ribbons include parameter

uncertainty only. The 0–2000 m thermal component uses the independently fitted NOAA coefficient, excluding unobserved deeper expansion.

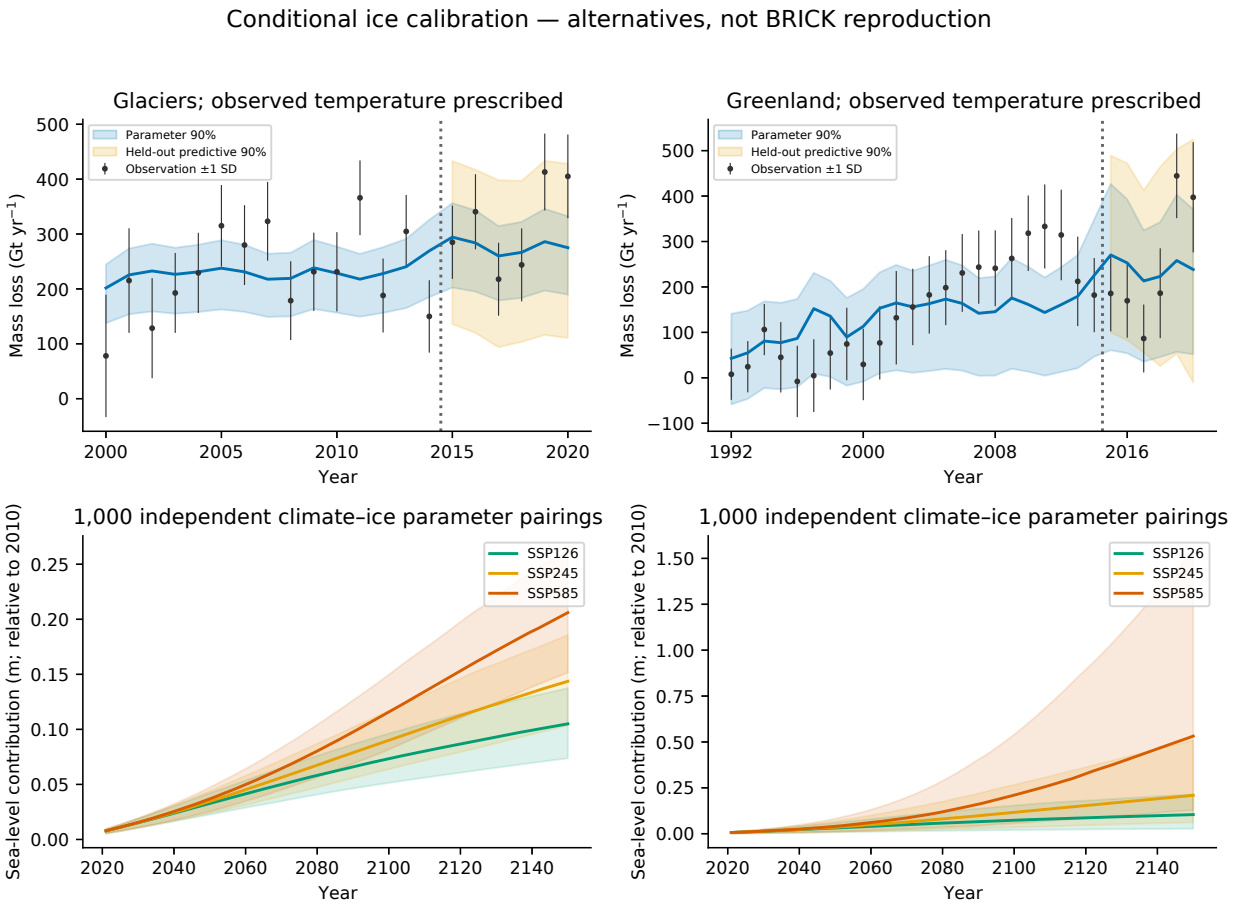


Figure 13: Conditional ice-component alternatives. Upper panels compare annual mass-loss rates with independent observations; the fit ends in 2014. Blue ribbons show parameter uncertainty, orange ribbons held-out predictive uncertainty including AR1 residuals and observational error. Lower panels show parameter-only projections relative to 2010 for three SSP forcing scenarios. These are separate component predictions, not a reproduced global sea-level distribution.

IMBIE gravimetry can contain peripheral glaciers also represented by GlaMBIE. Combining these series without an overlap correction could double count ice loss. The original component experiment therefore reports them separately. Otosaka and coauthors[21] quantify a4.1%/3.3% change in reconciled Greenland/Antarctic estimates when peripheral glaciers are removed from gravimetry over2002–2020, below reported uncertainty. This is a specific accounting limitation, rather than a prohibition on any alternative total. An explicitly uncorrected component budget and covariance sensitivities can therefore be investigated, while exact target component boundaries and joint calibration remain unresolved.

8.3 Geographic exposure and regional accounting

SLIIDERS v1.1 supplies independently published elevation-binned land, population, capital, wetland, surge and socioeconomic inputs[23]. Only its input product was retrieved; no pyCIAM implementa-

tion, construction scripts or model outputs were consulted. Aggregating 11980 subadministrative records by coast and country yields 9574 groups. Six have no coastline and 37 lack positive economic inputs; their inventories and exclusion reasons remain in a separate CSV. The experiment uses the remaining 9531 coast-country groups.

The alternative uses SSP2 and SSP4 OECD economics in 2019 USD PPP, annual interpolation between decadal nodes, and a 2010–2100 horizon. Sea level is spatially uniform and prescribed as $h[(t-2010)/90]^{1.5}$, with $h = 0.5, 1, 2\,\mathrm{m}$. It has no temperature-conditioned probability interpretation. The model selects the least-cost first-period policy, retains its strategy and return level, and updates height or retreat thresholds in 2051. Costs are discounted to 2010 at 4%. Incremental cost subtracts a paired zero-SLR calculation. The lowest retreat level follows mean sea level, an explicit alternative to unavailable DINAS annual-high-tide inputs. Wetland exposure follows its actual elevation profile, with the documented rate-dependent loss fraction. These conventions do not recover the target's segmentation, monetary basis, decadal accounting, 2150 horizon or spatial fingerprints.

The initial pilot exposed a 9.0% quadrature discrepancy in no-adaptation costs: narrow surge densities could lie between nodes even when elevation breaks were resolved. Splitting integration at both elevation breaks and hazard quantiles reduces the maximum 12-versus-24-node difference to 1.5×10^{-11} across 114 pilot cases, with no policy changes. Original 0.1 m profiles additionally test the separate six-bin exposure approximation. Numerical integration agreement does not validate exposure coarsening. Indeed, original-resolution profiles change two of the 114 pilot policy choices and increase summed incremental cost by 9.6%. The largest individual relative change is 71.6%. This deterministic pilot is not representative sampling and cannot supply a global correction factor. The coarse global experiment remains a sensitivity result while a full original-resolution SSP4, 1 m calculation tests the global effect.

Regional ratios use independent World Bank national GDP[24]. We value 2010 GDP at 2019 PPP by multiplying 2019 current-PPP GDP by the country's 2010/2019 real-GDP ratio. Available landlocked-country GDP is included in denominators; coastal costs lacking a matching positive GDP are excluded from the ratios and their coverage is recorded. Afghanistan and Pakistan are restored to their pre-July-2025 South Asia classification, as documented by the World Bank. These explicit choices support an inequality sensitivity experiment, not reproduction of the target's regional numerical bars.

SSP	Rise (m)	Selected cost	No-adapt. cost	Migrants (million)
SSP2	0.5	0.435	11.824	99.465
SSP2	1.0	1.019	25.776	127.280
SSP2	2.0	2.370	57.592	197.987
SSP4	0.5	0.421	10.131	98.755
SSP4	1.0	0.990	22.141	125.951
SSP4	2.0	2.305	50.112	195.002

Table 5: Geographic alternative totals. Cost columns are trillion 2019 USD PPP, discounted to2010. Migrants count movements in the selected policy over2010–2100 and are not unique individuals over a demographic lifetime. No warming probabilities are assigned.

8.4 Complete original-resolution coastal comparison

The earlier19-unit pilot cannot establish a global elevation-resolution error. A subsequent paired experiment evaluates every eligible unit with all200 original0.1 m exposure bins, retaining the same inputs, prescribed trajectory, quadrature and policy conventions. Figure15 and the computed table

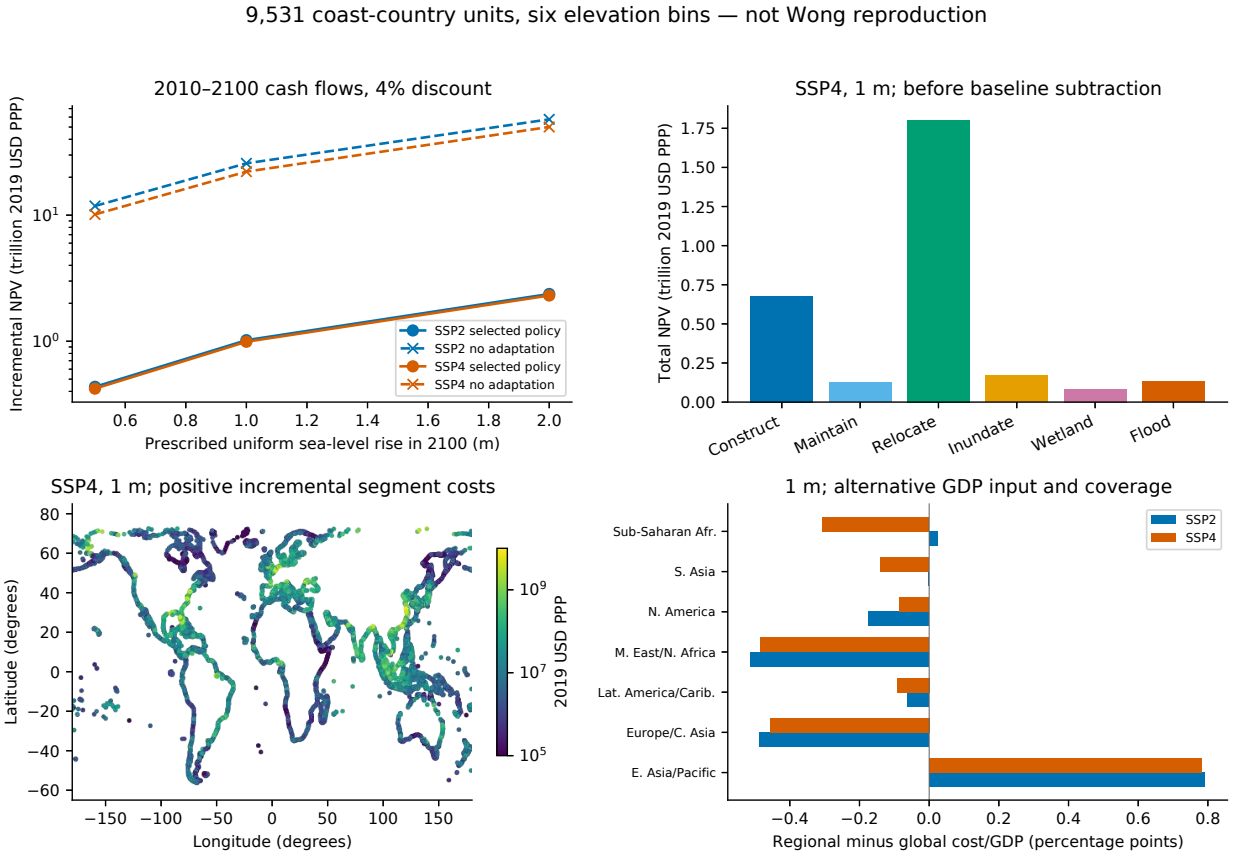


Figure 14: Geographic coastal alternatives. Upper left: global incremental costs for prescribed uniform sea-level paths, with paired no-SLR baselines for both selected policy and no adaptation. Upper right: unadjusted cost components for SSP4 and 1 m. Lower left: positive incremental costs at coast coordinates; points represent input coast-country groups, not population density. Lower right: regional minus global cost/GDP fractions at 1 m using independent GDP inputs and documented coverage exclusions. Every panel is a non-reproduction; the paths have no calibrated warming-bin or tail probability.

Ice component	Max $\widehat{R}$	Min bulk ESS	Min tail ESS	Holdout RMSE (Gt/yr)
Glaciers	1.0009	2752	1116	84.8
Greenland	1.0011	4449	2561	128.9

Table 6: Final conditional ice calibration diagnostics and held-out rate errors. Every coordinate and log density must pass $\widehat{R} < 1.01$ and both ESS$\geq$ 400. Observed temperature is supplied in held-out years.

report the complete comparison. Its cost difference is an effect of redistributing exposure within the original six bins, rather than adding population, capital or land to the domain. Independent binary and CSV readers agree exactly on the original-resolution pilot; doubled flood quadrature changes costs by less than 2×10^{-13} in relative terms and changes no pilot policy. The preferred

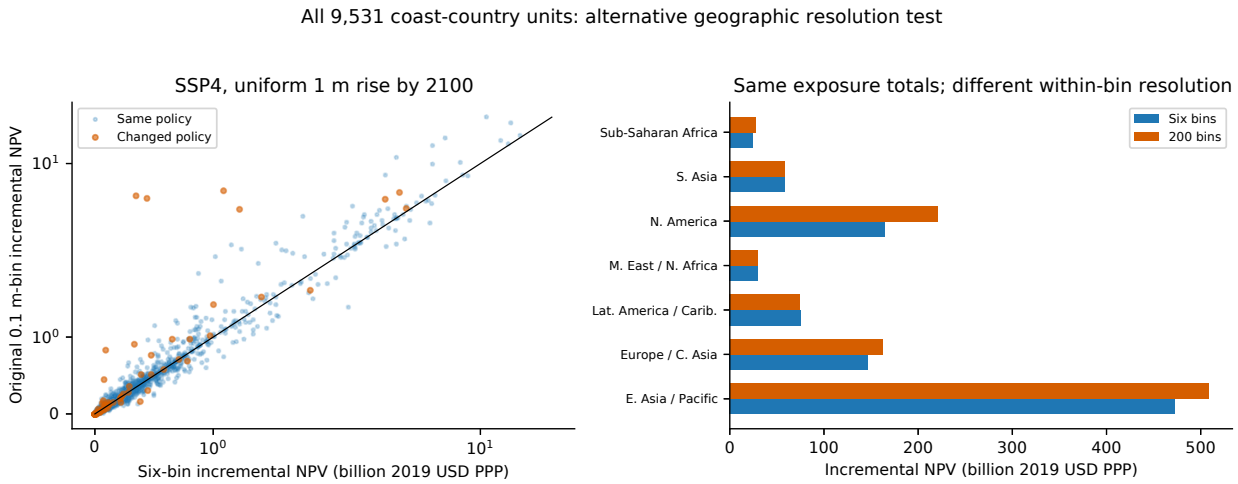


Figure 15: Matched SSP4,1 m coastal calculations for all9,531 eligible units. Left: incremental cost after subtraction of the paired zero-rise baseline; both axes use a symmetric logarithmic scale and the line marks equality. Right: World Bank regional totals under the same two elevation resolutions. Both calculations are geographic alternatives with uniform prescribed sea level, annual2010–2100 cash flows, and2019 PPP values. They are not recovered Wong warming-conditioned costs.

Quantity	Six bins	200 bins	Change (%)
Selected policy total NPV	2.9932	3.0576	+2.15
Paired-baseline incremental NPV	0.9900	1.0992	+11.03
Movement count (million)	125.9506	128.5904	+2.10

Table 7: Complete matched resolution test for 9,531 coast-country units, SSP4 and prescribed1 m rise by2100. Cost rows use trillion 2019 USD PPP. The input elevation range is unchanged; only within-bin exposure resolution changes.

first-period policy changes in 148 units (1.55%). These totals replace extrapolation from the small coastal pilot; they do not establish convergence of the physical exposure data themselves.

9 Recovered observation vintages and calibration tests

9.1 What the written calibration specification now supplies

The later Darnell manuscript's supplementary Table 2[33] identifies HadCRUT5.0.1.0 temperature for 1850–2021, Law Dome CO_2 for 1854–1953, Mauna Loa CO_2 for 1959–2017, and 0–3000 m ocean heat content for 1953–1996. It also identifies the IMBIE, glacier and global sea-level products used in the joint calibration. This recovers several input definitions that the initial assessment treated as missing. The later table is evidence about the methodological lineage; its identity with the target's unpinned 2024 manuscript remains unestablished. The 2024 text explicitly adds

autoregressive residuals to simulated observables, and Urban and Keller provide the stationary AR(1) likelihood[5]. Neither the statistical family nor all observation versions should therefore be described as inaccessible.

Original-provider archives supply the specified HadCRUT5 version and 83 irregular Law Dome readings within the stated period. NOAA supplies a current Mauna Loa annual series. Its 2026 release date is recorded; it is not silently equated with the target's historical release. Gouretski and Koltermann's written 3000 m heat estimate is a 12.8×10^{22} J increase between 1957–1966 and 1987–1996, with uncertainty of at least 8×10^{22} J [37]. We recover this aggregate constraint but not a verified annual 1953–1996 series. We do not digitize their curve or assign their uncertainty statement a Gaussian standard deviation. The existing HadCRUT4/700 m posterior predicts a median increase of 6.17×10^{22} J over those decades, with a parameter 5th–95th range of 4.86–12.37×10^{22} J. Different depths, error definitions and averaging conventions require explicit treatment in any subsequent joint fit.

9.2 Temperature-version sensitivity

HadCRUT5 differs from HadCRUT4 by an RMS 0.0526 K over the training years and is on average 0.1145 K warmer during the held-out 2006–2021 period. This is a consequential input difference rather than a numerical tolerance. Importance reweighting 1,000 retained HadCRUT4 posterior rows under the HadCRUT5 likelihood has effective sample size 14.1; it fails the declared minimum of 400. Those weights are retained as a failed diagnostic and do not define an accepted replacement posterior.

We instead run four fresh chains, retaining the same temperature/700 m heat likelihood, priors and 1850–2005 training boundary. Each chain has 8,000 warmup and 50,000 production iterations, thinned by ten. All coordinates and log density pass the declared diagnostics: maximum rank/folded split $\widehat{R}$ is 1.00954, minimum bulk ESS is 474 and minimum tail ESS is 430. The fitted ECS median rises from 1.485 to 1.660 K; its HadCRUT5 5th–95th range is 1.279–2.132 K. Held-out temperature RMSE is 0.0854 K and heat RMSE remains 0.879×10^{22} J. These checks establish convergence and predictive behavior for the restricted experiment, not recovery of the target joint fit.

Using 1,000 complete parameter rows, the SSP2-4.5 2081–2100 warming median rises from 1.973 to 2.223 K, with a 1.753–2.689 K parameter range in the new fit. The corresponding AR6 assessment is 2.7 (2.1–3.5) K. Under SSP5-8.5, the new estimate is 3.759 (2.971–4.476) K, compared with 4.4 (3.3–5.7) K. The temperature vintage explains part of the discrepancy but does not close it. The coupled coastal base ensemble below retains the earlier HadCRUT4 fit so that its successive numerical and socioeconomic tests remain matched.

The written residual convention also motivates an observable-classification test. Adding a stationary normal residual with variance $\sigma^2/(1-\rho^2)$ to annual 2100 temperature gives an exact Gaussian-mixture probability for each warming category. Under the 50% budget weights, the four probabilities change from 58.815/33.350/7.781/0.0534% to 59.424/32.933/7.571/0.0724%. This is an observation-error alternative: the residual does not drive carbon or ice dynamics, and no sea-level residual is silently added. It therefore tests classification separately from physical climate response.

9.3 Carbon constraints that fail to transfer

The initial atmospheric CO_2 and fertilization coefficient are fitted under the documented bounds 275–295 ppm and 0–1. The chemical reference equilibrium is recomputed at each initial concentration. Temperature sensitivity of respiration and thermocline exchange remain fixed at $Q_{10} = 2$ and 16.87 m yr^{-1}. An initial noise-free recovery case passes, but the subsequent audit below shows that this

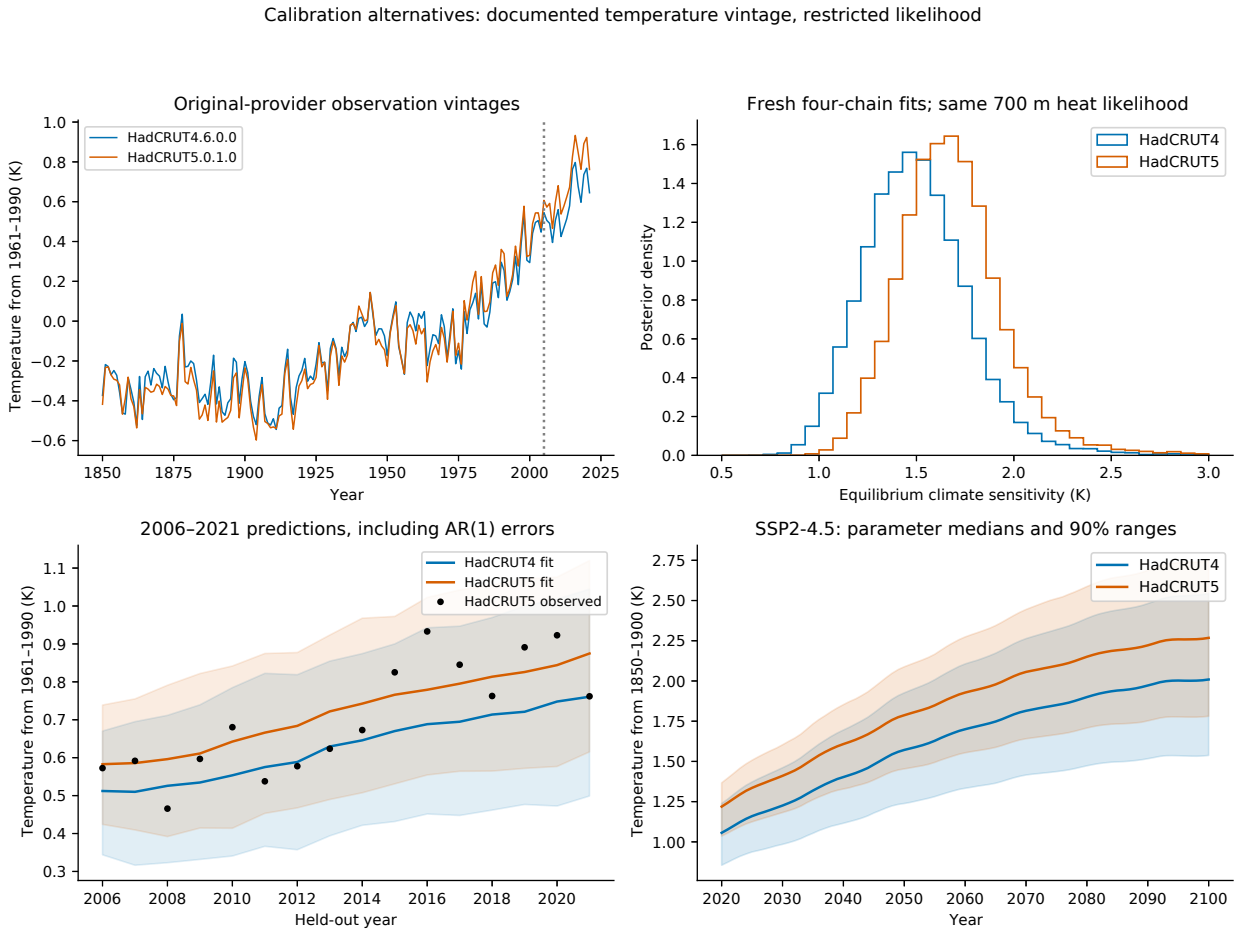


Figure 16: Observation-vintage alternatives. Upper left: original-provider temperature records on their common 1961–1990 datum; the dotted line marks the training cutoff. Upper right: independently sampled ECS posteriors under the same 700 m heat constraint. Lower left: held-out temperature predictions with stationary AR(1) errors and HadCRUT5 observations. Lower right: SSP2-4.5 parameter medians and 90% ranges relative to each model's 1850–1900 mean. The new fit changes the projection but does not recover the full target calibration.

does not establish optimizer convergence on real data. Zero-forcing and driven-mass tests verify the changed reference chemistry. The combined observational objective gives equal weight to the mean squared Mauna Loa and Law Dome residuals. This weighting is an explicit alternative, not the recovered target likelihood. It also omits the documented ice-core gas-age smoothing and dating uncertainty.

The original Mauna Loa-only fit reaches the lower initial-CO_2 bound and has a held-out RMSE of 1.79 ppm. Adding Law Dome reduces its early-record RMSE from 7.91 to 1.99 ppm but increases the held-out Mauna Loa RMSE to 4.59 ppm. Substituting HadCRUT5 gives a similar 4.19 ppm held-out error. Halving the integration step changes CO_2 by at most 2.13×10^{-4} ppm and carbon mass closes within 3.64×10^{-12} Gt C. The observational discrepancy is therefore much larger than the tested temporal error. These restricted fits are retained as failed predictive experiments and are not substituted into the coastal runs.

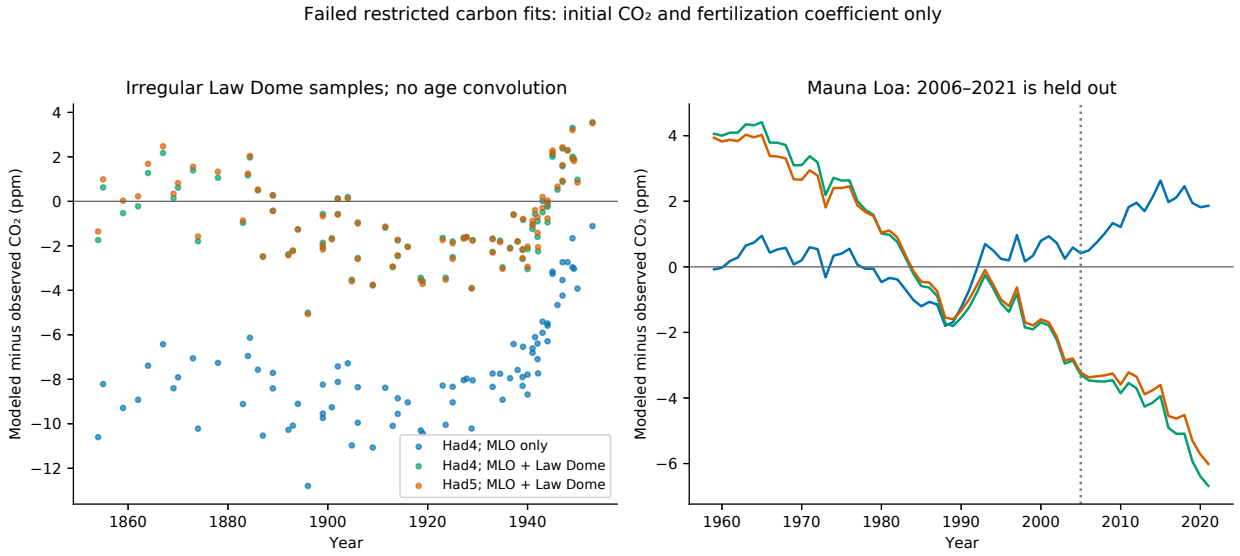


Figure 17: Original restricted carbon fits against independently acquired observations. Left: residuals at each irregular Law Dome gas age. Right: annual Mauna Loa residuals, with 2006–2021 held out beyond the dotted line. The combined fits improve the early record while degrading later predictions. The initial noise-free check passed, but later constrained-solver audits below distinguish computational limitations from predictive discrepancies. Numerical convergence alone cannot establish observational validation.

The first attempt to fit Q_{10} and thermocline transfer as additional coordinates failed an internal nested-model check: a clipped Gauss–Newton search stopped prematurely with training error above the restricted fit. Those trajectories and searches are preserved as failed computations. A second, bound-aware damped solver passes six analytic and nonlinear checks and includes the restricted fit as a starting point. Nevertheless, its real-data pilot reaches 100 iterations with projected gradient 0.0437 and objective 10.0613. This is an unconverged computation, not evidence for the fitted physics. The pilot is retained separately and its unfinished full run is interrupted. We therefore also evaluate the unchanged Julia residuals using the installed SciPy 1.5.4 bounded trust-region least-squares routine[1]. A pipe passes parameter and residual vectors; the library imports no climate-model implementation. Six cross-language tests check constrained analytical solutions and a nonlinear valley. Synthetic parameter recovery and comparison with the nested restricted model are required again before scientific interpretation. Four starting points test local-solution dependence

without proving a global optimum. The objective contains no held-out data or target-figure values.

Repeated carbonate roots dominate residual evaluation. We add an optional safeguarded Newton solution of the same equilibrium equation, retaining bisection as the default. Write $x = \log H$, $y = K_1/H$ and $z = K_1K_2/H^2$, where H is hydrogen concentration and K_1, K_2 are dissociation constants. If M is combined atmosphere–mixed-ocean carbon, atmospheric concentration is

$$C = \frac{M}{A + B(1 + y + z)}, \qquad \frac{dC}{dx} = \frac{CB(y + 2z)}{A + B(1 + y + z)}. \tag{11}$$

Here A converts atmospheric concentration to carbon and B combines mixed- ocean volume and solubility. Differentiating the existing alkalinity residual then supplies the Newton step; steps outside the log-hydrogen bracket are bisected. This changes the root algorithm, not the physical closure. The absolute alkalinity tolerance is 10^{-16} mol/L. Across 102 tests, isolated roots, mass/alkalinity balances and three historical integrations agree with bisection. The central 172-year benchmark differs by 7.43×10^{-11} ppm in CO_2 and 3.28×10^{-11} Gt C in stocks, with about sevenfold local speedup. The unfinished slower free-fit pilot is interrupted after this equivalence check; completed restricted bisection fits and all failure records are preserved.

9.4 Audited carbon fits and predictive disagreement

The SciPy pilot passes noise-free four-parameter recovery to 3.2×10^{-11} in scaled coordinates. The full experiment completes all twelve searches in 202 s, with no evaluation-limit stops. Three-point derivatives, Jacobian scaling and objective/step tolerances of 10^{-10} are recorded; the largest reported projected gradient is 1.55×10^{-4}. These are objective-based local stopping conditions, not a claim that every coordinate satisfies a 10^{-10} gradient criterion. Each selected training objective is below its corresponding re-audited restricted value. Replaying all twelve solutions verifies their objectives, agreement with bisection within 1.05×10^{-10} ppm, and time-step differences below 1.91×10^{-4} ppm. Selected fits conserve carbon within 3.96×10^{-12} Gt C.

Allowing respiration and thermocline transfer to vary improves the Mauna Loa-only held-out RMSE from 1.738 to 0.858 ppm, but its initial concentration remains at the 275 ppm lower bound (Table8). Adding Law Dome retains different predictive behavior: HadCRUT4 combined RMSE rises from 4.589 to 6.193 ppm and HadCRUT5 combined RMSE from 4.185 to 4.330 ppm. The best combined fits improve early-record RMSE to 1.736 and 1.979 ppm. Thus, lower training error does not resolve transfer to the later observational period. The likelihood weighting, gas-age mapping, initial disequilibrium and structural response remain distinct scientific questions.

Fit	C_0 ppm	β	Q_{10}	Exchange m/yr	CO_2 RMSE ppm	Ocean uptake Gt C
Had4, MLO	275.000	0.431	1.529	21.867	0.858	143.161
Had4, combined	282.899	0.721	1.064	11.313	6.193	94.417
Had5, combined	284.121	0.944	1.897	8.572	4.330	78.319

Table 8: Selected conditional carbon point fits. CO_2 RMSE uses withheld 2006–2021 observations; ocean uptake is modeled from the 1850 initial state to 1994. The Mauna Loa-only initial concentration reaches its lower bound. These are fitted alternatives, not a target posterior or identified constants.

The starting-point audit also exposes weak discrimination among physical responses. A second HadCRUT4 combined solution reaches the 200 m/year exchange bound, with objective 10.0805 instead of 10.0613, a difference of only 0.19%. Its modeled 1994 ocean uptake is 196.95 Gt C,

compared with 94.42 Gt C for the selected minimum, while held-out RMSE is 6.011 instead of 6.193 ppm. Selection uses the training objective throughout; held-out error is not used to choose a basin. The written 118 ± 19 Gt C inventory has a different preindustrial reference and is contextual evidence, not a matched likelihood or definitive rejection test. These differences prevent interpreting a stable minimum from several starts as a uniquely inferred carbon partition.

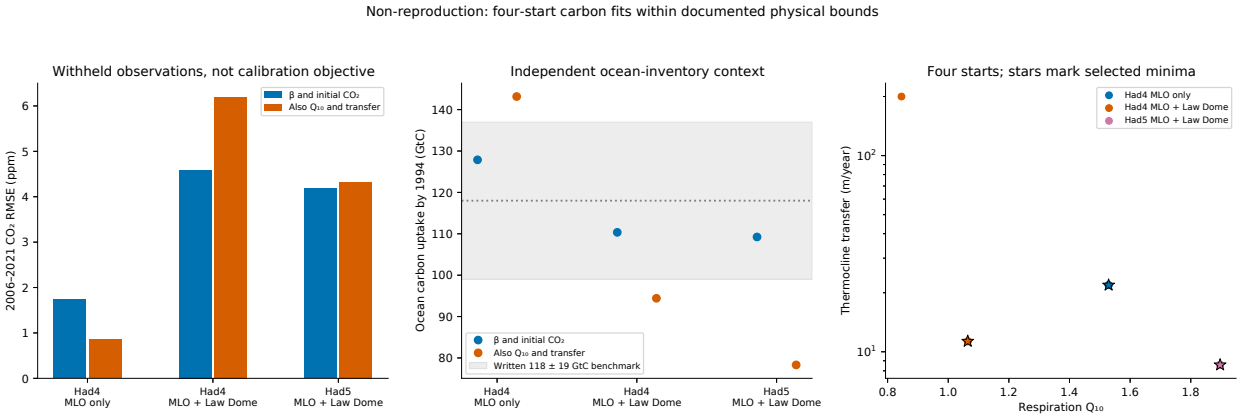


Figure 18: Audited carbon-parameter alternatives. Left: withheld CO_2 errors for restricted and four-parameter fits. Center: selected ocean uptake compared with the written inventory, whose time origin differs. Right: respiration and transfer values from all four starts; stars mark selected training minima. The alternative exchange-bound solution illustrates local-solution dependence. None of these fits replaces the carbon point parameters in the coastal ensembles.

9.5 An original-provider glacier record and a version discrepancy

The NSIDC archive supplies a 2005 observational supplement to Dyurgerov's glacier compilation [35]. Its Appendix 6 contains 43 annual global sea-level contributions for 1961–2003. We extract spreadsheet cells, verify annual-to-cumulative closure exactly, and cross-check the two populated worksheets. No figure is digitized. The global columns agree even though some regional entries differ. No annual global uncertainty column is supplied.

This archive is not confirmed identical to the cited Dyurgerov–Meier Occasional Paper 58. Its mean is 0.404 mm/year for 1961–2003 and 0.723 mm/year for 1994–2003, whereas the report's INSTAAR abstract states 0.51 and 0.93 mm/year. The old institutional download failed certificate validation, an earlier current link returned 404, and the government archive timed out. A subsequent search recovered the correctly capitalized institutional PDF URL; the printed observations are evaluated separately below. The NSIDC values remain unchanged.

A conditional least-squares experiment fits the glacier temperature sensitivity over 1961–1995 for 27 combinations of the documented volume, exponent and initial-state bounds, separately using HadCRUT4 and HadCRUT5 temperatures. The initial-state origin is explicitly assigned to 1961, and the observed annual increments are associated with intervals beginning in their labeled year. These choices, the rate objective and the absence of a joint AR(1) calibration make the experiment an alternative. A noise-free scalar fit recovers its known sensitivity to 4.4×10^{-18} m/(K year).

The selected HadCRUT4 fit has sensitivity 6.70×10^{-4} m/(K year) and withheld 1996–2003 RMSE 0.228 mm/year; the HadCRUT5 fit gives 6.10×10^{-4} and 0.241 mm/year. Both outperform a trend estimated from the same training years, whose withheld RMSE is 0.355 mm/year. The training minima occur on volume, exponent and initial-state grid boundaries; we do not interpret them as

identified physical parameters or a global optimum. Comparison with GlaMBIE in 2004–2020 gives RMSE 0.230 and 0.220 mm/year, but the observation products have different spatial coverage. They differ by 0.389 mm/year RMS during their four overlapping years. No fitted parameters from this sensitivity replace the accepted coupled sea-level calculations.

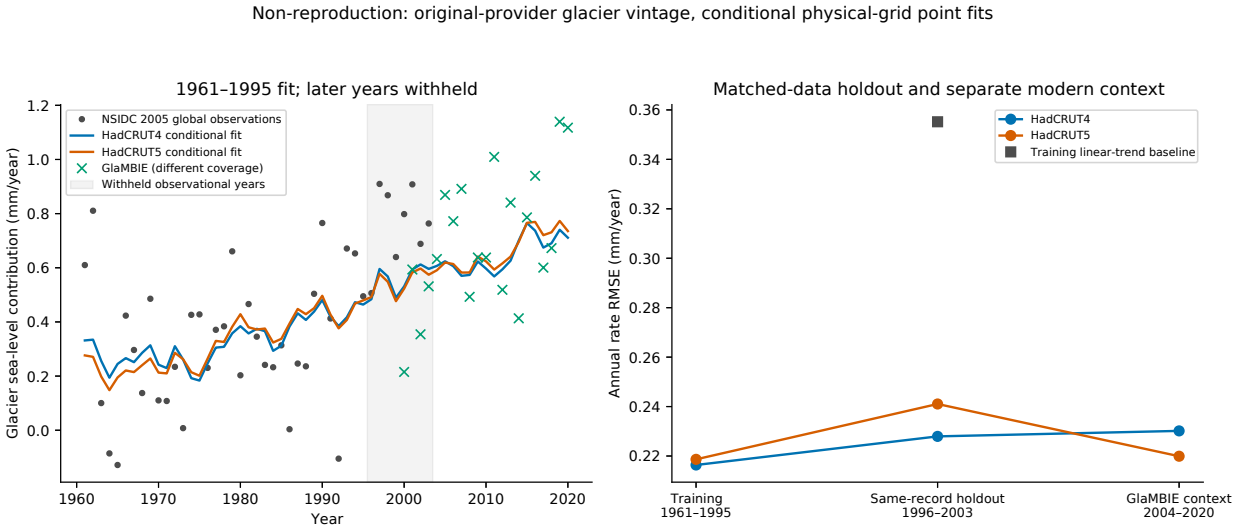


Figure 19: Glacier observation-vintage test. Gray points are the original NSIDC spreadsheet observations; colored lines are independently fitted conditional model rates. The shaded interval is withheld from fitting. Crosses show the newer GlaMBIE observations, which have different coverage. Right: training and withheld errors, with the later product treated as separate context. The archived NSIDC series differs from the cited report's written average; this is a non-reproduction.

9.6 Recovery of the cited glacier report

The institutional OP58 report is now available at a correctly capitalized URL [34]. Appendix 3, printed page 117, supplies annual global observational SLR increments and cumulative values for 1961–2003. Text-table extraction gives means of 0.511 and 0.927 mm/year over the two stated intervals. Annual increments and cumulative values are independently rounded to 0.01 mm, so individual closure discrepancies reach 0.01 mm. We preserve that rounding and verify it against the printed volume conversion; no figure is digitized. The source includes peripheral glaciers and supplies no annual global uncertainty column in this table. Recovering it therefore resolves the unavailable-record claim, while leaving the target likelihood, processed input identity and component coverage conventions to be established.

Repeating the same conditional 54-point glacier experiment on this record gives HadCRUT4 and HadCRUT5 sensitivity estimates of 8.25×10^{-4} and 7.52×10^{-4} m/(K year). The withheld 1996–2003 rate RMSE is 0.400 and 0.412 mm/year, compared with 0.578 for the training linear trend. The selected physical grid points again lie on all three boundaries. Their 2004–2020 GlaMBIE contextual RMSE is 0.201 and 0.202 mm/year, with unmatched coverage. The original observational products differ by 0.730 mm/year RMS in 2000–2003. This is a reproducible observation-version comparison, not a fitted error bar or recovery of the target glacier posterior.

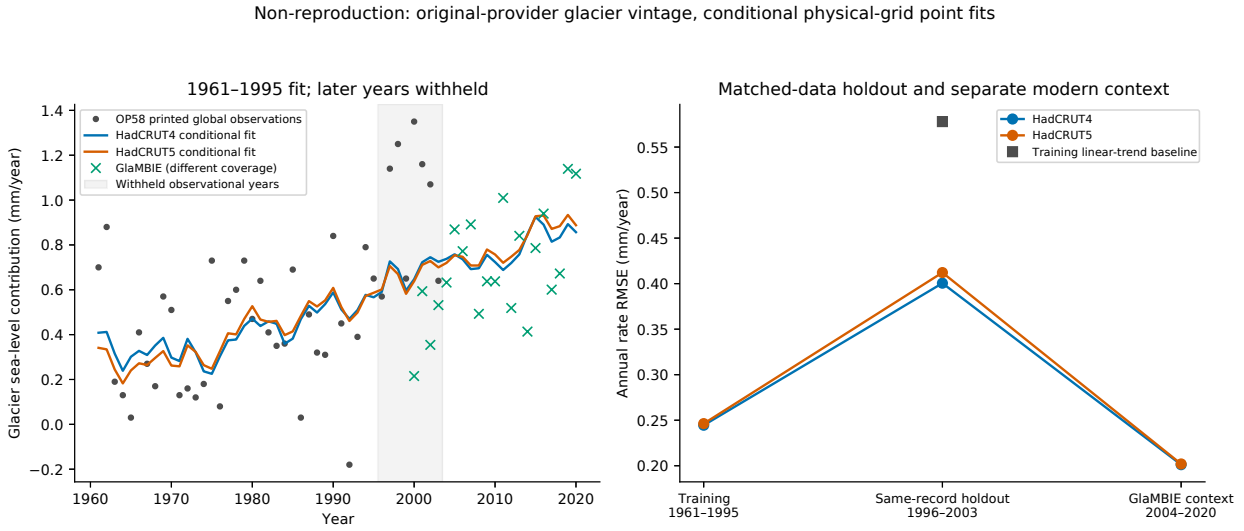


Figure 20: Conditional glacier fit using the recovered OP58 printed observations. The same grid, training boundary and withheld years as the NSIDC experiment isolate the source-version change. The later GlaMBIE record is separate context. No fitted glacier parameters in this figure replace the base sea-level ensemble.

9.7 Propagation of the temperature-version change

The independently calibrated HadCRUT5 posterior now drives 250 emissions states and 2,500 sea-level cases under the same five closure choices and two non-CO_2 rules. Its pilot checks give maximum CO_2 and temperature differences of 4.82×10^{-4} ppm and 9.05×10^{-6} K when halving the step; Antarctic refinement differs by 2.45×10^{-13} m. The carbon hindcast and point parameters remain fixed so that these runs isolate the posterior change.

Under the explicit 50% budget-event weights, median annual 2100 warming rises from 1.264 to 1.484 K. Its 5th–95th range changes from 0.838–2.606 K to 0.996–2.879 K. The corresponding median uncorrected sea-level rise from 2010 changes from 0.232 to 0.237 m, while its upper percentile rises from 0.641 to 0.736 m. The fixed Antarctic model's triggering probability rises from 13.99% to 18.77%. Category counts become 101, 56, 71 and 22; their weighted masses are 54.59, 34.26, 10.76 and 0.393%. These are conditional alternatives, not estimates under the recovered expert-density likelihood or the target posterior. The additional high-warming states motivate a separately validated coastal propagation rather than assuming that the earlier cost distribution transfers.

9.8 Separating heat constraints from expert-emissions assumptions

The temperature-only posterior supplies a third, deliberately unfavorable predictive comparison. We propagate the same 250 emissions states through each independently sampled climate posterior, holding carbon and ice choices fixed. Parameter values are not paired between posteriors; only the emissions design and random row indices are shared. Figure21 therefore compares distributions, not the response of a single physical state to changing an observation.

Omitting the upper-ocean heat constraint raises median annual 2100 warming under the 50% budget-event alternative to 2.156 K, with a 5th–95th range of 1.261–3.935 K. Warming-category counts are 25, 84, 63 and 78, with weighted masses 13.02, 47.39, 29.71 and 9.87%. The corresponding uncorrected sea-level median is 0.450 m from 2010, with range 0.244–1.423 m, and the fixed Antarctic

fast-loss model activates in 43.91% of weighted paths. The high-category mass is closer to the target's printed 7.78% than either heat-constrained result. However, this posterior has withheld temperature RMSE 0.209 K, bias +0.189 K and upper-ocean heat RMSE 16.46×10^{22} J. Its failure to predict these independent observations prevents using category resemblance to select it as a recovered calibration.

The three climate ensembles pass the same time-step and carbon-accounting checks. The temperature-only pilot's maximum changes are 5.10×10^{-4} ppm CO_2 and 8.98×10^{-6} K; carbon mass closes within 1.28×10^{-11} Gt C over the full ensemble. Thus, these distribution differences exceed the measured numerical time-step effects. The experiment tests the influence of an alternative 700 m constraint. It does not recover the missing annual 3000 m observational record.

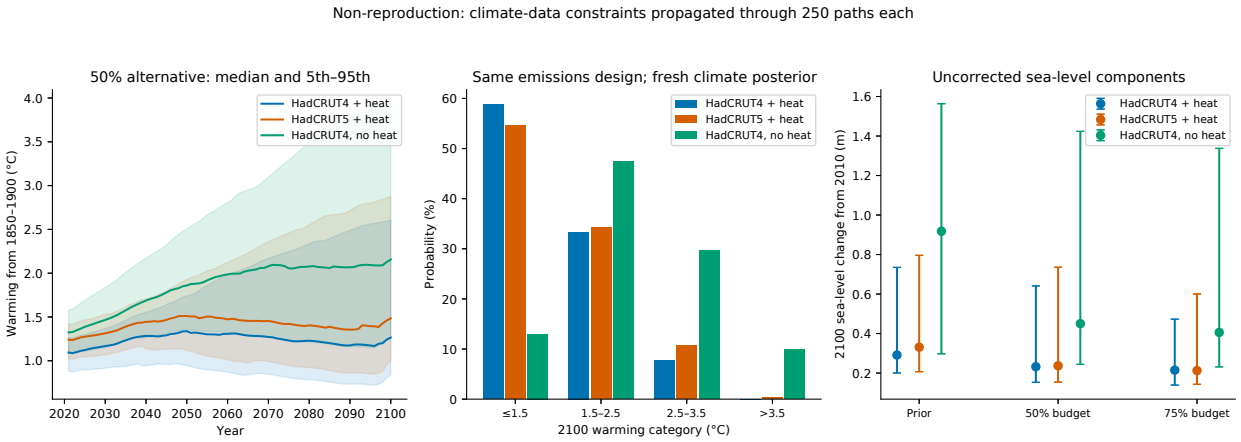


Figure 21: Climate-observation constraints propagated through independent emissions-driven alternatives. Left: annual warming medians and 5th–95th percentiles under explicit 50% budget-event weights. Center: 2100 category masses under those weights. Right: uncorrected global sea-level medians and 5th–95th percentiles under three alternative emissions measures, relative to 2010. Colors distinguish independently calibrated climate posteriors. The temperature-only result fails predictive checks; none is the target posterior.

Applying the recovered prior-times-expert-density inversion changes the distributions again (Figure22). Across the four explicitly assumed densities, HadCRUT4 with heat gives only 0.0072–0.0399% probability above 3.5 K; HadCRUT5 with heat gives 0.193–0.450%. Temperature-only fits give 17.09–18.09%, despite all four expert densities assigning the same 50% mass to the cumulative-emissions budget event. Median sea-level changes range from 0.277–0.291 m, 0.321–0.334 m and 0.867–0.903 m, respectively. These are different likelihood assumptions from directly imposing a posterior budget-event probability. The spread identifies two consequential unresolved choices: the climate likelihood and the complete expert density. It supplies no basis for choosing either to force a match to the target.

For comparison with the target's reference period, each sea-level member is also shifted by its own modeled 1995–2014 arithmetic mean. Under 50% budget-event weights the HadCRUT5 median becomes 0.250 m, with 17th–83rd percentiles 0.193–0.457 m. The temperature-only median becomes 0.465 m, with interval 0.322–1.178 m. The datum corrections are 0.0065–0.0180 m and 0.0119–0.0262 m across the five closure alternatives. These shifts reconcile reference periods only; they do not make the emissions ensemble commensurate with an individual AR6 SSP scenario or remove component-coverage overlap.

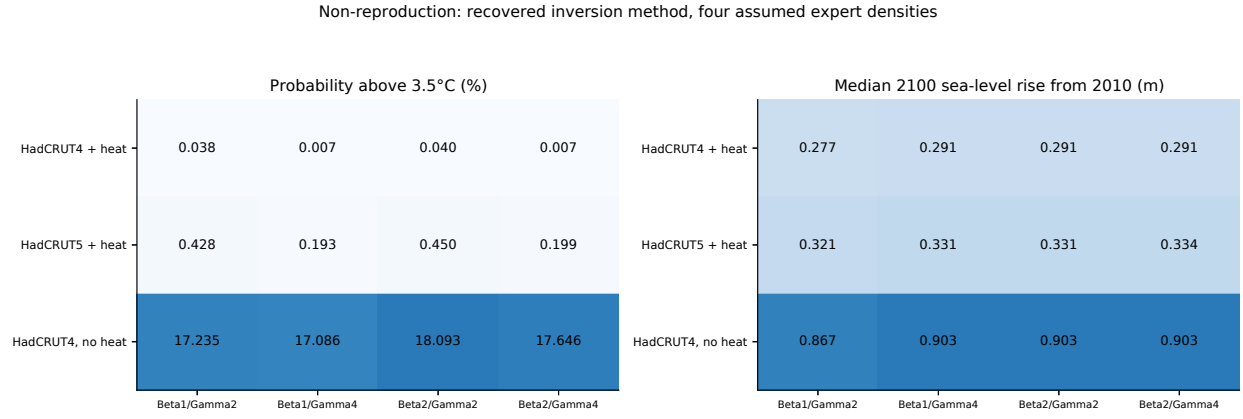


Figure 22: Recovered inversion method with four alternative expert probability densities and three independent climate calibrations. Entries show probability above 3.5 K and median uncorrected sea-level change from 2010. The lower-budget density is uniform or Beta(2,2); the upper-budget tail uses Gamma shape two or four, with the documented conditional mode and budget mass imposed. The families and widths are assumptions. No target posterior is inferred from this comparison.

10 Antarctic boundary recovery and conditional hindcasts

10.1 Recovered initial state and alternative forcing

The initial Antarctic state is specified more fully than the first assessment recognized. Shaffer's hindcasts begin at 240 kyr before CE 2000 with radius 1.8636×10^6 m and convert 24.78×10^{15} m^3 to 57 m of sea-level equivalent[9]. We now integrate this initial-value problem using independent proxy and instrumental inputs. NOAA Dome C and Spratt records use years before 1950; they are converted to calendar years before comparison with the model's different age origin. The W17F calibration dates of 118000,18000 and 4000 BCE correspond to 120,20 and 6 kyr before CE 2000[25].

Antarctic air temperature combines Dome C anomalies, the global land–ocean Mann EIV reconstruction, and HadCRUT4. We retain Shaffer's −0.1 K Dome C adjustment, independently rebase Mann to 1961–1990, and apply the documented modern amplification factor 1.2 around −18°C. The independently recomputed 500–1500 CE overlap adjustment would instead be −0.168 K; the reconstruction variant and datum therefore remain explicit uncertainties. The main alternative sea-level boundary subtracts Spratt's youngest 8.49 m value and connects its CE 950 point to the first available Church–White annual value at 1880.5. This interpolation is a declared assumption, not the recovered Lambeck splice. An unshifted Spratt case and both archived Waelbroeck chronologies test boundary sensitivity. These are observational proxy products, not recycled DAIS hindcasts[26, 27, 28, 29, 30].

The original transient subsurface-ocean history remains unrecovered. Shaffer's written steady-state approximation supplies a useful alternative:

$$T_o = 0.00690 T_a^2 + 0.439 T_a + 6.39 \quad [^\circ\mathrm{C}]. \tag{12}$$

Here T_a is Antarctic air temperature reduced to sea level. Applying this formula during a transient omits ocean thermal memory. It permits a test of the other equations without claiming recovery of the original DCESS forcing.

10.2 Volume integration and switching dynamics

The implementation advances physical ice volume by accumulation plus marine discharge and independently inverts the verified geometric relation to obtain radius at each integration stage. A separate radius-coordinate integration, including the prescribed sea-level derivative, agrees with the volume solution. All forcing knots and temperature gates split the integration grid; endpoint stages use the branch valid within that interval.

The additional fast-loss flux follows the recovered temperature threshold and stops below 18×10^{15} m^3[25]. Its cutoff creates a discontinuous differential equation. When slow accumulation would increase volume at the cutoff but full fast loss would decrease it, the mathematical sliding solution holds volume at the cutoff by balancing those fluxes. Slow mass loss can still lower volume below the cutoff. This continuation is derived and tested on exact piecewise trajectories; it is not an assertion about the original authors' timestep convention. Nineteen checks cover the geometry inversion, coordinate equivalence, constant forcing, hitting times, sliding and departure from the boundary.

For the central fixed-parameter hindcast, reducing the maximum step from 10 to 5 years changes saved sea-level equivalent by less than 8×10^{-8} m. That numerical agreement does not remove the physical mismatch: the predicted Last Interglacial contribution is -0.17 m, outside the written 3.6–7.4 m window. Figure23 compares input histories and their effects on this result. It also tests initialization and the sharply different responses obtained by changing the fast-loss threshold.

10.3 Conditional fitting and withheld annual rates

A separate experiment varies discharge exponent γ, ocean partition α, fast-loss rate λ and threshold T_c over their documented uniform ranges, fixing the remaining Shaffer parameters. Independent Latin hypercube draws assess parameter coverage before a bounded coordinate search minimizes the four standardized squared window residuals. The Last Interglacial hard range is also checked. This restricted point fit omits additional residual-process variance, uncertain forcing and the target joint calibration; it is neither a posterior nor an independent test of those four fitted values.

Annual IMBIE mass-loss observations after 2002 provide a separate conditional check. Observed temperatures remain prescribed in these years, and the calibration's modern cumulative constraint ends in 2002. Figure24 reports both parameter coverage and the later annual-rate errors. A good fit to four widely separated volume constraints need not predict short-period mass-loss variability.

Constraint	Written mean (m)	Written SD (m)	Point fit (m)
Last Interglacial	5.5	0.95	5.5
Last Glacial Maximum	-11.35	2.23	-11.392
Mid-Holocene	-2.63	0.68	-2.7276
1992–2002 increase	0.00197	0.00046	0.001951

Table 9: Conditional Antarctic predictions under the explicitly alternative forcing. These four quantities are fitted, not held-out validation. The Last Interglacial hard acceptance interval is 3.6–7.4 m.

The conditional solution uses $\gamma = 1.93747$, $\alpha = 0.02077$, $\lambda = 0.01161$ m yr^{-1} and $T_c = -15.62980\,°C$. The search begins with 512 Latin hypercube points; 1 draws satisfy all four windows before local optimization. The final four-window χ^2 is 0.02267. Halving the maximum timestep from 10 to 5 years changes these predictions by at most 2.7e-09 m. The conditional 2003–2008 mass-loss RMSE is 32.35 Gt yr^{-1}, with bias -25.22 Gt yr^{-1}. The observed and predicted mean rates are 84.17 and 58.96 Gt yr^{-1}, respectively.

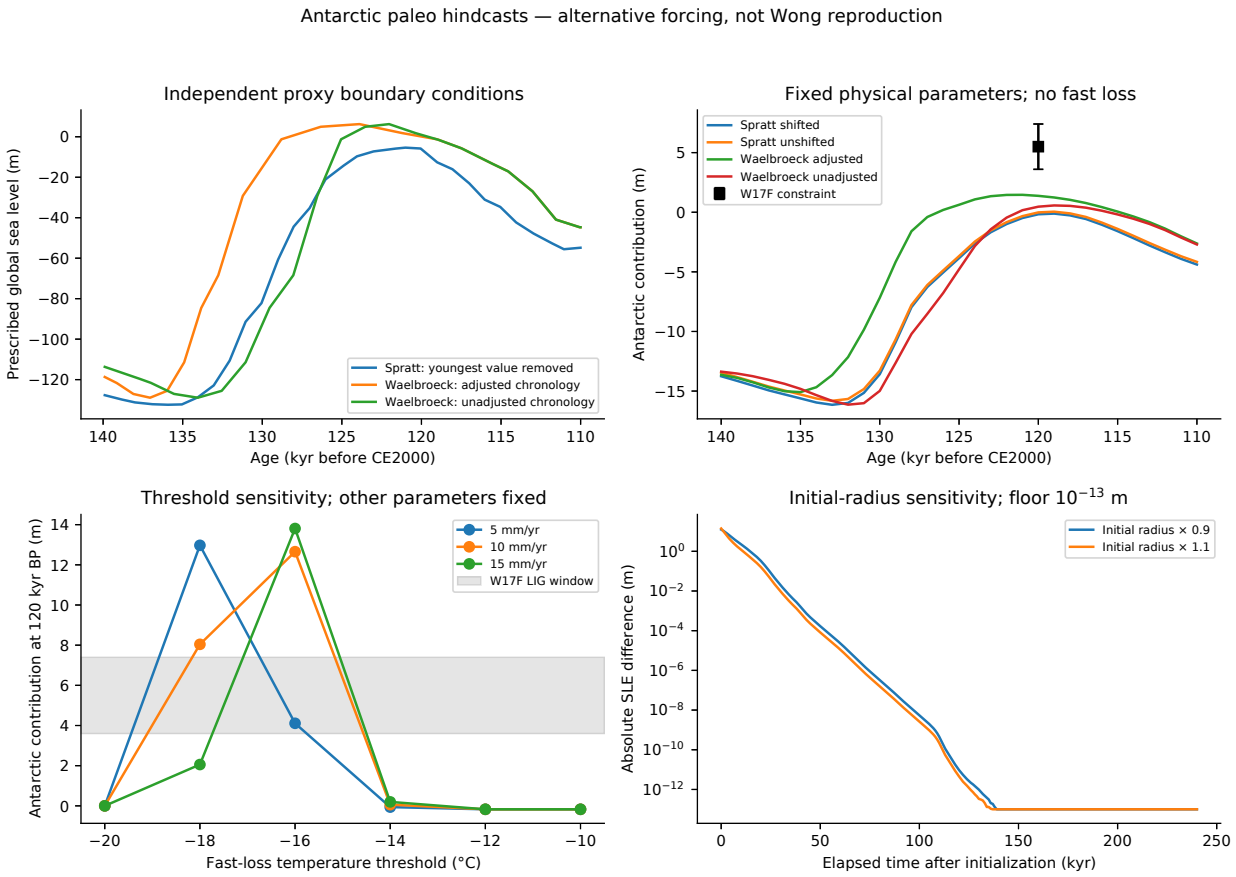


Figure 23: Antarctic hindcasts under explicit alternative forcing. Upper left: independent sea-level proxy inputs on a common calendar. Upper right: fixed Shaffer case 4 parameters without fast loss, with the W17F Last Interglacial constraint shown at 120 kyr before CE 2000. Lower left: fast-loss rate and threshold sensitivities with other parameters fixed. Lower right: differences after changing initial radius by 10%; the plotting floor denotes roundoff, not persistent physical differences. All panels are non-reproductions of Wong's calibrated Antarctic ensemble.

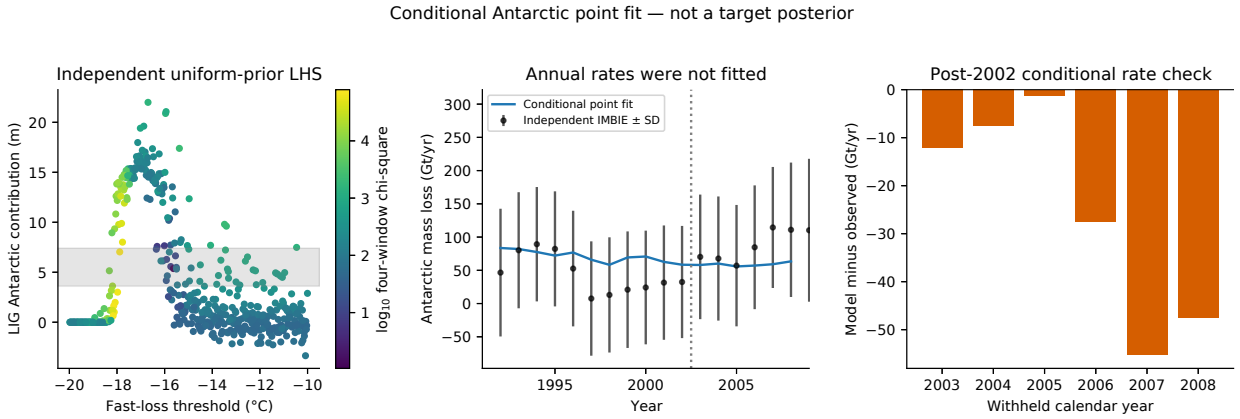


Figure 24: Conditional Antarctic fitting under alternative forcing. Left: independent uniform-prior Latin hypercube draws, coloured by four-window misfit; shading marks the Last Interglacial range. Centre: fitted-model annual mass loss and independent IMBIE observations with reported uncertainty. Right: model minus observed rates after the 2002 cumulative constraint. No posterior interval or recovered Wong result is implied.

The earlier64-point pilot reaches a different solution, with fast-loss rate 0.015 m yr^{-1} and threshold$-15.484\,^{\circ}\mathrm{C}$, while its four-window $\chi^2 = 0.02354$ is almost unchanged. The larger search gives a lower misfit without reaching that rate bound. These two fits illustrate conditional parameter ambiguity; neither is presented as a uniquely inferred fast-loss mechanism.

BRICK's written shore coupling is also recoverable: the Antarctic contribution has fingerprint-1, while the other contributions have factor$+1$[7]. Consequently, if S_o is the sum of other components and S_A is Antarctic global sea-level change, the local boundary changes as $S_o - S_A$. A volume integrator applies this feedback through geometric inversion at each stage. An independent radius-coordinate calculation includes its derivative in $\dot{S}$; ten checks cover that equivalence, equilibrium, timestep refinement and component accounting. A lagged-input alternative follows the documented previous-year forcing convention, while retaining an explicitly independent within-year numerical integrator.

11 Emissions probabilities and coupled future alternatives

11.1 Recovered priors and unresolved expert density

The written emissions constraints are more specific than the initial blocker assessment implied. Darnell's manuscript and its published successor specify the cumulative interval as 2022–2100 and describe annual selection of the nearest SSP CO_2 trajectory[4]. The prior peak year can occur after the simulation horizon. The implementation now retains such draws as growing trajectories through the last simulated year, rather than rejecting valid prior mass. Ten thousand independent draws contain eight such peaks after 2150. The starting net emissions, including the documented cement-carbonation adjustment to Global Carbon Budget 2023, are 40.663 $\mathrm{GtCO_2\ yr^{-1}}$ in 2022.

Only 2.35% of these prior draws satisfy the 1223 $\mathrm{GtCO_2}$ cumulative budget. The written 50% and 75% expert assessments therefore imply substantial reweighting. They do not specify a unique density, and the baseline requirement that SSP4-6.0 be the conditional mode above the budget imposes information beyond a single event probability. To isolate this ambiguity, two explicit

alternatives minimize relative entropy with respect to the empirical sample from the documented prior subject only to the budget-event probability. If A denotes draws within the budget, n_A its size and p the imposed probability, the weights are

$$w_i = \begin{cases} p/n_A, & i \in A, \\ (1-p)/(n-n_A), & i \notin A. \end{cases} \tag{13}$$

This formula follows from Lagrange multipliers and leaves conditional distributions within each class unchanged. It does not recover the elicited expert density or the SSP4-6.0 mode. The effective weighted sample sizes are 918 and 417 for $p = 0.5$ and 0.75. Changing peak-year rounding to floor or ceiling changes budget membership for 28 or 21 draws; trapezoidal rather than annual-mean summation changes 26 memberships. These sensitivities are retained explicitly.

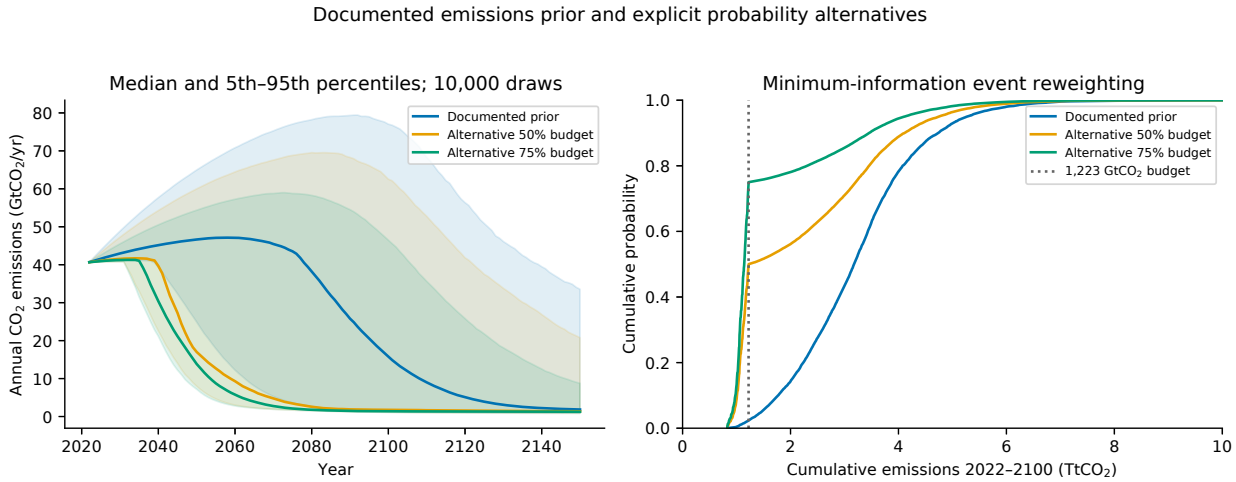


Figure 25: Documented emissions priors and explicit alternative weighting. Left: weighted medians and 5th–95th percentiles of annual emissions. Right: cumulative budget distributions and the 1223 $GtCO_2$ threshold. All distributions use the same 10,000 independent draws; only their weights differ. The two reweighted distributions are non-reproductions of the published emissions posterior.

11.2 Independent SSP inputs and carbon–climate coupling

The RCMIP protocol 5.1.0 supplies independent exogenous SSP emissions rather than Wong model outputs[31]. Total global CO_2, fossil and industrial CO_2, and land-use CO_2 are extracted for six SSPs and linearly interpolated between supplied annual nodes. Their partition identity closes to 2.9×10^{-14} $GtCO_2$ yr^{-1}. No extrapolation is needed through 2150. The interpretation tested here selects, separately each year, the SSP with smallest absolute total-emissions difference. CEDA non-CO_2 forcing from that scenario is then linearly interpolated between annual nodes. A paired fixed-SSP2-4.5 calculation measures the consequence of this selection. The written nearest-trajectory statement alone does not settle all interpolation and distance conventions.

The future calculation starts from the conditional carbon hindcast and the accepted climate posterior's physical temperature state in 2021. Carbon feedback temperature is transformed to the same preindustrial datum used in the carbon fit. The future CO_2 forcing uses $3.7 \log_2(C/280)$ W m^{-2}; its difference from the historical CEDA forcing is reported as a restart discontinuity, without compensating adjustment. The carbon fertilization parameter is fixed at the independently fitted

central value; uncertainty in this parameter and its dependence on climate parameters are omitted. The resulting distribution is therefore conditional and cannot stand in for the target joint posterior.

A balanced conditional sample contains 128 draws on each side of the budget. Complete accepted climate-posterior rows are paired independently with these emissions draws. Each probability measure retains its own class weights. An eight-member pilot with both forcing rules verifies carbon conservation and compares eight with sixteen coupling substeps per year. The maximum differences are 1.45×10^{-5} K in temperature and 3.96×10^{-4} ppm in CO_2; these checks precede the 256-member calculation. Warming categories use the annual 2100 temperature relative to 1850–1900, with boundary values assigned to the lower category exactly as in the publication.

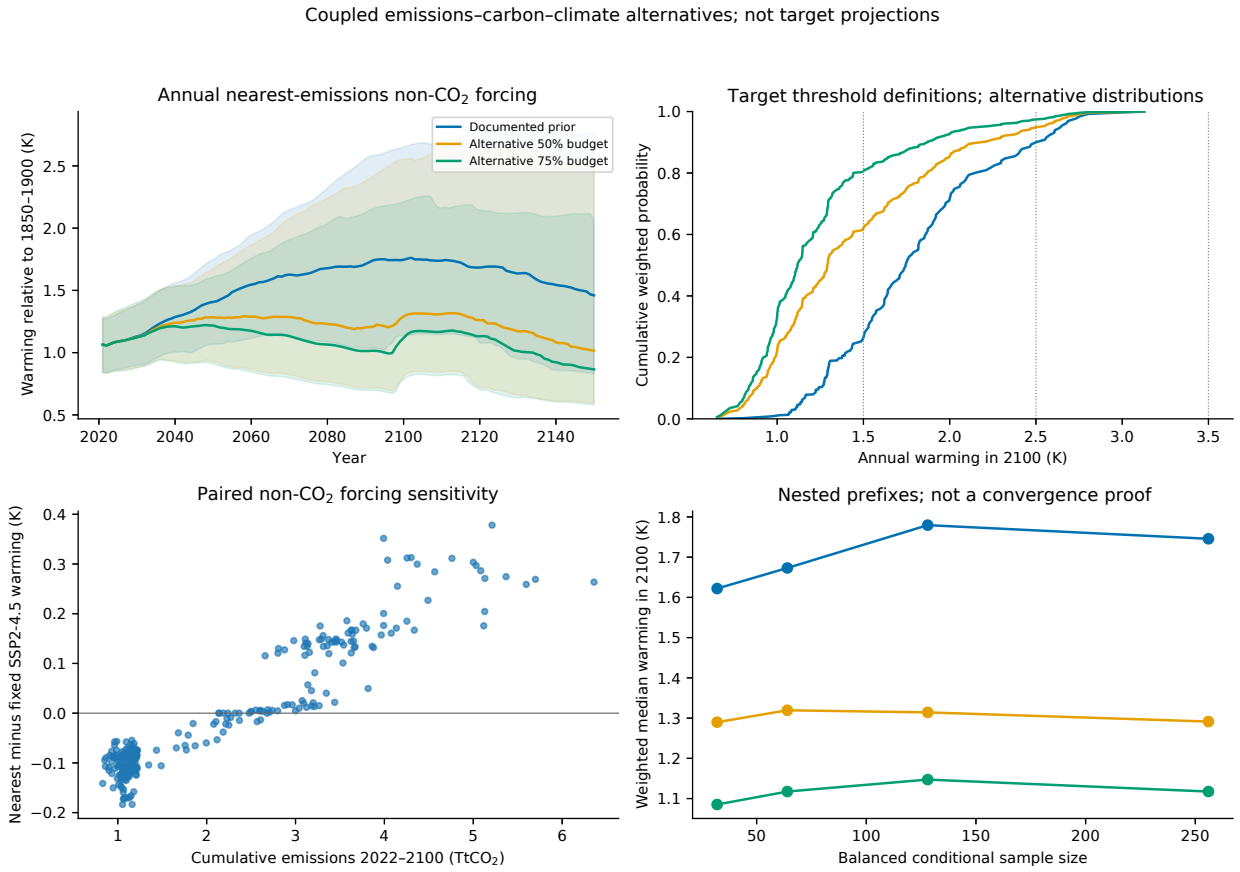


Figure 26: Coupled alternative future experiments. Upper left: median and 5th–95th temperature ranges for the three weighting measures. Upper right: weighted 2100 temperature distributions and the published category boundaries. Lower left: paired sensitivity to annual nearest-emissions versus fixed-SSP2-4.5 non-CO_2 forcing. Lower right: medians for nested balanced sample sizes; stability of these medians alone is not evidence of tail convergence. Carbon parameters are conditionally fixed and climate rows are independently paired, so no panel is a recovered published projection.

The largest paired change in 2100 warming from nearest-scenario versus fixed-SSP2-4.5 non-CO_2 forcing is 0.378 K. The restart changes the CO_2 forcing by -0.1762 W m^{-2}, without compensation. Maximum carbon accounting error is 8.3e-12 Gt C.

No member of this256-member calculation exceeds3.5 K. The zero entry in the fourth category therefore denotes absence in this finite sample, not proof of zero probability. A separate tail-stratified emissions experiment is motivated by this limitation; it preserves empirical prior class masses with

Distribution	$T_{.05}$	$T_{.50}$	$T_{.95}$	$P_1(\%)$	$P_2(\%)$	$P_3(\%)$	$P_4(\%)$
Prior	1.136	1.746	2.650	26.0	63.3	10.7	0.0
50% alternative	0.811	1.291	2.552	61.7	32.8	5.5	0.0
75% alternative	0.784	1.117	2.211	80.5	16.8	2.7	0.0

Table 10: Alternative annual warming in 2100 relative to 1850–1900, in K. Probabilities refer to the four published threshold definitions. The balanced 256-member subsample has distinct weights for each probability measure; it is not the published 5,000-member ensemble.

explicit weights rather than changing parameters to create high warming.

11.3 Recovering the probabilistic-inversion convention

The cited method of Fuller et al. [36] uses the expert assessment as a likelihood evaluated at the model output. For emissions parameters θ, prior $p_0(\theta)$, cumulative emissions $C(\theta)$ and expert density f, the resulting density is

$$p(\theta \mid f) \propto p_0(\theta) f[C(\theta)]. \tag{14}$$

Consequently, assigning half the expert density below a carbon budget does not require half the posterior model outputs to fall below that budget. The earlier event-weighting experiment is a different statistical assumption.

We test four explicit expert-density alternatives. Each assigns half its mass below $B = 1223$ $GtCO_2$, using either a uniform or scaled Beta$(2, 2)$ density on $(0, B)$. Above B, a shifted gamma density has shape 2 or 4 and scale $(M - B)/(k - 1)$, where k is its shape. Thus its conditional mode is $M = 3263.330$ $GtCO_2$, the independently computed annual SSP4–6.0 sum for 2022–2100. The two documented assessment constraints are respected, but the density families and shapes are alternatives, not recovered target inputs. Integration and mode tests verify all four densities.

Reweighting the 10,000 independent prior draws yields posterior budget probabilities of 4.56–11.31%, despite the common 50% assessment mass. Reweighting the 250 propagated paths gives 4.36–11.05%; their normalizing constant differs from the larger prior sample by 2.42–5.22%. Effective sample sizes are 65–83 for the propagated paths. Median 2100 warming is 1.78–1.85 K, and unconditional median incremental coastal cost is 0.266–0.277 trillion 2019 USD PPP. These calculations identify a consequential statistical distinction without identifying the missing target density.

12 Coupled sea-level alternatives and tail coverage

12.1 Sampling rare emissions without changing their probability

The balanced 256-member calculation contains no warming above 3.5 K. To examine rare emissions without altering their probability, we partition the same 10,000 prior trajectories into five strata. These contain 235 budget-compliant draws, 9,265 remaining draws through the overall 95th percentile, 400 between the 95th and 99th percentiles, 90 between the 99th and 99.9th percentiles, and 10 above the 99.9th percentile. We sample 96, 64, 48, 32 and 10 trajectories, respectively, without replacement. The prior and both budget alternatives retain their empirical stratum masses. Complete climate-parameter rows are paired independently with these emissions draws.

The initial eight-member tail pilot fails the declared 0.001 ppm CO_2 refinement criterion: doubling eight to sixteen substeps changes CO_2 by 0.001908 ppm. We retain that failed result

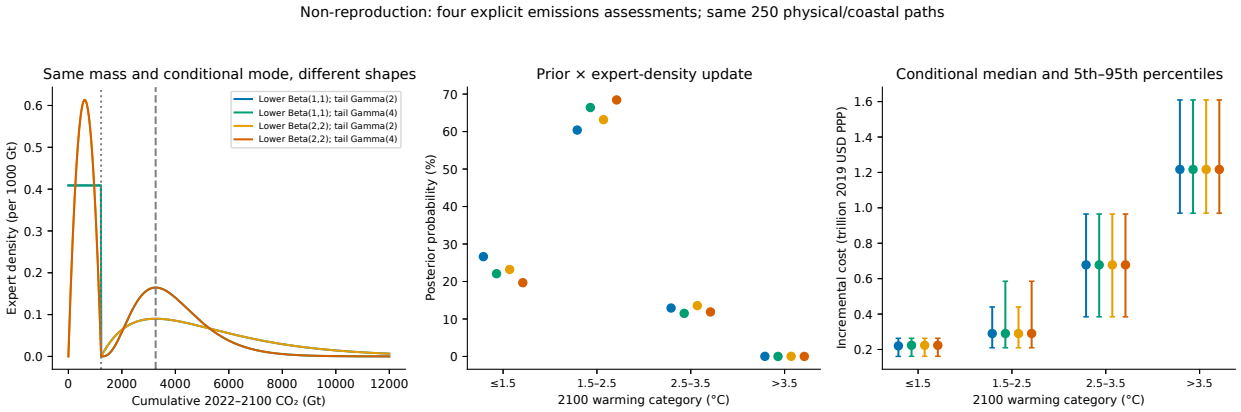


Figure 27: Four explicitly assumed expert densities and their propagation through the same independent climate, sea-level and coastal calculations. Left: equal assessment mass below the budget and the same conditional above-budget mode do not determine a unique density. Middle: warming probabilities after multiplying the parameter prior by the expert density. Right: conditional coastal cost medians and 5th–95th order statistics. Five physical paths remain in the highest warming category; reweighting does not resolve its tails. All panels are non-reproductions.

and double the working resolution. The revised sixteen-versus-thirty-two-substep pilot changes temperature by at most 9.08×10^{-6} K and CO_2 by 4.77×10^{-4} ppm. Carbon mass closes within 8.19×10^{-12} Gt C. The subsequent 250-member experiment uses sixteen substeps and takes 512 wall seconds on one thread.

The four warming categories contain 107, 65, 73 and 5 members. Under the 50% budget alternative, their probabilities are 58.81%, 33.35%, 7.78% and 0.0534%. The last category's conditional effective sample size is only 4.23; its endpoints cannot establish converged 5th or 95th percentiles. Tail sampling therefore resolves the earlier absence of high warming as a finite-sample limitation, but does not recover the published warming probabilities.

12.2 Component assembly and Antarctic shore feedback

BRICK's written description supplies Antarctic-shore factors of -1 for the Antarctic contribution and $+1$ for other components, together with previous-year forcing in its annual coupling[7]. These local factors permit a shore feedback calculation without inventing global fingerprint maps. If A is the Antarctic global contribution and O is the sum of other contributions, the shore level is

$$S_{\text{shore}}(t) = S_{\text{shore}}(t_0) + O(t) - A(t), \qquad O(t_0) = A(t_0) = 0. \tag{15}$$

We evolve Antarctic volume and invert its verified geometric relation at each stage. A separate radius-coordinate integration includes the resulting sea-level derivative and agrees with the volume calculation. Continuous stagewise feedback is an explicitly derived alternative; freezing the forcing at the previous annual knot tests the documented lag convention.

The future budget combines conditional glacier and Greenland parameter rows, the Antarctic point fits, thermosteric change and land-water storage. The Antarctic splice begins at the final paleo-forcing knot, 2009.5, and subsequent shore changes follow the equation above. Glacier and Greenland histories use observed temperature through 2021. Future forcing retains the observed boundary temperature before joining the modeled series. Storage has a documented mean rate of 0.30 mm yr^{-1} and standard deviation 0.18 mm yr^{-1}; Gaussian independent annual rates and a

common uncertain rate are separate assumptions. Thermosteric alternatives apply the independently fitted 0–2000 m coefficient either to that depth alone or to the full modeled ocean heat change. Neither alternative recovers a full-depth target posterior.

Five closure choices and two non-CO_2 forcing rules give 2,500 member cases. The initial pilot checks halved Antarctic timesteps and component sums before the 250-member extension. The base experiment produces 0.111–1.448 m of global change by 2100 relative to 2010. This is an *uncorrected alternative budget*: the observational peripheral-glacier overlap, conditional parameter fits and omitted cross-component dependence remain explicit limitations.

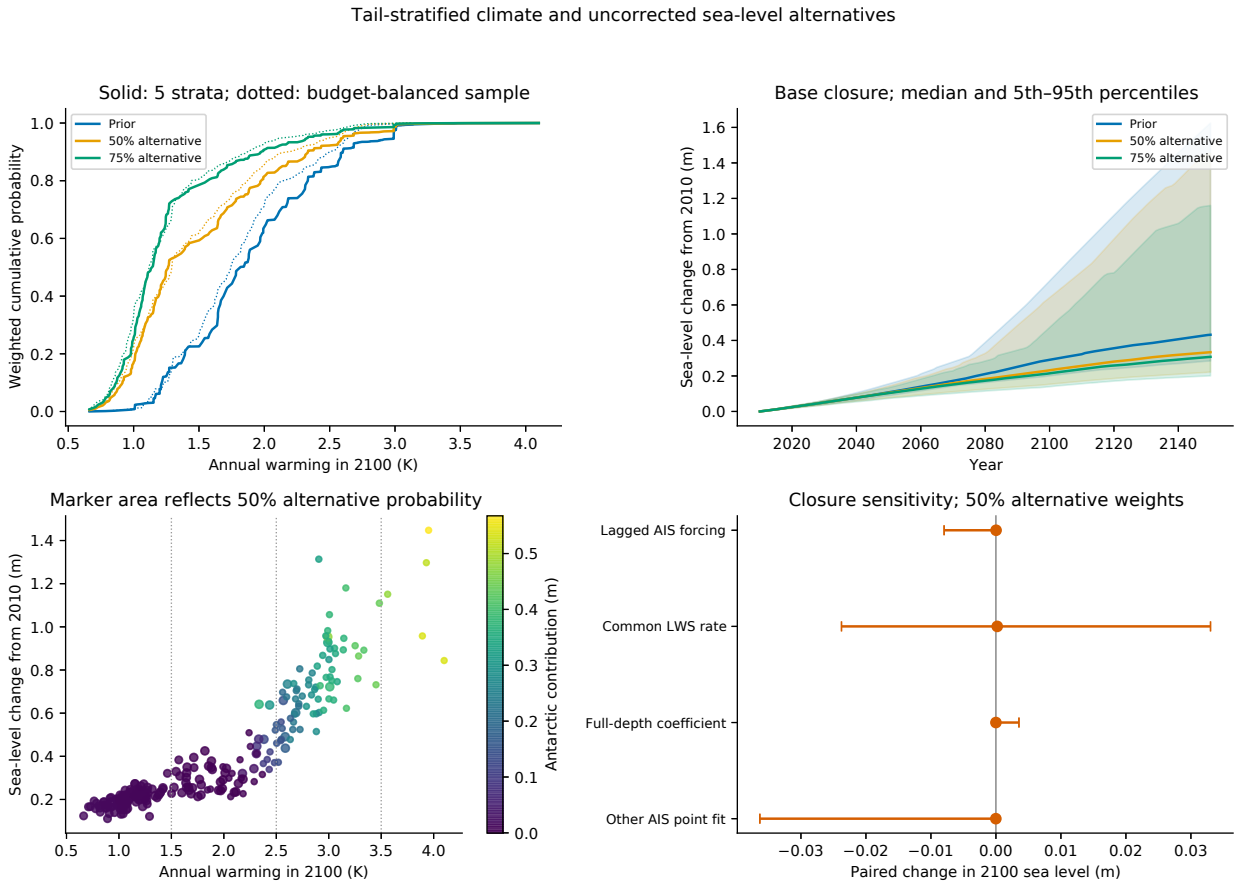


Figure 28: Tail sampling and coupled sea-level alternatives. Upper left: weighted warming distributions for the five-stratum experiment (solid) and the earlier balanced experiment (dotted). Upper right: base sea-level medians and 5th–95th percentiles. Lower left: member-level warming and sea level; color denotes the Antarctic contribution and marker size reflects alternative probability weight. Lower right: paired 2100 closure differences, using 50% budget weights. These are independent sensitivity experiments, not recovered published distributions.

The two Antarctic fits satisfy almost identical paleo constraints but differ by as much as 0.363 m in their future contribution. The fitted fast-loss threshold and rate therefore remain consequential beyond the calibration windows. Changing the storage covariance shifts individual 2100 totals by up to 0.050 m; the full-depth thermal alternative changes them by up to 0.025 m; lagging Antarctic forcing changes them by up to 0.011 m. These are sensitivities to named choices, not bounds on total scientific uncertainty.

12.3 Reference-period recovery and AR6 context

We recover the model's own 1995–2014 reference by combining its historical component calculations with the saved 2010–2014 forward segment. Each member's twenty annual values are averaged before subtracting that mean. The datum shift is 0.0071–0.0183 m and is not fitted to an external projection. On this datum, the 50% budget alternative has a 2100 median of 0.244 m and a 17th–83rd range of 0.194–0.363 m. The corresponding prior values are 0.304 m and 0.242–0.491 m.

AR6 Table 9.9 reports a 2100 median and likely range of 0.56 (0.44–0.76) m for SSP2-4.5[32]. The independent alternatives remain lower after matching the reference period. However, their emissions distributions are different; this is a contextual comparison, not a matched-scenario validation. Conditional ice fits, the low inferred climate sensitivity and missing joint uncertainty can all affect this comparison. We do not adjust any component to match AR6.

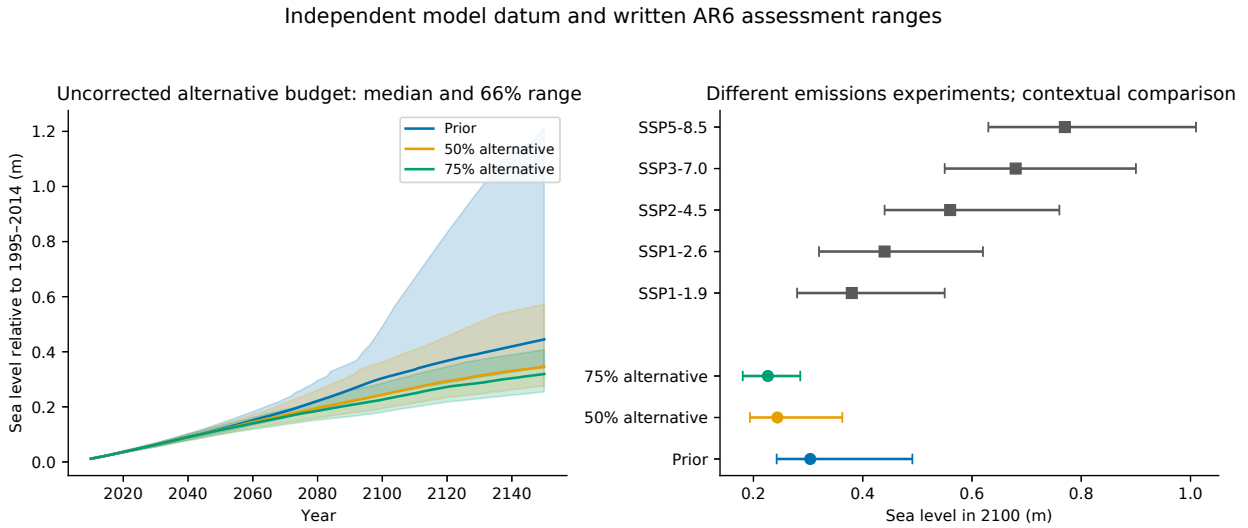


Figure 29: Datum-consistent sea-level context for supplementary Figure 4. Left: independent alternative medians and 17th–83rd percentiles relative to each member's 1995–2014 mean. Right: 2100 alternative ranges and printed AR6 Table 9.9 SSP ranges on the same datum. The SSP scenarios and probabilistic emissions alternatives are distinct experiments. The written AR6 summary values are comparison evidence only and are not inputs to the model.

12.4 Consistency of the historical component sum

The historical sea-level budget supplies a further constraint on interpretation (Figure30). We sum the independently fitted components for 1990.5–2009.5 and compare them with the recovered Church–White observations [29]. Each model series and the observations are centered on their own 1995.5–2004.5 mean. Annual states are linearly interpolated to year midpoints. Historical land-water storage is assigned the same 0.3 mm/year trend used before 2021 in the coupled experiment. We retain the 250 saved physical combinations and emissions-prior design weights for each climate calibration, testing heat integrated to 2000 m and to the full 4057.5 m column.

The observed linear trend is 2.602 mm/year. With the 2000 m heat response, HadCRUT4 and HadCRUT5 joint-calibration alternatives have median trends of 2.048 and 2.027 mm/year; their 5th–95th physical-parameter ranges are 1.635–2.460 and 1.653–2.415 mm/year. The temperature-only median is 3.006 mm/year, with range 2.340–3.469. Median-trajectory RMSE values are 4.65, 4.74

and 4.11 mm, respectively. Full-depth heat changes the no-heat median trend to 3.074 mm/year but changes the two heat-constrained medians by less than 0.007 mm/year. Refining 80 to 160 ocean cells changes the tested centered thermal trajectories by at most 0.105 mm; Antarctic 10/5-year step refinement differs by 3.10×10^{-12} m. The observed budget mismatch therefore requires physical, observational or statistical investigation beyond those numerical refinements.

This is a consistency diagnostic, not an independent validation. Church–White sea level already enters the Antarctic boundary forcing, known temperatures force the ice hindcasts, and the empirical thermal coefficient uses observations through 2015. The intervals exclude stochastic discrepancy and observation error, retain peripheral-ice overlap and omit the missing joint dependence. No rejection-sampling weights are manufactured from this comparison. Its closer temperature-only sea-level RMSE does not outweigh that calibration's failed temperature and heat predictions.

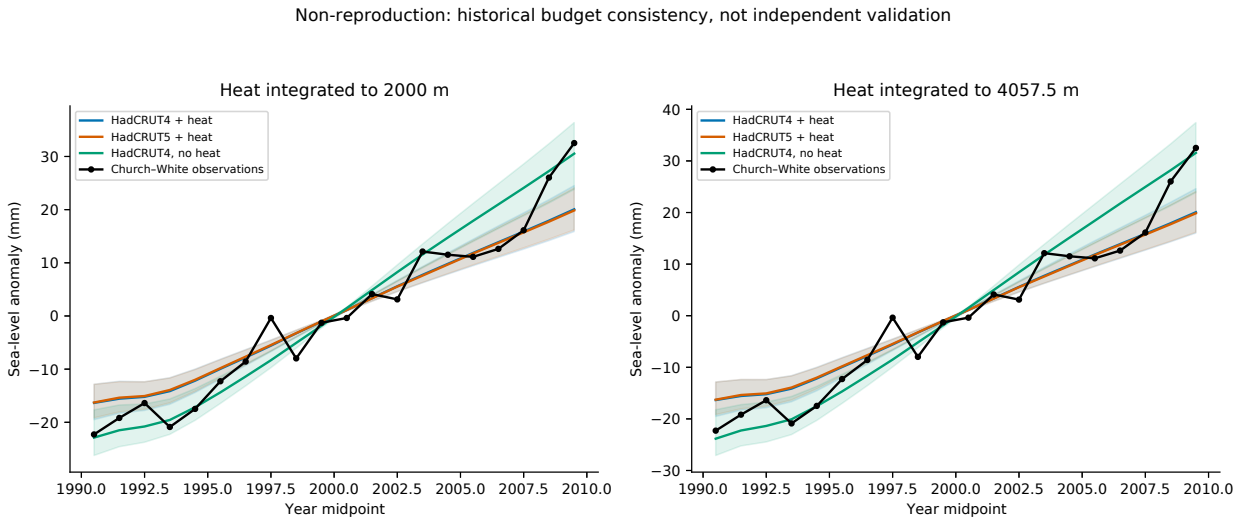


Figure 30: Historical consistency of the uncorrected sea-level component sum. Black points show Church–White observations; colored lines and bands show medians and 5th–95th physical-parameter percentiles under saved prior design weights. Model and observations are centered separately over 1995.5–2004.5. Panels differ only in the heat depth used for thermal expansion. Observation and stochastic model errors are not included in the bands. The observational record also supplies an Antarctic boundary condition, so agreement cannot be counted as independent validation or recovery of the target joint calibration.

13 Coupled geographic cost alternatives

13.1 A probability sample with an independently known control

Each alternative climate member now drives sea-level components and a coastal cost calculation. We retain the original 200 elevation bins, annual cash flows, SSP4 OECD economics and the paired no-rise baseline. Sea-level change is uniform across coasts because the target fingerprint and land-motion inputs have not been recovered. Economic response parameters remain fixed at their declared central values. Thus this experiment propagates the sampled climate and sea-level variation but not the full target scientific or socioeconomic uncertainty.

The complete 9,531-coast calculation under a prescribed 1 m path supplies an auxiliary cost x_i and exact regional sum X_r. We sample coastal units independently within seven World Bank

regions and an unmapped stratum, assigning

$$p_i = 0.9\frac{|x_i|}{\sum_{k\in r}|x_k|} + \frac{0.1}{N_r}. \tag{16}$$

Here N_r is the number of eligible coasts in the stratum. All coasts have positive inclusion probability. For a new path's cost y_i, we estimate

$$\widehat{Y}_r = \beta X_r + \frac{1}{n_r}\sum_{j=1}^{n_r}\frac{y_{I_j} - \beta x_{I_j}}{p_{I_j}}, \qquad \beta = \frac{S(2100) - S(2010)}{1\,\mathrm{m}}. \tag{17}$$

The sampled indices I_j are independent draws with replacement. Taking the expectation cancels the control term and gives $\sum_i y_i$ for any fixed β. This identity does not assume that coastal costs are linear in sea level. The sample variance of the summands, divided by n_r, estimates the conditional numerical variance. Variances add across independent strata; using the same coasts for every climate member induces covariance between their numerical errors.

Exact enumeration tests verify the estimator and its variance for a finite three-coast population. A further 5,000 design replicates per test compare estimates with complete independent 0.5, 1 and 2 m calculations. Biases lie within 1.84 Monte Carlo standard errors, and empirical-to-estimated variance ratios span 0.966–1.027. Normal 95% intervals cover only 93.3–94.0% at sixteen draws per stratum. The variance check therefore supports the numerical error estimate without establishing exact normal-interval coverage.

A six-path pilot covers the minimum and maximum sea-level paths and each occupied warming category. Doubled flood quadrature changes selected cost by less than 2×10^{-16} relatively and preserves all tested policies. The full calculation uses sixteen draws per stratum, 119 distinct coasts and 250 climate members: 29,750 climate–coast combinations. It takes 3,016 wall seconds, including controlled pauses for serial validation. Median conditional global cost SE is 6.26% of estimated cost; the maximum is 9.22%.

13.2 Global distributions and comparison limits

Figure 31 shows the warming-conditioned cost estimates. Under 50% budget weights, the four median incremental costs are 0.208, 0.290, 0.664 and 1.217 trillion 2019 USD PPP. The corresponding 5th–95th ranges are 0.137–0.260, 0.214–0.440, 0.385–0.965 and 0.808–1.609 trillion. The highest category contains only five members, with conditional ESS 4.23; its plotted percentiles identify order statistics, not converged tail estimates.

The target reports medians of 1.2, 1.5, 2.1 and 2.6 trillion 2010 USD over a 2150 horizon. Our estimates are numerically smaller, but the monetary inputs, reference-cost subtraction and time horizon differ. Applying the independently documented CPI factor 0.852924 lowers the alternative values further; that price-year conversion does not reconcile international PPP, exposure or accounting. We do not report the remaining difference as a commensurate model error. The target's threshold probabilities also differ substantially from our alternative weights (published Table 1 and the figure inventory).

13.3 Regional burdens and planned movement

The regional ranking differs from the target's written interpretation. For category 3, East Asia and the Pacific has a median burden 0.591 percentage points above the covered global cost/GDP ratio. Latin America and the Caribbean, South Asia and Sub-Saharan Africa instead lie 0.218,

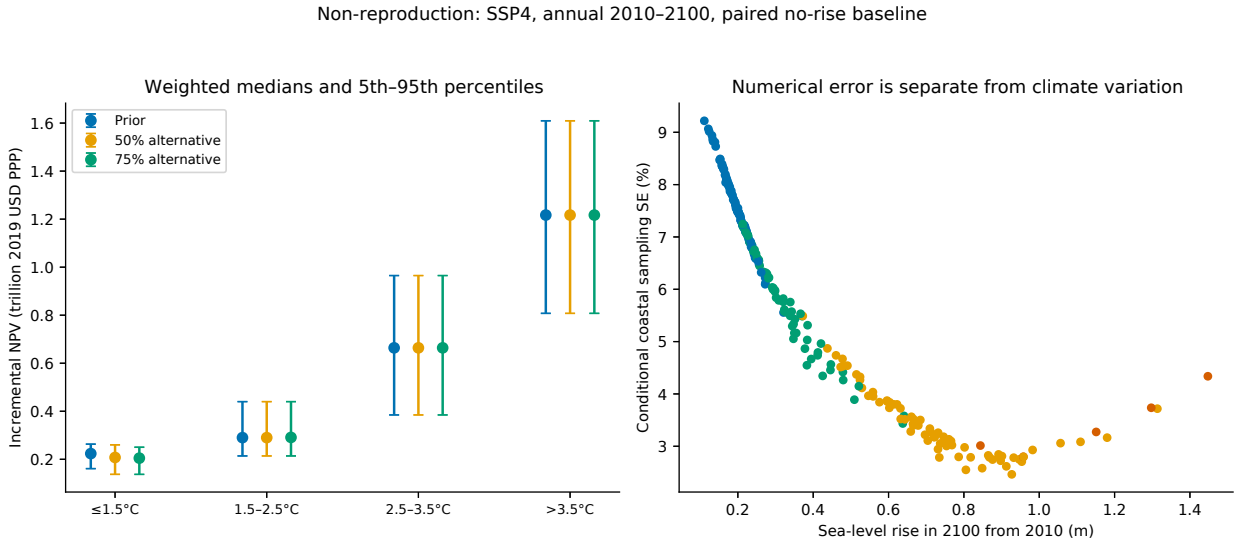


Figure 31: Global cost alternatives corresponding to the quantities in main Figure 1. Left: weighted medians and 5th–95th order statistics under three emissions probability measures. Costs use annual 2010–2100 accounting, 4% discounting to 2010, 2019 USD PPP and a paired no-rise counterfactual. Right: conditional coast-sampling SE for each climate member, separate from the variation shown on the left. Errors are correlated across members because coast draws are shared. Only five members occupy the highest warming category; its tail percentiles remain unresolved. These are non-reproductions.

0.132 and 0.187 percentage points below it. The target identifies disproportionate burdens across several Global South regions. Agreement in the direction of increasing global costs with warming therefore does not establish agreement in geographic inequality. The denominator is independently acquired 2010 regional GDP in 2019 PPP units; unmapped costs are excluded consistently from the covered global comparator.

Under the fixed Antarctic point fit, 94 of 250 paths trigger rapid ice loss before 2100, carrying 13.99% of the alternative weight. Estimated median planned movement is 79.93 million without triggering and 104.10 million with triggering. Median conditional coast-sampling SE is 12.50 and 13.60 million, respectively. The median paired difference from no-rise planned movement is 8.31 versus 32.48 million. These are movement counts under optimal modeled policies, including initialization; they are neither counts of unique persons nor forecasts of actual migration. Conditional association is not a causal estimate of the Antarctic trigger, since triggering covaries with warming.

The alternative weighted 95th-percentile cost threshold is 0.597 trillion 2019 PPP. Its retained upper-tail members include one category-2 path, 48 category-3 paths and all five category-4 paths. No retained path exceeds the target's literal 2.7 trillion threshold even after the CPI-only conversion. Neither observation establishes a zero population probability: the target threshold uses different units and accounting, and the high-warming tail is poorly resolved.

13.4 Cost-baseline sensitivity

The written CIAM lineage subtracts initial-interval calibration costs, whereas the alternative above subtracts an independently optimized no-rise trajectory. These operations need not agree. A six-path pilot holds the annual policies fixed and separately retains gross costs, removes the 2010

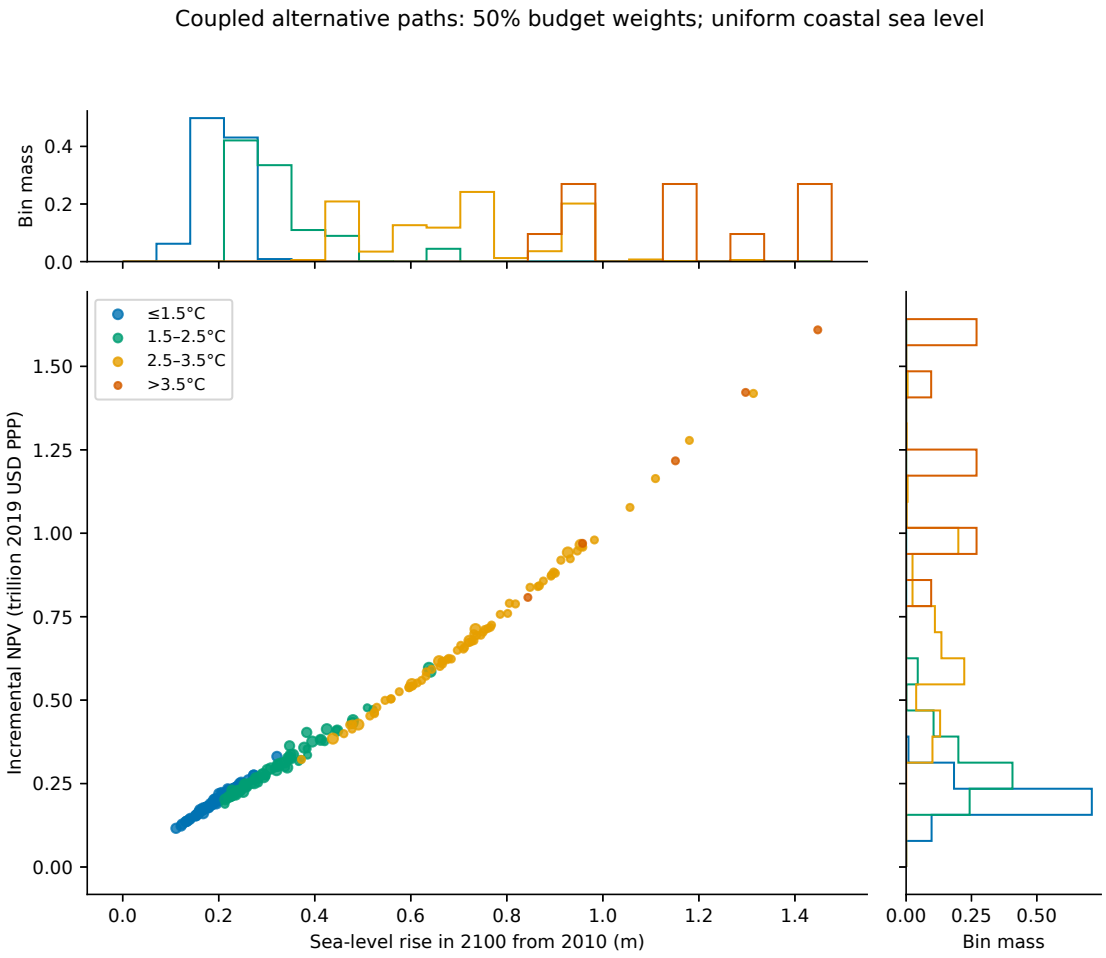


Figure 32: Coupled sea level and coastal cost with empirical conditional marginals, corresponding to main Figure 2. Symbol area reflects 50% budget weights and color identifies annual 2100 warming. Marginal histograms normalize within each warming category and do not invent a smooth tail density. Sea level is relative to 2010; its separate 1995–2014 transformation is documented in Figure 29. Coastal units, accounting and uncertainty differ from the target, so the plotted relationship is an independent alternative.

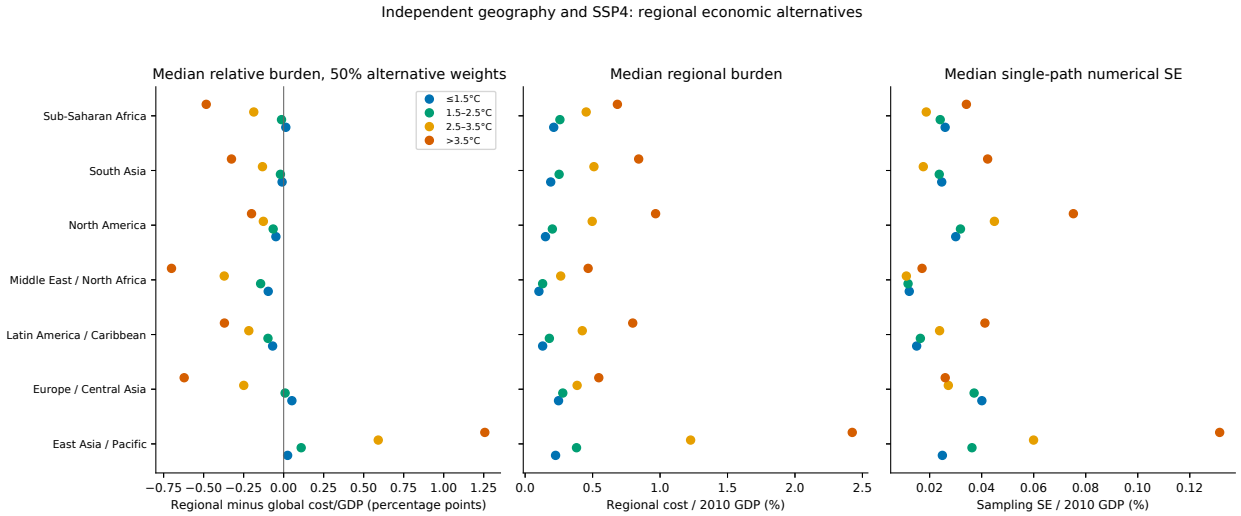


Figure 33: Regional alternative for main Figure 3. Left: the median of each member's regional-minus-covered-global cost/GDP ratio. Middle: median regional incremental cost divided by 2010 GDP. Right: median conditional single-path coastal sampling SE on the same GDP scale. The last quantity is not the SE of the median. Symbols distinguish warming categories; the highest category has only five members. Independent geography and economic accounting produce a different regional pattern from the target's written findings.

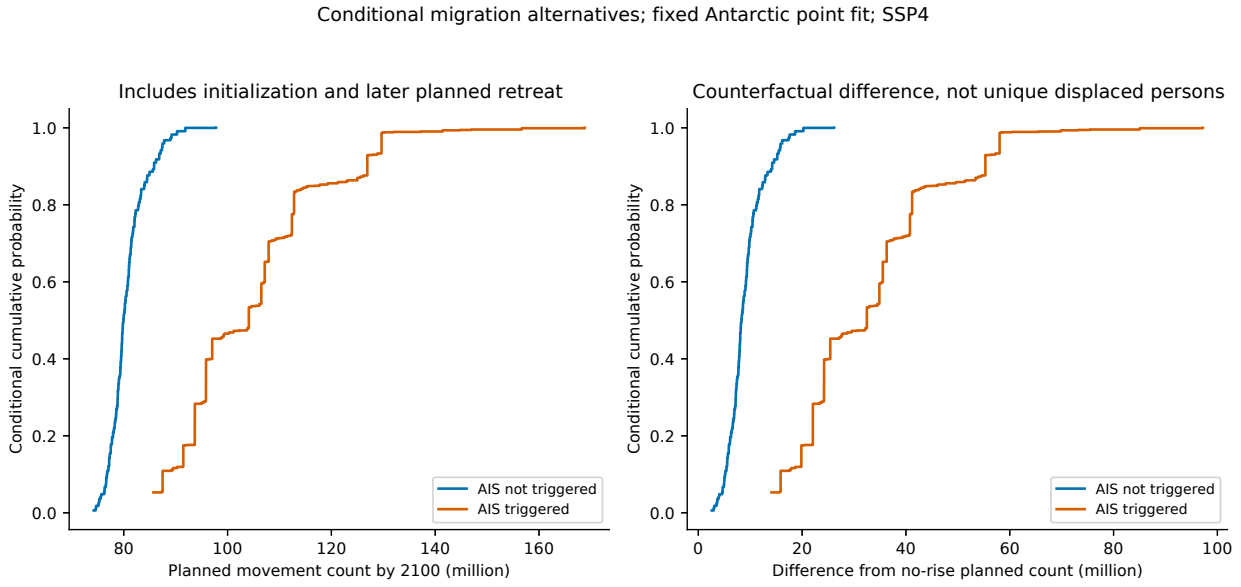


Figure 34: Planned movement alternatives for supplementary Figure 10. The empirical weighted distributions condition on whether the fixed Antarctic model triggers by 2100. Left: all planned relocation, including initial adaptation. Right: the paired difference from the no-rise planned count, which can be negative when the optimal policy changes. These estimated distributions omit target socioeconomic parameter uncertainty and are subject to the separately reported coastal sampling error.

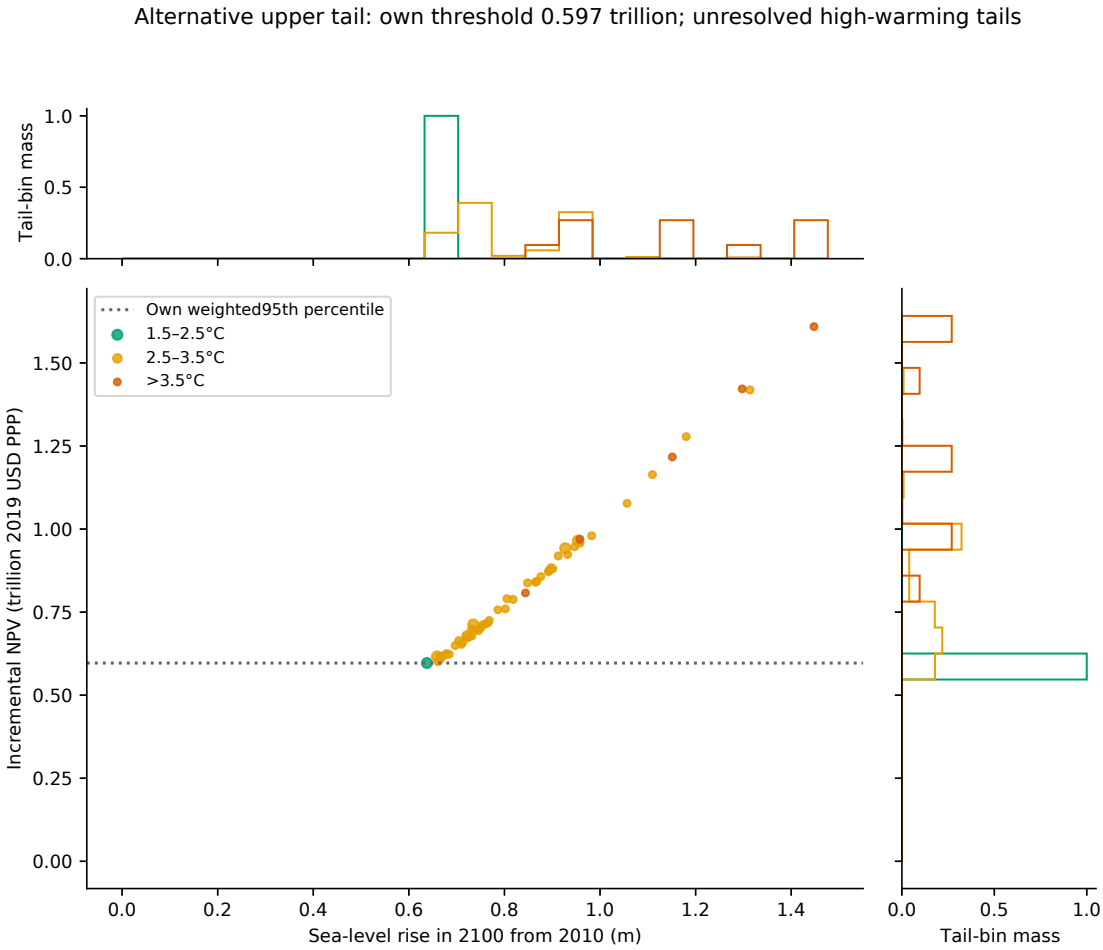


Figure 35: Upper-tail alternative corresponding to supplementary Figure 7, using the independent distribution's own weighted 95th-percentile threshold. Empirical marginals normalize within each retained tail category. The plot uses the same sea-level datum and cost conventions as Figure 32; it does not replace the target's 2.7 trillion 2010 USD selection. Tail category sizes and coastal sampling errors limit its interpretation.

payment, removes all 2010–2019 payments, or subtracts the full no-rise trajectory. The complete 9,531-coast no-rise present value is 1.958 trillion 2019 USD PPP. Adding this known constant to the four baseline-subtracted medians gives 2.166, 2.249, 2.622 and 3.175 trillion. Thus the accounting convention can reverse the sign of an apparent cost-level difference; gross costs are not an alternative claim of agreement with the publication.

We also reject a literal annual interpretation that repeatedly subtracts the 2010 payment. That payment contains one-off investments; treating it as a recurring reference rate produces values as low as −44.23 trillion. The diagnostic exposes an invalid translation between annual investments and decadal reference costs. It is not retained as a physical cost model.

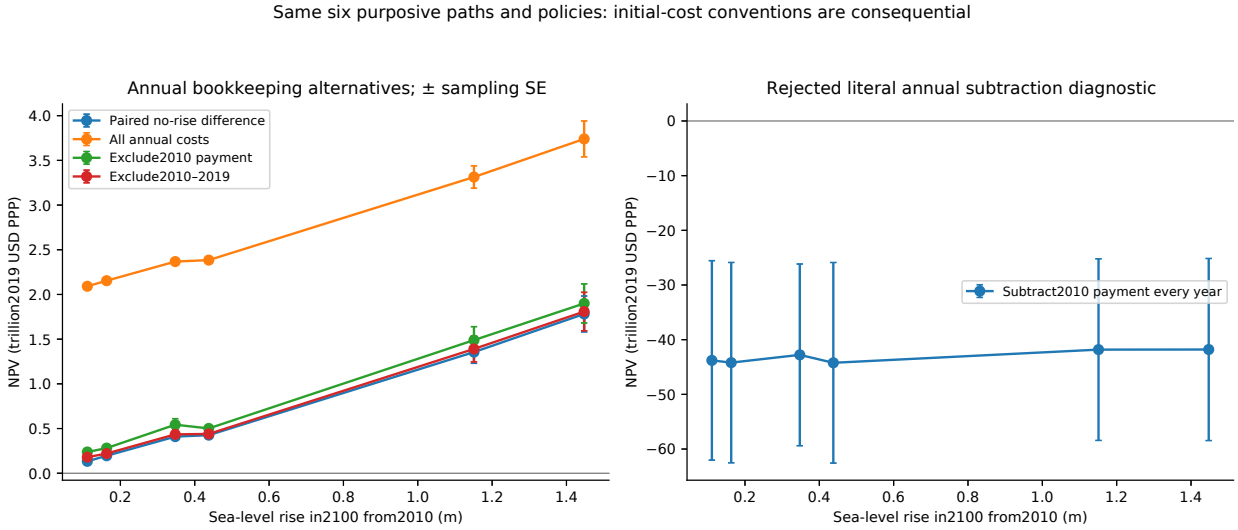


Figure 36: Accounting alternatives for six purposively selected paths and four coastal draws per region. Left: different removals of initial or no-rise costs, with conditional sampling SE. Gross totals use the known complete no-rise sum as an additional control. Right: a rejected repeated subtraction of a one-off initial payment. None of these conventions is asserted to recover the target decadal spin-up accounting.

13.5 Finite-ensemble and geographic sampling tests

Supplementary Figures 1 and 2 ask whether the cost distributions depend on the number of retained climate states. We test that question within the alternative ensemble without claiming the target's 5,000 versus 4,000 comparison or its 202 members per warming category. Five hundred independent subsamples retain 200 of the 250 states, preserving the total weight of each emissions stratum. The category-4 median difference has a 5th–95th range of −0.247 to +0.205 trillion 2019 PPP. The category-2 upper quantile also changes substantially: the median difference in its 95th percentile is +0.145 trillion. Such behavior is incompatible with a claim that all tails are converged.

A second test selects five distinct states per warming category. Its apparently zero category-4 variability merely retains all five available states. The category-3 median difference ranges from −0.279 to +0.278 trillion across subsamples. Separately, nested coastal probability samples increase from four to eight to sixteen draws per region. Increasing eight to sixteen changes the four cost medians by −0.0050, −0.0048, +0.0052 and −0.0400 trillion. These are dependent numerical samples, not independent posterior replications. Neither this comparison nor a reduced conditional sampling SE establishes convergence of the climate distribution or its rare-event probabilities.

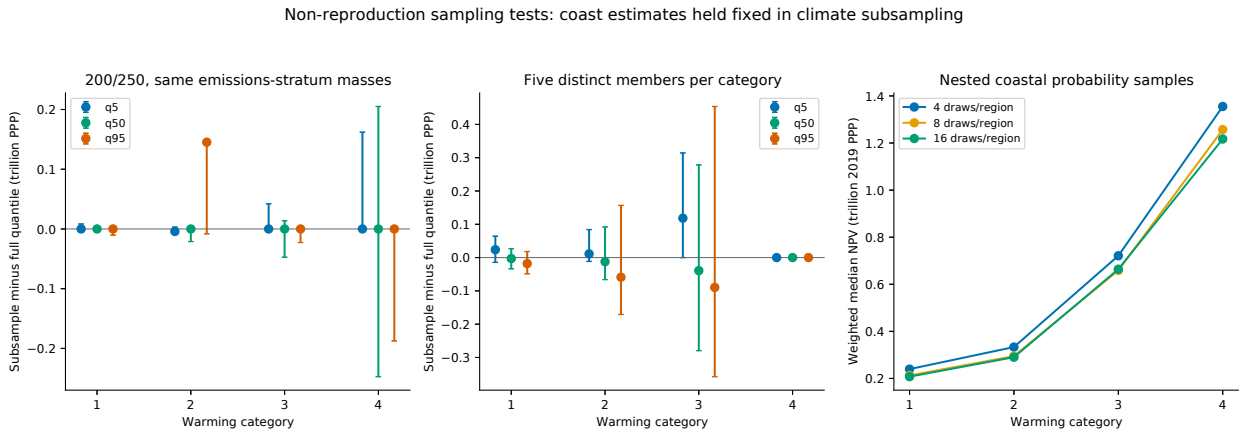


Figure 37: Alternatives addressing supplementary Figures 1 and 2. Left and middle: median changes across 500 subsamples, with their 5th–95th ranges. Coastal cost estimates are held fixed when subsampling climate states. Right: conditional cost medians under nested coastal samples. The zero fourth-category range in the equal-size test is a finite-population identity, not evidence of convergence.

13.6 Socioeconomic scenario, horizon and income contrasts

A matched SSP2 calculation uses the same 250 climate–sea-level states and the same eight coastal draws per region as a separately aggregated SSP4 calculation. It evaluates 14,750 state–coast combinations after twelve pilot quadrature and policy checks pass. The comparison isolates the scenario change within the alternative geography, monetary units and annual accounting. East Asia and the Pacific remains above the global cost/GDP burden in the high-warming categories; the broader target Global South pattern is still not reproduced.

Six purposively selected paths test assumptions separately. SSP2 increases their estimated global incremental costs by 0.66–1.78% relative to SSP4. Extending the horizon through 2150 increases costs by 4.40–33.18% when post-2100 population and income are frozen. Continuing their final growth rates gives 4.38–33.01%. These are two transparent extrapolations, not recovered post-2100 socioeconomic specifications. Geometric interpolation between the provider's nodes changes these six costs by less than 0.10%. The results show that horizon and initial-cost accounting deserve more attention than within-node interpolation in these cases; they do not rank sensitivities for the population of possible states.

For supplementary Figure 9, the explicit contrast is SSP4 minus SSP2 of the change between category-3 and category-2 median regional cost/GDP. The income contrast is SSP4 minus SSP2 regional 2100 GDP per person. Regional populations combine independent 2010 observations with provider scenario growth factors; the least complete region covers 80.5% of its reference population. Across seven regions, the descriptive slope is 0.00219 percentage points per thousand 2019 PPP dollars, with $R^2 = 0.664$. The caption's precise target contrast, coverage, cost inputs and monetary convention remain unmatched. Correlated coastal sampling errors and seven aggregated observations preclude treating this regression as a causal result or a validated target coefficient.

13.7 Conditional economic-parameter uncertainty

The six recovered truncated-normal socioeconomic marginals are now sampled independently, using sixteen draws plus the central parameter vector. Each parameter draw reoptimizes adaptation and

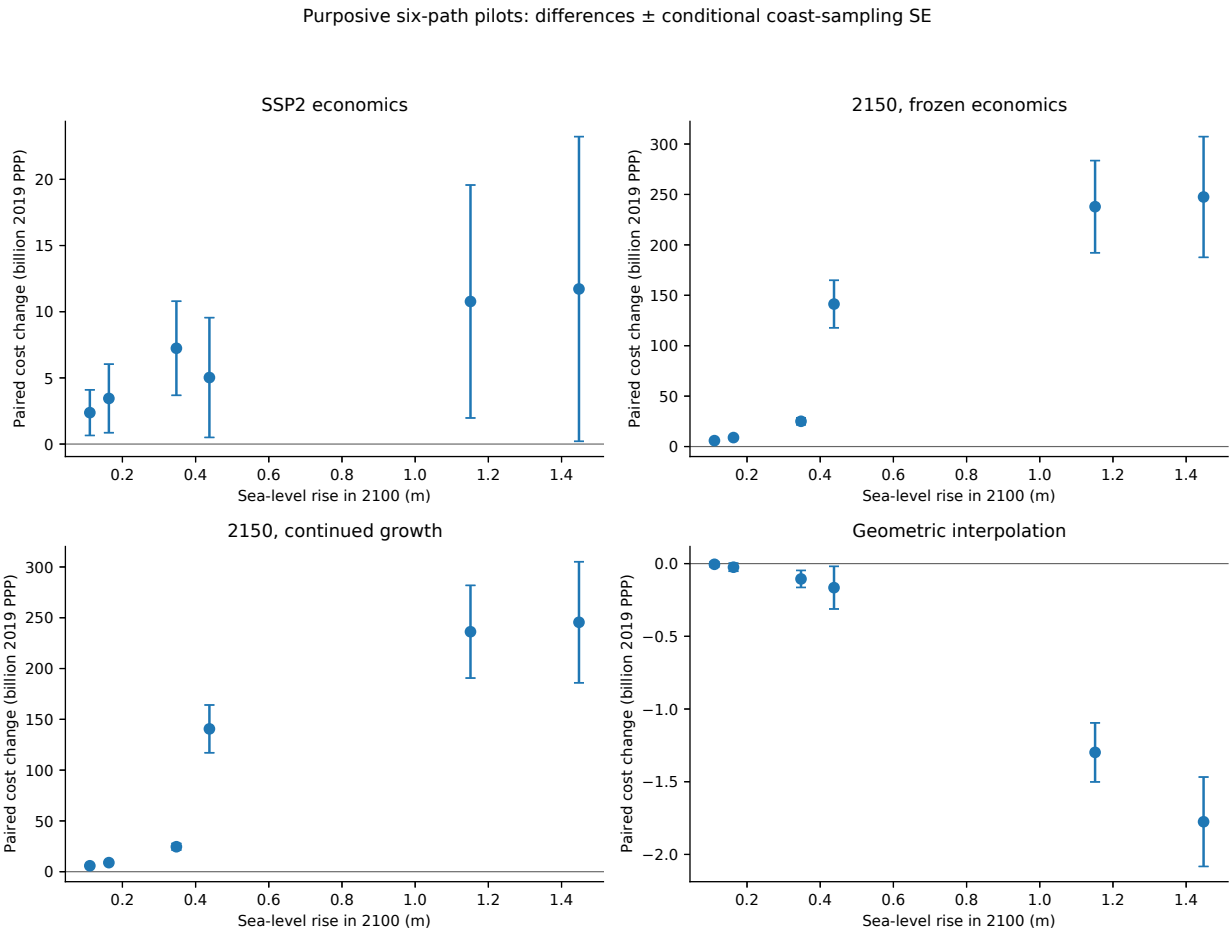


Figure 38: Paired six-path sensitivities. Error bars give conditional coastal sampling SE of each difference. The horizon extensions retain the same first-period policy rule and independently optimize each no-rise baseline. They are explicit alternatives to unresolved target conventions.

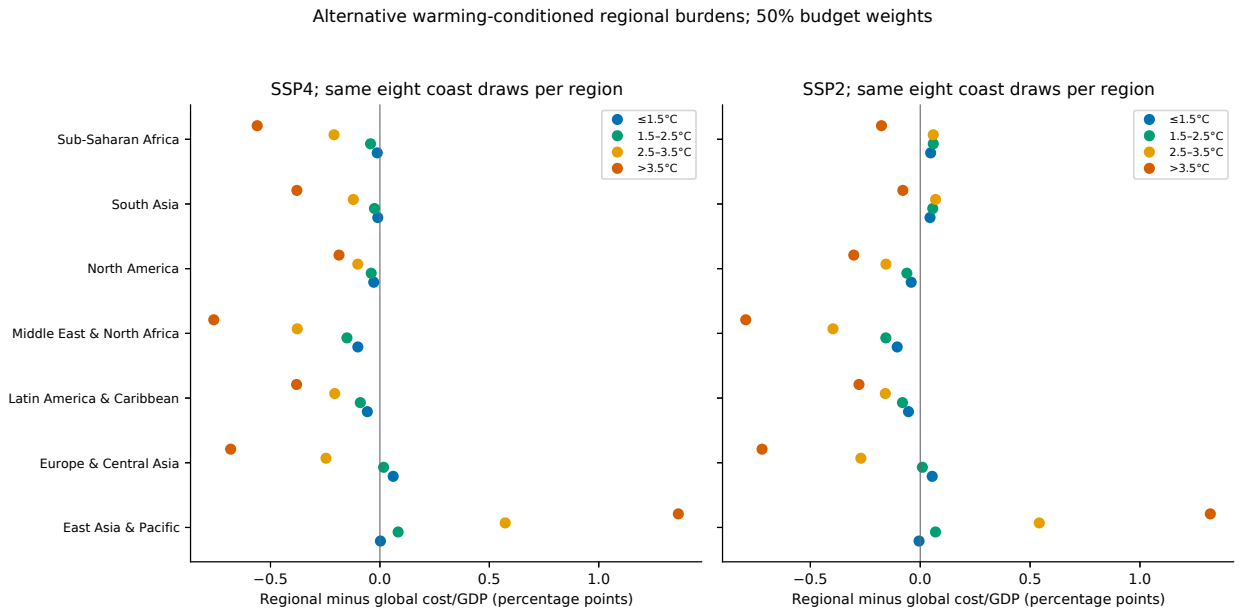


Figure 39: Supplementary Figure 8 alternative. Both scenarios use the same 250 climate states, eight coastal draws per region and alternative 50% budget weights. Points show the median regional-minus-global cost/GDP burden. The panels compare socioeconomic inputs within the reconstruction; neither is a published-result reproduction.

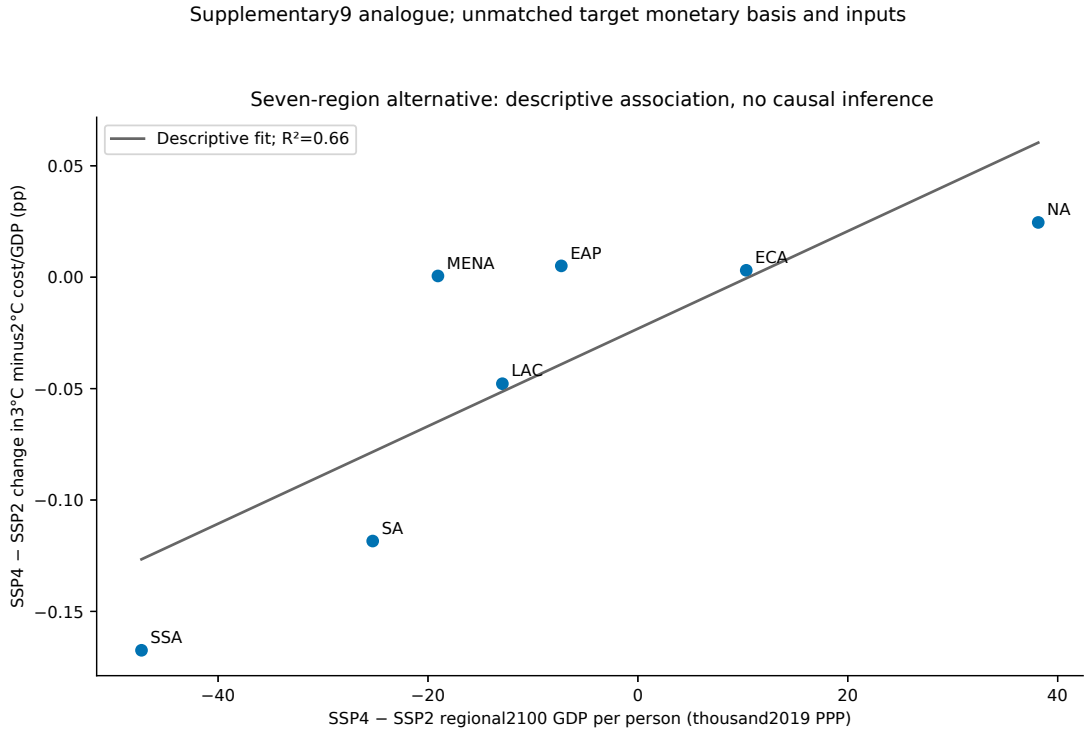


Figure 40: Explicit interpretation of the supplementary Figure 9 comparison. The line is a descriptive least-squares fit to seven alternative regional contrasts. No significance test or causal interpretation is assigned.

its own no-rise baseline across six purposive climate paths and four common coastal draws per region. Migration and statistical-life parameters enter directly. Land values are scaled relative to the documented central coefficient; wetland values receive relative income and population elasticity changes after 2010. These latter mappings and independent parameter correlations are explicit assumptions required by the different input valuation, not recovered target conventions.

The calculation completes 3,162 climate–coast–economic cases. All 68 additional quadrature/policy comparisons pass. The central vector reproduces the earlier pilot's coastal costs within 2.2×10^{-15} relative error. Across the six paths, the sampled cost differences range from −35.2 to +85.5% relative to central parameters. Median paired coast-sampling SE ranges from 4.65 to 9.40% of each path's central cost. The largest sampled total is 3.304 trillion 2019 PPP, on the purposive highest-sea-level path. This is not a calibrated tail probability or a target-threshold exceedance result. Six chosen paths and sixteen independent marginal draws cannot represent the full joint probabilistic experiment, but they show that holding these economic parameters fixed can exclude substantial conditional variation.

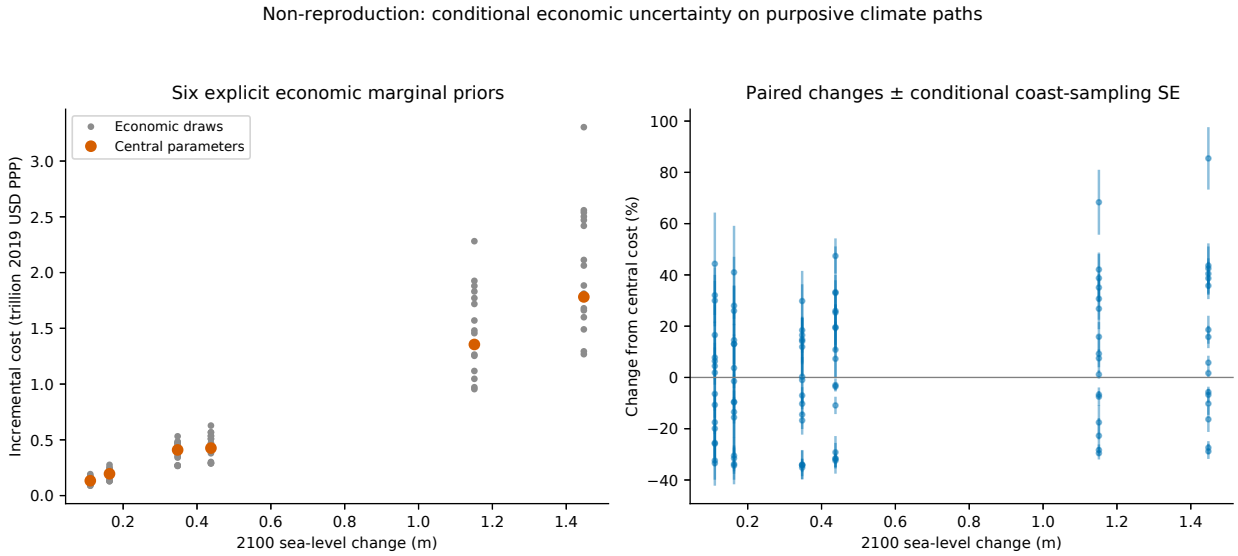


Figure 41: Economic-parameter sensitivity under recovered marginal prior bounds and explicit independent sampling. Left: estimated costs for sixteen economic draws, with central-parameter totals marked separately. Right: paired changes and conditional coastal sampling SE. Sample extrema are not credible limits; climate paths are purposively selected and the initial-value mappings differ from the target model.

13.8 Coastal propagation of the temperature-version change

A separate HadCRUT5 coastal pilot passes the same policy and doubled-quadrature checks before 250 paths are propagated through 59 distinct coast units, representing eight draws per region. The full run contains 14,750 cases and takes 1508 s, including controlled pauses for serial verification. We compare it with the HadCRUT4 aggregate using the same eight coastal draws, retaining the original sixteen-draw base results separately. The 250 regional/global cost and variance identities are checked after aggregation.

Figure42 compares these matched geographic designs. Under 50% budget-event weights, the unconditional cost median rises from 0.235 to 0.241 trillion 2019 USD PPP, and the 95th percentile

rises from 0.591 to 0.677 trillion. Conditional medians are 0.216, 0.308, 0.599 and 1.117 trillion for HadCRUT5, compared with 0.213, 0.295, 0.659 and 1.257 for the matched HadCRUT4 design. The two upper conditional medians decrease because the independently sampled posterior changes category composition; this is not evidence that warming reduces a fixed coast's losses. HadCRUT5 conditional 5th–95th ranges are 0.149–0.285, 0.197–0.567, 0.407–1.142 and 0.800–1.591 trillion. These bars omit the separate coastal sampling error, whose median relative value is 7.39% and maximum is 13.07%.

The highest category grows from five states, with conditional effective size 4.23, to 22 states, with effective size 13.61. Its weighted probability rises from 0.0534% to 0.393%. This is improved coverage, not converged tail resolution. All 22 high-category states and 35 third-category states lie at or above the HadCRUT5 ensemble's own weighted 95th cost percentile; none from the lower two categories does. This threshold is 0.677 trillion 2019 PPP, not the target's 2.7 trillion 2010 USD threshold. No computed case exceeds the latter after a CPI-only rescaling, which remains an incommensurate diagnostic and cannot imply a zero population probability.

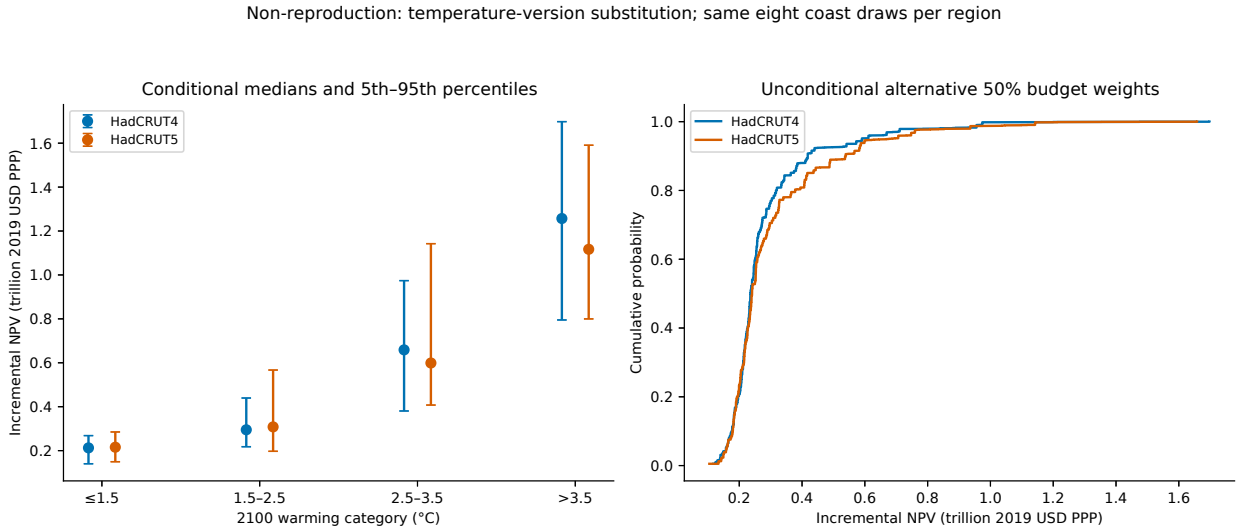


Figure 42: Temperature-version substitution using matched coastal probability samples. Left: warming-conditioned cost medians and 5th–95th percentiles. Right: unconditional weighted cumulative distributions. Both use eight draws per region, 200-bin profiles, SSP4 and annual 2010–2100 accounting with a paired no-rise control. Physical posterior states are independently sampled; category composition changes. Coast-sampling errors are separate from the shown physical-parameter variation. Neither distribution is a target reproduction.

The new joint plot uses each member's modeled 1995–2014 mean sea level (Figure43). The tight cost–sea-level relation reflects uniform geographic sea-level forcing and fixed economic parameters. It omits the published joint uncertainty rather than recovering a sharply determined physical relation. Separate economic-prior experiments quantify part of the omitted variation; their purposive paths cannot simply be merged into this probability ensemble.

The regional ranking remains discrepant. In category three, the median regional minus global cost/GDP contrast is +0.498 percentage points in East Asia and the Pacific, but −0.187 in Latin America and the Caribbean, −0.111 in South Asia and −0.179 in sub-Saharan Africa. Changing temperature version therefore does not reproduce the target's broad regional inequality pattern (Figure44). Planned-movement medians are 72.31 million without Antarctic triggering and 98.39

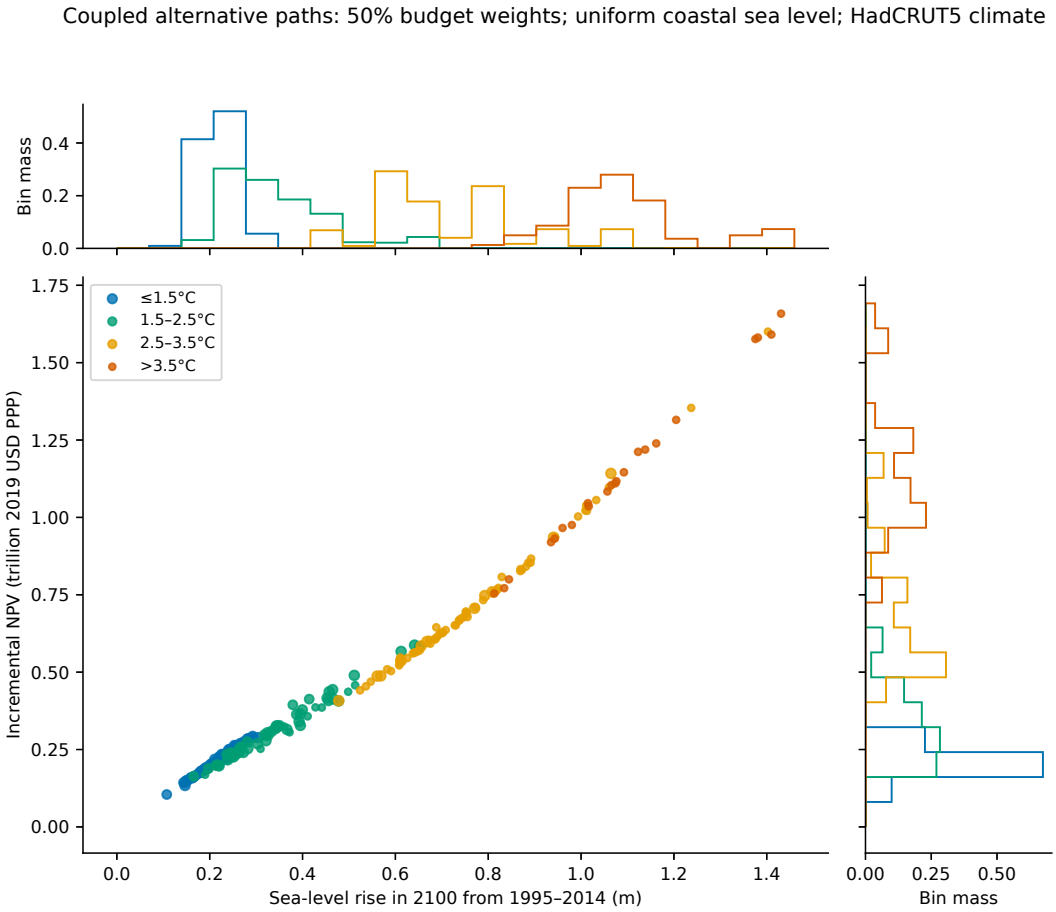


Figure 43: HadCRUT5 coupled sea-level and cost alternative with conditional empirical marginals. Sea level uses each member's own 1995–2014 reference. Symbol area reflects 50% budget-event weight; colors indicate annual 2100 warming categories. Marginals are normalized within each category. Fixed valuation and uniform geographic sea-level forcing limit the modeled dispersion. The monetary basis and horizon remain different from target Figure 2.

million with it; paired no-rise increments are 8.89 and 34.98 million. Their median conditional coast-sampling standard errors are 12.48 and 14.75 million for total movement. These counts include initialization and successive planned retreats, not unique lifetime migrants. The trigger groups differ in climate as well as ice behavior, so the contrast is not a causal estimate of migration attributable to Antarctic loss.

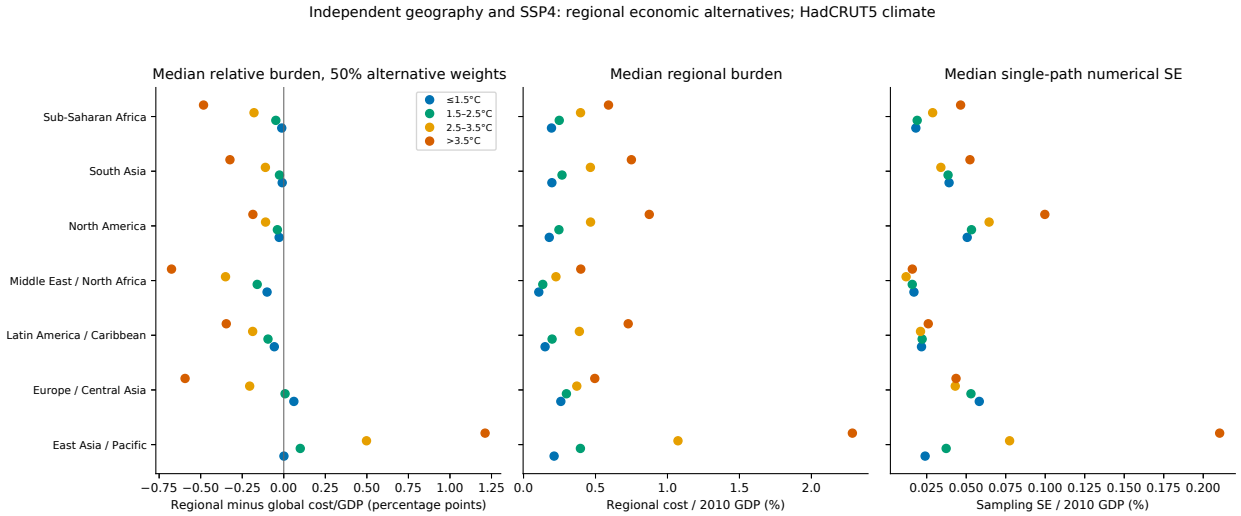


Figure 44: HadCRUT5 regional cost alternatives. Left: median regional minus global cost/GDP by warming category. Center: regional burden relative to 2010 GDP. Right: median conditional numerical standard error expressed in the same GDP units. Independent exposure and valuation inputs, uniform sea-level change and annual accounting remain unmatched to the target. The geographic ranking is not repaired by substituting the documented temperature version.

13.9 Sampling sensitivity with increased high-warming coverage

The HadCRUT5 ensemble permits a second finite-sample test with 22 high-category states, instead of five. Five hundred 200-of-250 subsamples preserve each emissions stratum's probability mass. They retain 12–22 high-category states. Across these subsamples, the 5th–95th range of high-category median changes is −0.071 to +0.029 trillion 2019 PPP; the corresponding change in the third category's upper cost quantile is −0.206 to zero. Equal samples of 16 distinct states per warming category produce a high-category upper-quantile change from −0.276 to +0.068 trillion. In category three, the median upper-quantile change under equal sampling is −0.275 trillion.

Figure45 describes sensitivity within the saved states. Unlike the base equal-five test, equal-sixteen sampling now varies the retained high-category membership. It still cannot sample unobserved climate states or missing physical uncertainty. Coastal estimates remain fixed across replicates; their errors are not included in these bars. Neither a zero median difference nor a narrow range for one statistic proves convergence of the population tail.

14 Unresolved inputs and implications

The joint parameter distribution is a central missing scientific object. The cited methodology combines paleo Antarctic constraints, modern climate and component-sea-level observations, uncertain emissions, and autoregressive residuals. The later written supplement identifies several data vintages,

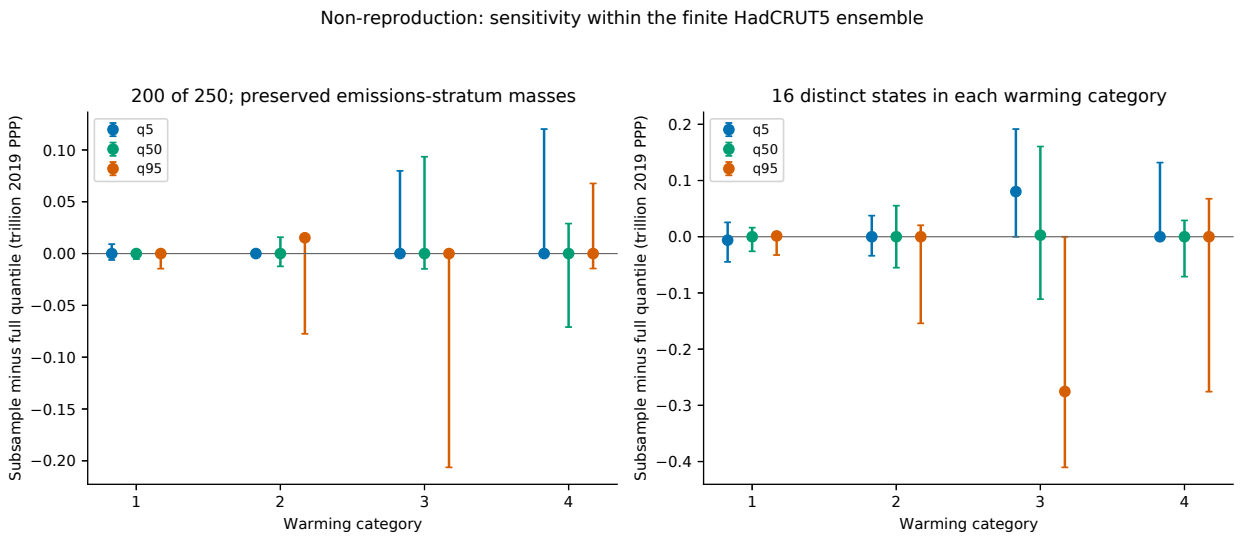


Figure 45: HadCRUT5 finite-ensemble subsampling alternatives to the questions in S1 and S2. Points and bars give medians and 5th–95th percentiles of changes in each cost quantile across 500 subsamples, relative to the saved full ensemble. Left: 200 states with preserved emissions-stratum masses. Right: sixteen states per warming category. The bars describe subsample variation, not population confidence or posterior credible intervals.

and the cited climate methods specify AR(1) residuals. The exact target combination of annual 3000 m heat data, cross-component measurement errors, initial states and joint calibration conventions has not been recovered. Reusing published posterior samples is prohibited here. Independent conditional calibrations now use versioned exogenous observations and explicitly alternative likelihood choices. Independent draws from older marginal quantiles would discard physically relevant dependence.

The historical DIVA exposure bundle is another independent obstacle. The CIAM supplement reports that its v1.5.5 database was no longer supported and that v1.8 was not public at that time [10]. The EEA provider directory currently reports unavailable software [13]. This does not prove that every DIVA dataset is inaccessible; it establishes that a complete permissible target bundle was not recovered in this investigation. The independent SLIIDERS segmentation is used only in labeled alternative experiments, not substituted as if it were the 12,148 target segments.

Independent GTSR data and six SSP forcing inputs have been acquired [14, 15]. The GTSR release supplies coastal features and a later release supplies Gumbel parameters, but invalid records, repeated station labels and return-level discrepancies require explicit treatment. The 2022 academic supplement was recovered in the scholarly preprint, resolving the six socioeconomic marginal priors and several decision conventions. Independent country economics and geographic exposure now support an alternative calculation. Original exposure, component fingerprints and the precise target input vintages remain unresolved.

Independent paleo records and the documented initial Antarctic radius have now been recovered. These are no longer categorical barriers to hindcasts. The remaining paleo uncertainty concerns the exact splice, chronology and transient ocean forcing, rather than absence of all boundary information. Emissions priors and the2022–2100 budget interval are also recovered; a unique expert density and its full joint calibration remain unspecified.

The independent pathways now drive a complete alternative sea-level budget and a sampled geographic cost calculation. Fresh temperature-vintage calibration and restricted carbon fits test

consequential input choices. Further extensions must address the remaining carbon/ice calibration and target dependence rather than merely increasing sample size. The regional alternatives do not establish quantitative agreement with target inequality.

15 Assessment of the research-agent process

The agent recovered much of the mathematical interface from distributed scholarly sources, distinguished versions of a mutable preprint, and implemented several independently checkable numerical components. The analytical tests separate solver accuracy from calibration and data availability. Cross-language flood integration supplies an additional check against transcription or implementation errors in the elevation primitive and GEV convention.

Two interpretation errors were avoided through source comparison. First, the older BRICK temperature-based expansion emulator is not the heat-based component described in the target lineage. Second, a later revision behind the same DOI cannot automatically be treated as the manuscript available before publication. Direct calculation then identified two tensions between printed equations and their prose descriptions. These checks do not reveal the original implementation.

Target identity remains limited by evidence recovery, while numerical auditing is also essential. An agent can write functioning code while failing to identify the scientific experiment if it supplies convenient values for missing inputs. This project does not count such code as a successful reproduction. The initial assessment stopped too early: further scholarly reading recovered major mechanisms and independent inputs. Observation-based calibration, carbon mass closure and DAIS integration are now demonstrated, while a global adaptation ensemble remains unreproduced. The failed holdout checks show why solver verification and posterior convergence cannot substitute for validation of the predicted observable.

Two further checks change the interpretation of apparent blockers. A renewed source search finds the cited glacier table at a case-sensitive institutional URL after earlier download failures; an unavailable link was not an unavailable scientific record. Conversely, a free-parameter carbon search stops prematurely and fits training data worse than its nested restricted model. A second solver passes manufactured tests but fails to converge on real observations. Both failures are preserved. A general numerical-library solver, using the unchanged independent scientific residuals, supplies an additional audit. An equivalent carbonate root solution reduces computational cost after direct comparison with the original solver. These cases require source, convergence and performance audits in addition to mechanism tests. The recovered expert-inversion likelihood also shows why a 50% expert budget mass cannot be assumed to yield 50% posterior budget probability. Each correction is preserved in the assumption register.

The research log records source access failures, incidental search exposure, interpretation decisions, and verification outcomes. No subagents were used. The work followed Brad Lipovsky's supplied research instructions and academic style guide. This assessment concerns the recorded execution; it is not a general benchmark of all AI research systems.

16 Reproducibility and current scope

Run `bash scripts/run_all.sh` from the project root to execute the offline verification, generate numerical diagnostics and figures, and compile this report. Julia uses standard libraries only. Python 3.6.8, NumPy 1.19.5, SciPy 1.5.4 and Matplotlib 3.3.4 were available locally; pdfLaTeX comes from TeXLive 2018. The driver verifies executable availability and imports, requires no installation, and

limits numerical libraries to one thread. Calculations run at scheduling priority 19 on a heavily loaded shared server. All artifacts, temporary files and tool caches are directed into the designated project. Source acquisition is a separate optional, hash-checked command. The scientific build does not require network access or downloaded authors' outputs.

The independent calculations verify response and cost mechanisms, quantify written-specification discrepancies, and test alternative observation-conditioned climate scenarios. An original conceptual Figure 4 is complete. Main Figures 1–3, supplementary Figures 1–10 and modeled Table 1 remain unreproduced. A printed-table arithmetic agreement is not promoted to numerical-model agreement. The missing calibration and coastal inputs must be resolved before an assessment of the target scientific conclusions can be claimed. These remaining tasks do not invalidate the completed mechanism tests or authorize treating this progress report as project completion.

References

[1] SciPy Developers (2020). *SciPy 1.5.4 reference guide: scipy.optimize.least_squares*. Versioned numerical documentation.

[2] Roquet, F., Madec, G., Brodeau, L., and Nycander, J. (2015). Defining a simplified yet "realistic" equation of state for seawater. *Journal of Physical Oceanography*, 45, 2564–2579. doi:10.1175/JPO-D-15-0080.1.

[3] Wong, T. E., Dake, S., Feke, K., Darnell, C., and Srikrishnan, V. (2025). Coastal adaptation and damage costs at different global warming thresholds. *npj Natural Hazards*, 2, 35. doi:10.1038/s44304-025-00089-0. Includes seven-page supplementary information.

[4] Darnell, C., and coauthors (2024 manuscript revision). The interplay of future emissions and geophysical uncertainties for future sea level rise. doi:10.31219/osf.io/j47ts. Original scholarly file kt9av, revision 9, deposited 11 December 2024; versioned manuscript. The citation's exact manuscript revision is unresolved.

[5] Urban, N. M., and Keller, K. (2010). Probabilistic hindcasts and projections of the coupled climate, carbon cycle and Atlantic meridional overturning circulation system: a Bayesian fusion of century-scale observations with a simple model. *Tellus A*, 62,737–750. doi:10.1111/j.1600-0870.2010.00471.x.

[6] Ricciuto, D. M., Davis, K. J., and Keller, K. (2008). A Bayesian calibration of a simple carbon cycle model: the role of observations in estimating and reducing uncertainty. *Global Biogeochemical Cycles*, 22. doi:10.1029/2006GB002908.

[7] Wong, T. E., and coauthors (2017). BRICK v0.2, a simple, accessible, and transparent model framework for climate and regional sea-level projections. *Geoscientific Model Development*, 10, 2741–2760. doi:10.5194/gmd-10-2741-2017.

[8] Wong, T. E., and coauthors (2022). MimiBRICK.jl: a Julia package for the BRICK model for sea-level change in the Mimi integrated modeling framework. *Journal of Open Source Software*. doi:10.21105/joss.04556. Written article consulted; software was not consulted.

[9] Shaffer, G. (2014). Formulation, calibration and validation of the DAIS model (version 1), a simple Antarctic ice sheet model sensitive to variations of sea level and ocean subsurface temperature. *Geoscientific Model Development*, 7, 1803–1818. doi:10.5194/gmd-7-1803-2014.

[10] Diaz, D. B. (2016). Estimating global damages from sea level rise with the Coastal Impact and Adaptation Model (CIAM). *Climatic Change*. doi:10.1007/s10584-016-1675-4. Mathematical specification in supplementary information, pp.9–17.

[11] Vega-Westhoff, B., Sriver, R. L., Hartin, C. A., Wong, T. E., and Keller, K. (2019). Impacts of observational constraints related to sea level on estimates of climate sensitivity. *Earth's Future*, 7,677–690. doi:10.1029/2018EF001082.

[12] Wong, T. E., and coauthors (2022). Sea level and socioeconomic uncertainty drives high-end coastal adaptation costs. *Earth's Future*. doi:10.1029/2022EF003061.

[13] European Environment Agency, Climate-ADAPT. Dynamic Interactive Vulnerability Assessment model (DIVA), provider directory, accessed 12 October 2026. Provider record.

[14] Muis, S., Verlaan, M., Winsemius, H. C., Aerts, J. C. J. H., and Ward, P. J. (2016). A global reanalysis of storm surges and extreme sea levels. *Nature Communications*, 7, 11969. doi:10.1038/ncomms11969. Dataset provenance checked at Vrije Universiteit Amsterdam and 4TU.

[15] IIASA. Shared Socioeconomic Pathways Scenario Database. Provider documentation, accessed 12 October 2026. Multiple GDP interpretations and data vintages require explicit selection before a target reconstruction.

[16] Tanaka, K. (2008). Inverse estimation for the simple Earth system model ACC2. Doctoral thesis. Chapter 2 supplies physical equations; software appendices were not consulted. Scholarly thesis.

[17] Siegenthaler, U., and Oeschger, H. (1987). Biospheric CO_2 emissions during the past 200 years reconstructed by deconvolution of ice core data. *Tellus B*, 39, 140–154. Published methods and reservoir diagram.

[18] Hooss, G., and coauthors (1999). A nonlinear impulse response model of the coupled carbon cycle–climate system (NICCS). MPI scientific report, carbon chemistry section. Scholarly report.

[19] Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., and Buerkner, P.-C. (2021). Rank-normalization, folding, and localization: an improved $\widehat{R}$ for assessing convergence of MCMC. *Bayesian Analysis*. doi:10.1214/20-BA1221.

[20] Independent input providers. Met Office HadCRUT4.6.0.0 (Morice et al., 2012); CEDA radiative forcing (Smith, 2023); CMIP6 concentrations (Meinshausen et al., 2020); NOAA GML CO_2 (Lan et al.); NOAA NCEI ocean heat and thermosteric sea level (Levitus et al., 2012); Global Carbon Budget 2023 (Friedlingstein et al., 2023). Exact citations, units, licenses, subset definitions and versioned hashes are in `inputs/README.md` and `docs/EXTERNAL_INPUTS.csv`.

[21] Otosaka, I. N., and coauthors (2023). Mass balance of the Greenland and Antarctic ice sheets from 1992 to 2020. *Earth System Science Data*, 15, 1597–1616. doi:10.5194/essd-15-1597-2023. IMBIE2021 data: British Antarctic Survey archive.

[22] GlaMBIE Team (2025). Community estimate of global glacier mass changes from 2000 to 2023. *Nature*. doi:10.1038/s41586-024-08545-z. Observational data: GlaMBIE2024.

[23] Bolliger, I., and coauthors (2023). SLIIDERS v1.1 input dataset. doi:10.5281/zenodo.7693868, CC BY-NC 4.0. Methods: Depsky and coauthors (2023), written coastal input and model documentation.

[24] World Bank (2026 retrieval). World Development Indicators API v2, national population, PPP GDP and real GDP. Exact query URLs and hashes are in the external-input manifest. Historical grouping correction follows the July2025 release note.

[25] Wong, T. E., Bakker, A. M. R., and Keller, K. (2017). Impacts of Antarctic fast dynamics on sea-level projections and coastal flood defense. *Climatic Change*. Scholarly manuscript and written calibration constraints.

[26] Jouzel, J., and coauthors (2007). Orbital and millennial Antarctic climate variability over the past800,000 years. *Science*. doi:10.1126/science.1141038.

[27] Spratt, R. M., and Lisiecki, L. E. (2016). A Late Pleistocene sea level stack. *Climate of the Past*,12,1079–1092. doi:10.5194/cp-12-1079-2016.

[28] Mann, M. E., and coauthors (2008). Proxy-based reconstructions of hemispheric and global surface temperature variations over the past two millennia. *PNAS*. doi:10.1073/pnas.0805721105.

[29] Church, J. A., and White, N. J. (2011). Sea-level rise from the late19th to the early21st century. *Surveys in Geophysics*. doi:10.1007/s10712-011-9119-1.

[30] Waelbroeck, C., and coauthors (2002). Sea-level and deep water temperature changes derived from benthic foraminifera isotopic records. *Quaternary Science Reviews*. doi:10.1016/S0277-3791(01)00101-9. Only independent proxy sea-level data and documented chronologies are used.

[31] Nicholls, Z., and coauthors (2020–2021). Reduced Complexity Model Intercomparison Project, protocol5.1.0. Versioned exogenous input archive. Only input emissions are used; model results and implementations are not consulted.

[32] Fox-Kemper, B., and coauthors (2021). Ocean, Cryosphere and Sea Level Change. In *Climate Change 2021: The Physical Science Basis*, Chapter 9, Table 9.9. IPCC AR6 WGI Chapter 9. Printed summary ranges are comparison evidence only.

[33] Darnell, C., and coauthors (2025). The interplay of future emissions and geophysical uncertainties for projections of sea-level rise. September 2025 version; supplementary Table 2 and written methods only. Archived written manuscript.

[34] Dyurgerov, M. B., and Meier, M. F. (2005). Glaciers and the changing Earth system: a 2004 snapshot. INSTAAR Occasional Paper 58, Appendix 3, p.117. `https://www.colorado.edu/instaar/sites/default/files/attached-files/OP58_dyurgerov_meier.pdf`.

[35] Dyurgerov, M. (2002, updated 2005). Glacier Mass Balance and Regime Measurements and Analysis, 1945–2003. G10002, Version 1, NSIDC; Supplement 2005, Appendix 6. `https://doi.org/10.7265/N52N506F`. The cited report's written abstract is cataloged at `https://www.colorado.edu/instaar/research/publications/instaar-occasional-papers`.

[36] Fuller, R. W., Wong, T. E., and Keller, K. (2017). Probabilistic inversion of expert assessments to inform projections about Antarctic ice sheet responses. PLOS ONE, 12, e0190115. `https://doi.org/10.1371/journal.pone.0190115`.

[37] Gouretski, V., and Koltermann, K. P. (2007). How much is the ocean really warming? *Geophysical Research Letters*, 34, L01610. doi:10.1029/2006GL027834.